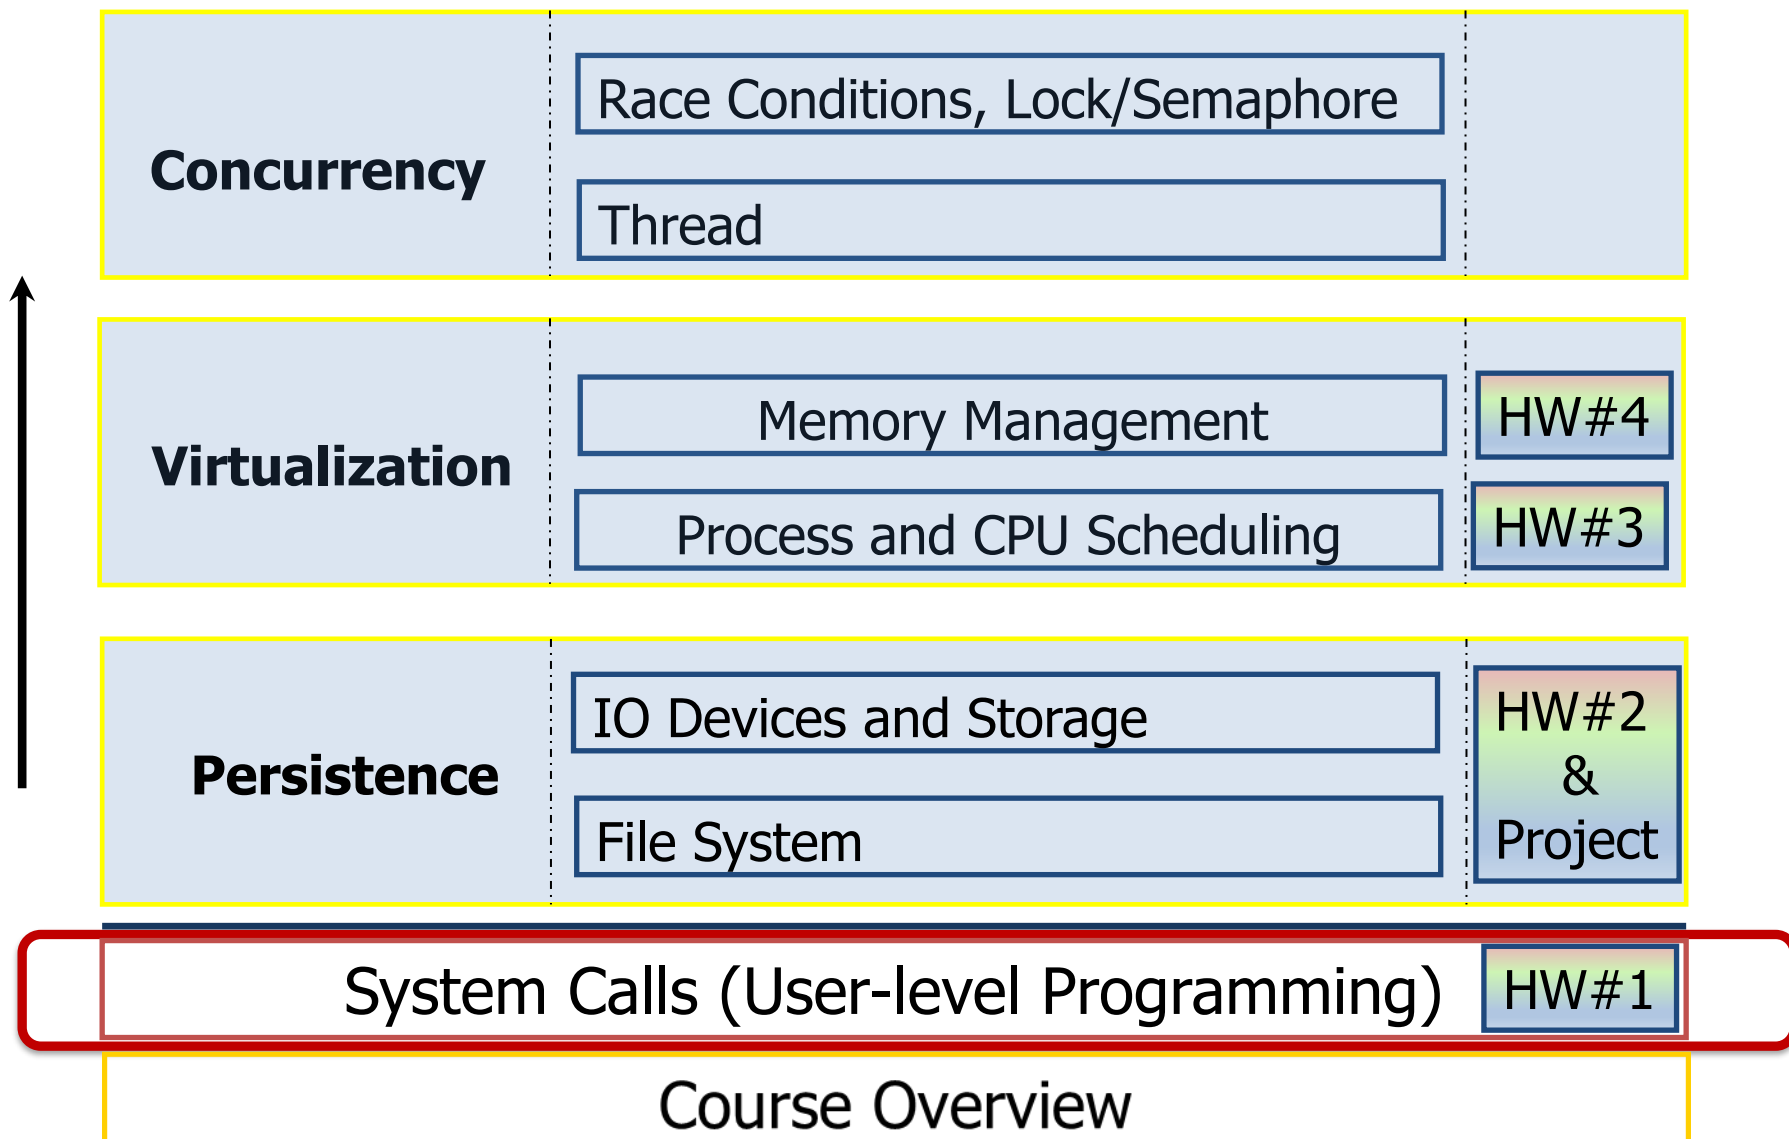


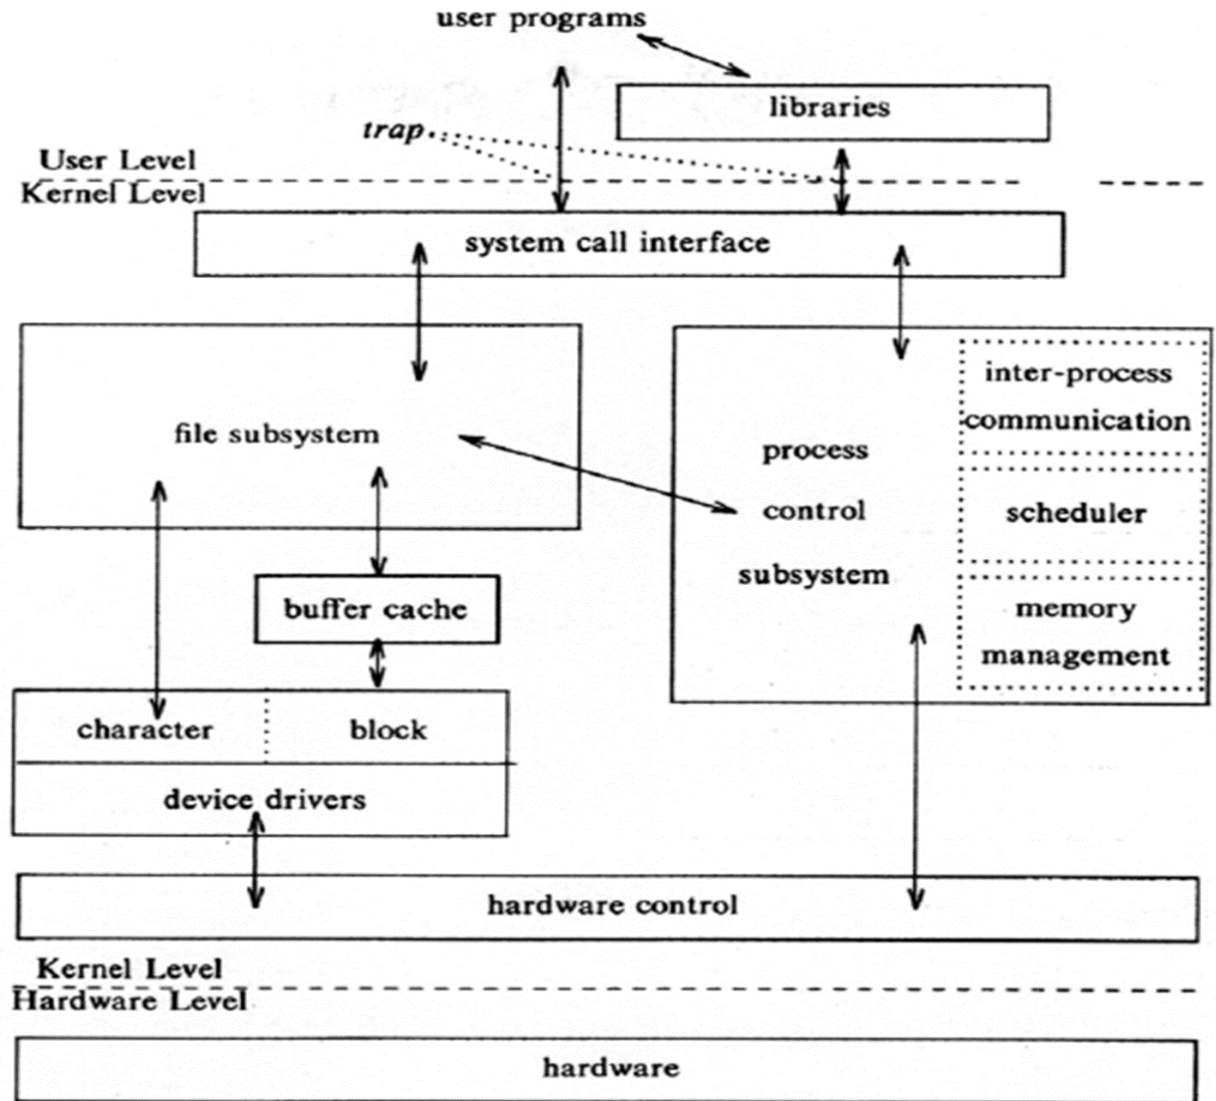
Lecture 4: User-level Programming via System Calls (File & Directory)

The Course Organization (Bottom-up)



System call

OS provides services via **System Call** (typically a few hundred) to run **process**, access memory/**devices/files**, etc.



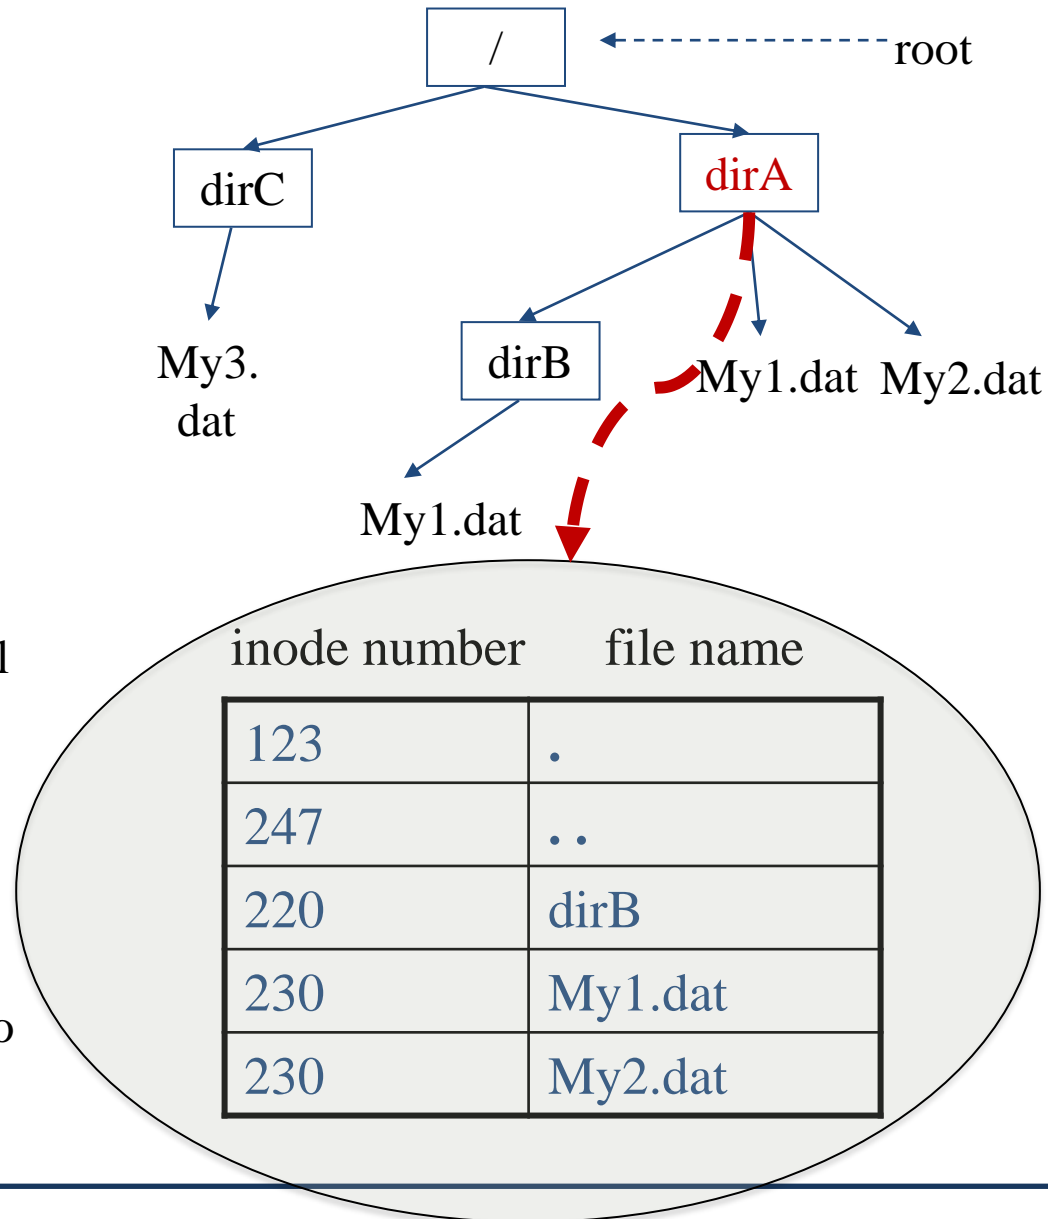
The Design Of The Unix Operating System (Maurice Bach, 1986)

Persistent Storage

- ❑ Keep a data **intact** even if there is a power loss.
 - ◆ Hard disk drive
 - ◆ Solid-state storage device
- ❑ Two key abstractions in the virtualization of storage
 - ◆ File
 - ◆ Directory

File and Directory

- ❑ File – A container to contain data of a file (a linear array of bytes)
 - ◆ Each file has low-level name (**inode number**)
- ❑ Directory – Implement directory tree (directory hierarchy)
 - ◆ Like a file, it also has a low-level name (**inode number**).
 - ◆ It contains a list of (**file name**, **inode number**) pairs.
 - ◆ Each entry in a directory refers to either *files* or other *directories*.



Creating Files

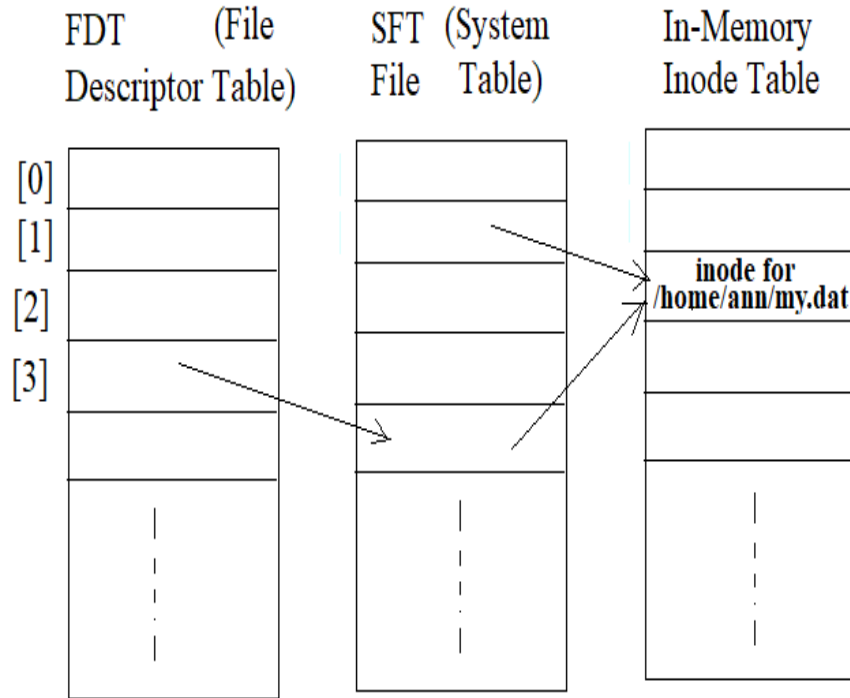
- Use `open()` system call with `O_CREAT` flag.

```
int myfd = open("/home/ann/my.dat", O_CREAT | O_WRONLY | O_TRUNC);
```

- `O_CREAT` : create file.
 - `O_WRONLY` : only write to that file while opened.
 - `O_TRUNC` : make the file size zero (remove any existing content).
- `open()` system call returns **file descriptor**.
 - ◆ *File descriptor* is an integer, and is used to access files.
- `read` (file descriptor, buffer pointer, the number of bytes to read from)
 - ◆ Return the number of bytes it read
- `write` (file descriptor, buffer pointer, the number of bytes to write to)
 - ◆ Return the number of bytes it write

The Story behind open() (Unix System V)

```
int myfd = open("/home/ann/my.dat", O_CREAT | O_WRONLY | O_TRUNC);
```



myfd (the file descriptor for the newly-created file) is 3

When a process is created, file descriptor 0, 1, 2, are opened by default for standard input/output/error.

- A file descriptor specifies the index into file descriptor table (FDT) of the process.
- Entries of FDT contain pointers to entries in system file table (SFT).
- When a file is opened, an entry is created in both EDT and SFT.
- SFT entry contains information about whether a file is open for read or write, protection, and lock, the file offset, where the next data is read from or written to in the file, etc.
- Several entries in SFT may point to the same physical file

Reading And Writing, But Not Sequentially

- ❑ An open file has a **current offset**.
 - ◆ Determine **where** the next read or write will begin reading from or writing to within the file.
- ❑ Update the current offset
 - ◆ **Implicitly**: A read or write of N bytes takes place, N is added to the current offset.
 - ◆ **Explicitly**: `lseek()`

Reading And Writing, But Not Sequentially (Cont.)

```
off_t lseek(int fd, off_t offset, int whence);
```

- ◆ `fd` : File descriptor
- ◆ `offset` : Position the file offset to a particular location within the file
- ◆ `whence` : Determine how the seek is performed

From the man page:

If `whence` is `SEEK_SET`, the offset is set to offset bytes.
If `whence` is `SEEK_CUR`, the offset is set to its current location plus offset bytes.
If `whence` is `SEEK_END`, the offset is set to the size of the file plus offset bytes.

I/O redirection

- ❑ To access a file, a process uses a *file descriptor*, which is an index into the process *file descriptor table*, which in turn points to an entry in the *system file table*.
- ❑ *Redirection* means that the process modifies its file descriptor table entry so that it points to a different entry in the system file table.
 - ◆ Consider the command `cat`, which reads from a file and echoes to standard output.

The following command redirects standard output to **my.file**

`ls -l > my.file`

Use “dup()” to implement “redirection”

```
int dup(int fd);
```

DESCRIPTION

dup() is a “smart” function that can duplicate the file descriptor, “fd”, to the lowest-numbered unused file descriptor in the file descriptor table.

RETURN VALUE:

On success, return the new file descriptor.

On error, -1 is returned, and errno is set to indicate the error.

Use “dup()” to implement “redirection”

```
int dup(int fd)
```

- **dup()** is a “smart” function that can duplicate the file descriptor, “fd”, to the **lowest-numbered unused file descriptor** in the file descriptor table.

```
#include <sys/types.h>
#include <sys/wait.h>
#include <unistd.h>
#include <stdio.h>
#include <stdlib.h>
#include <sys/stat.h>
#include <fcntl.h>

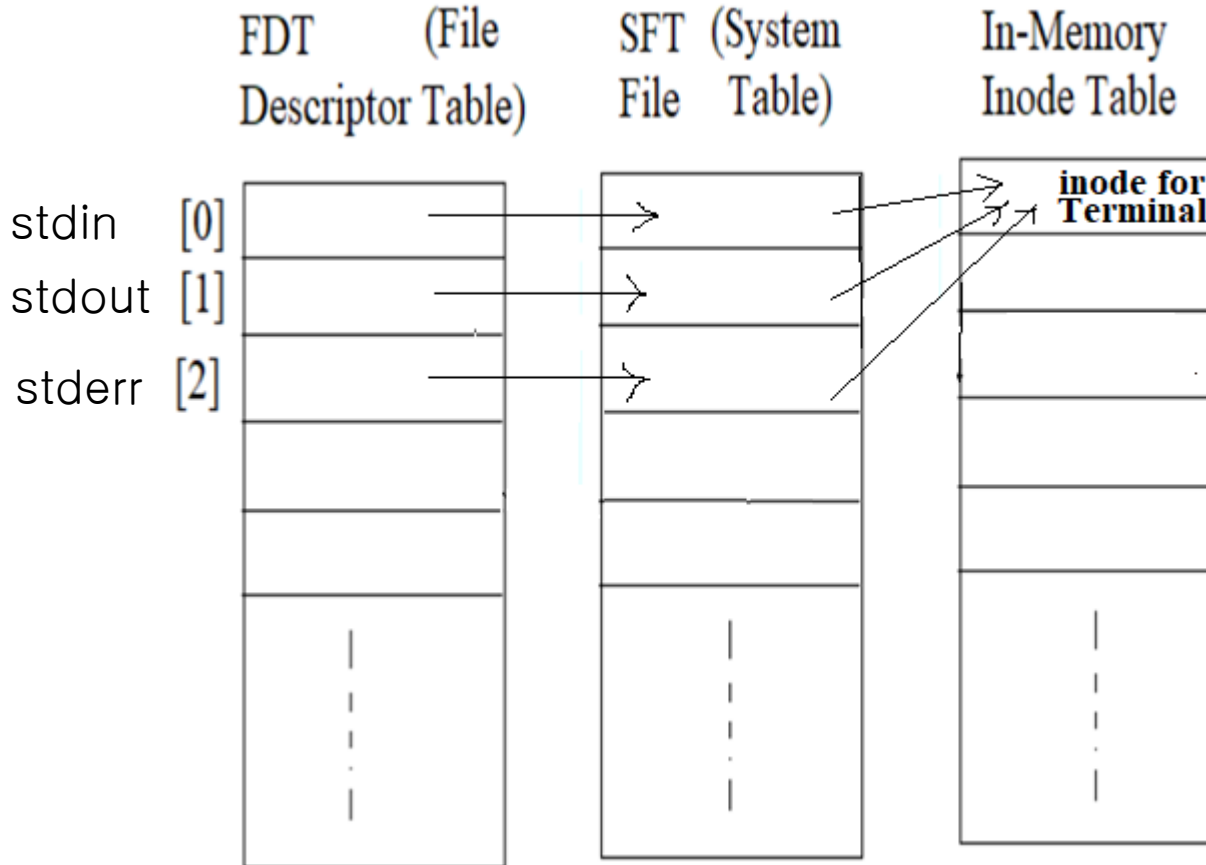
int main(void)
{
    int fd;

    fd = open("my.file", O_CREAT | O_TRUNC | O_WRONLY, S_IRUSR | S_IWUSR );
    close(1);
    dup(fd);
    close(fd);

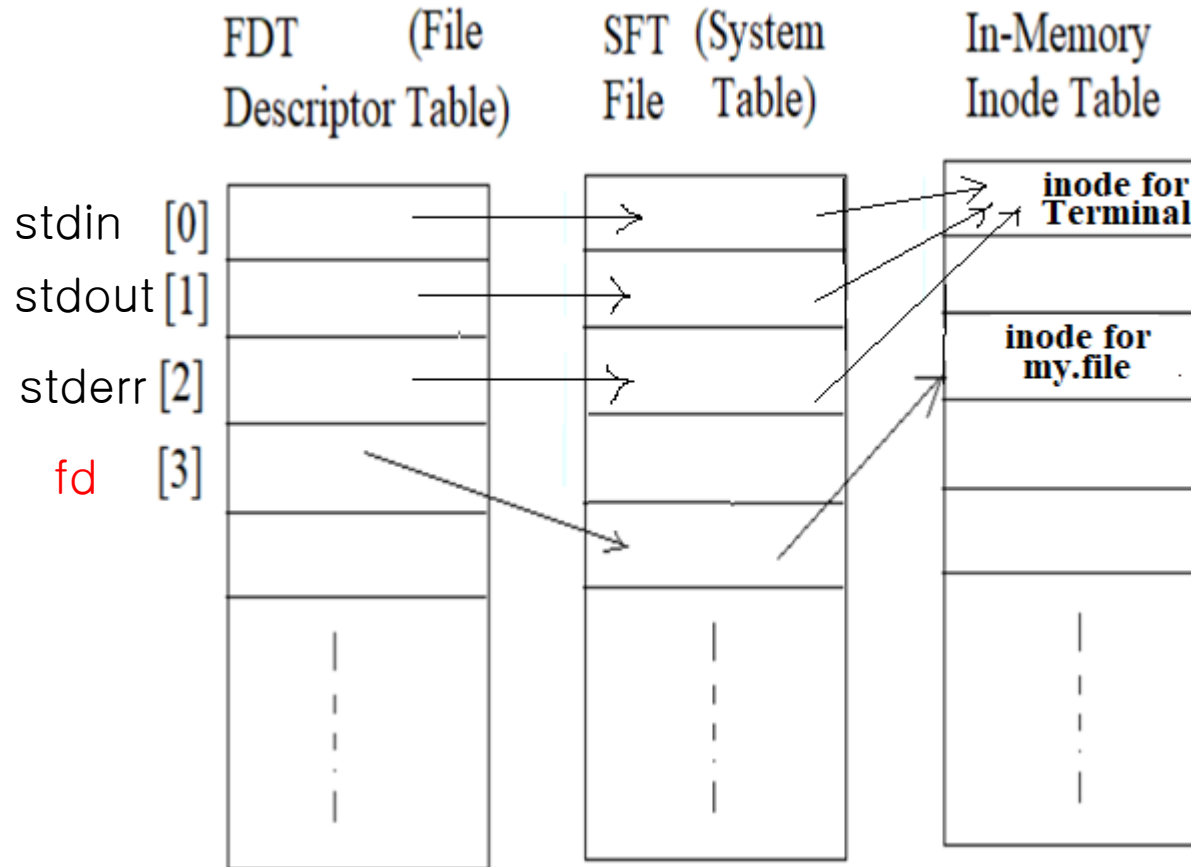
    char *cmd = "ls";
    char *argv[3];
    argv[0] = "ls";    argv[1] = "-l";    argv[2] = NULL;

    execvp(cmd, argv);
}
```

I/O redirection

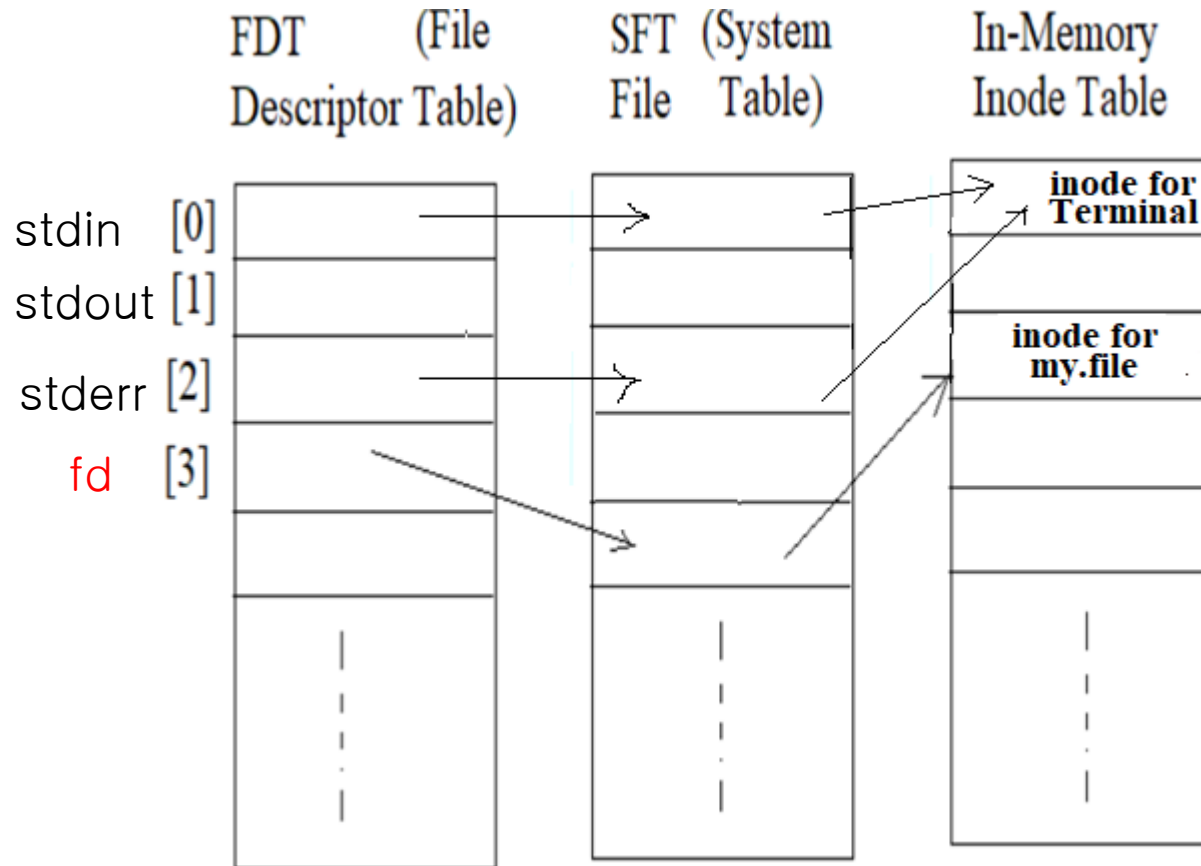
[illegible]

I/O redirection (continue)



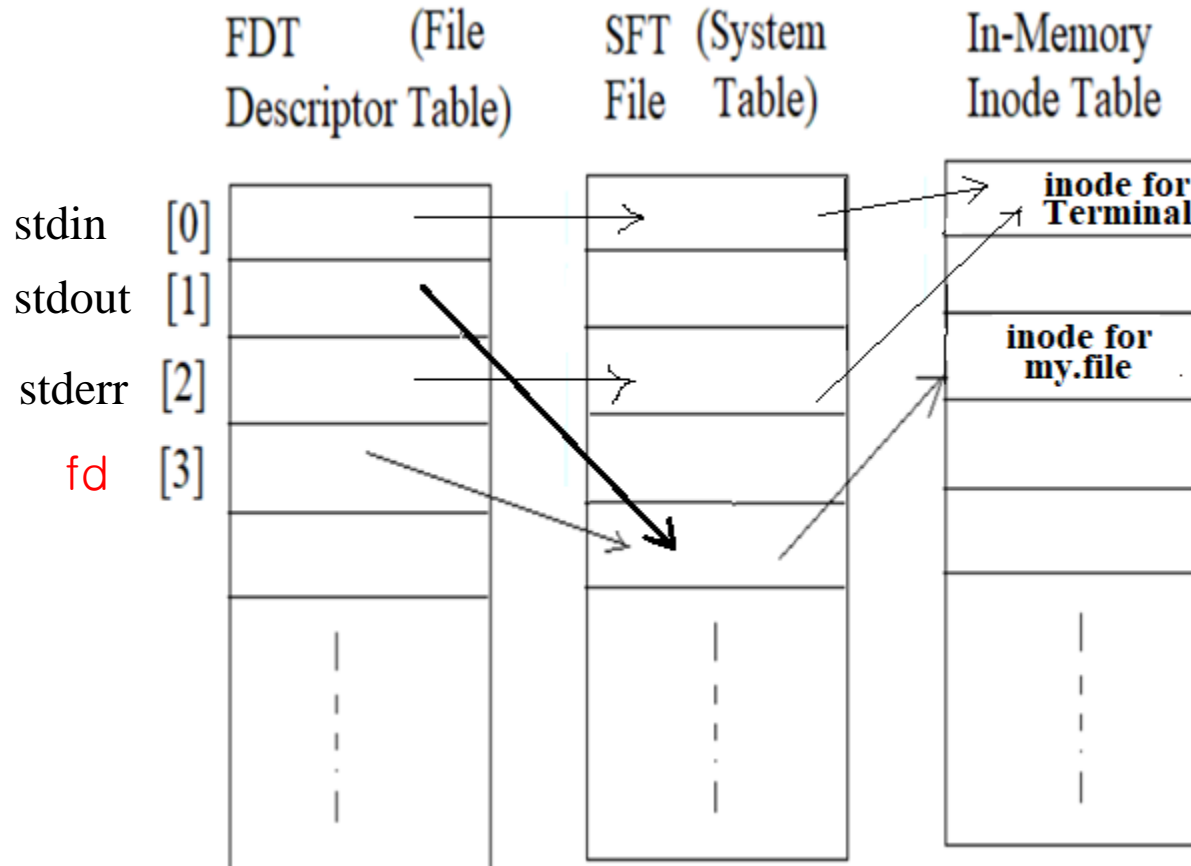
→ **`fd = open("my.file", O_CREAT, S_IRUSR|S_IWUSR); /* Create the file – my.file */`**
`close(1); /* Close 1 (Standard output)*/`
`dup(fd); /* Duplicate fd */`
`execvp(cmd, argv); /* Execute the command ls -l */`

I/O redirection (continue)



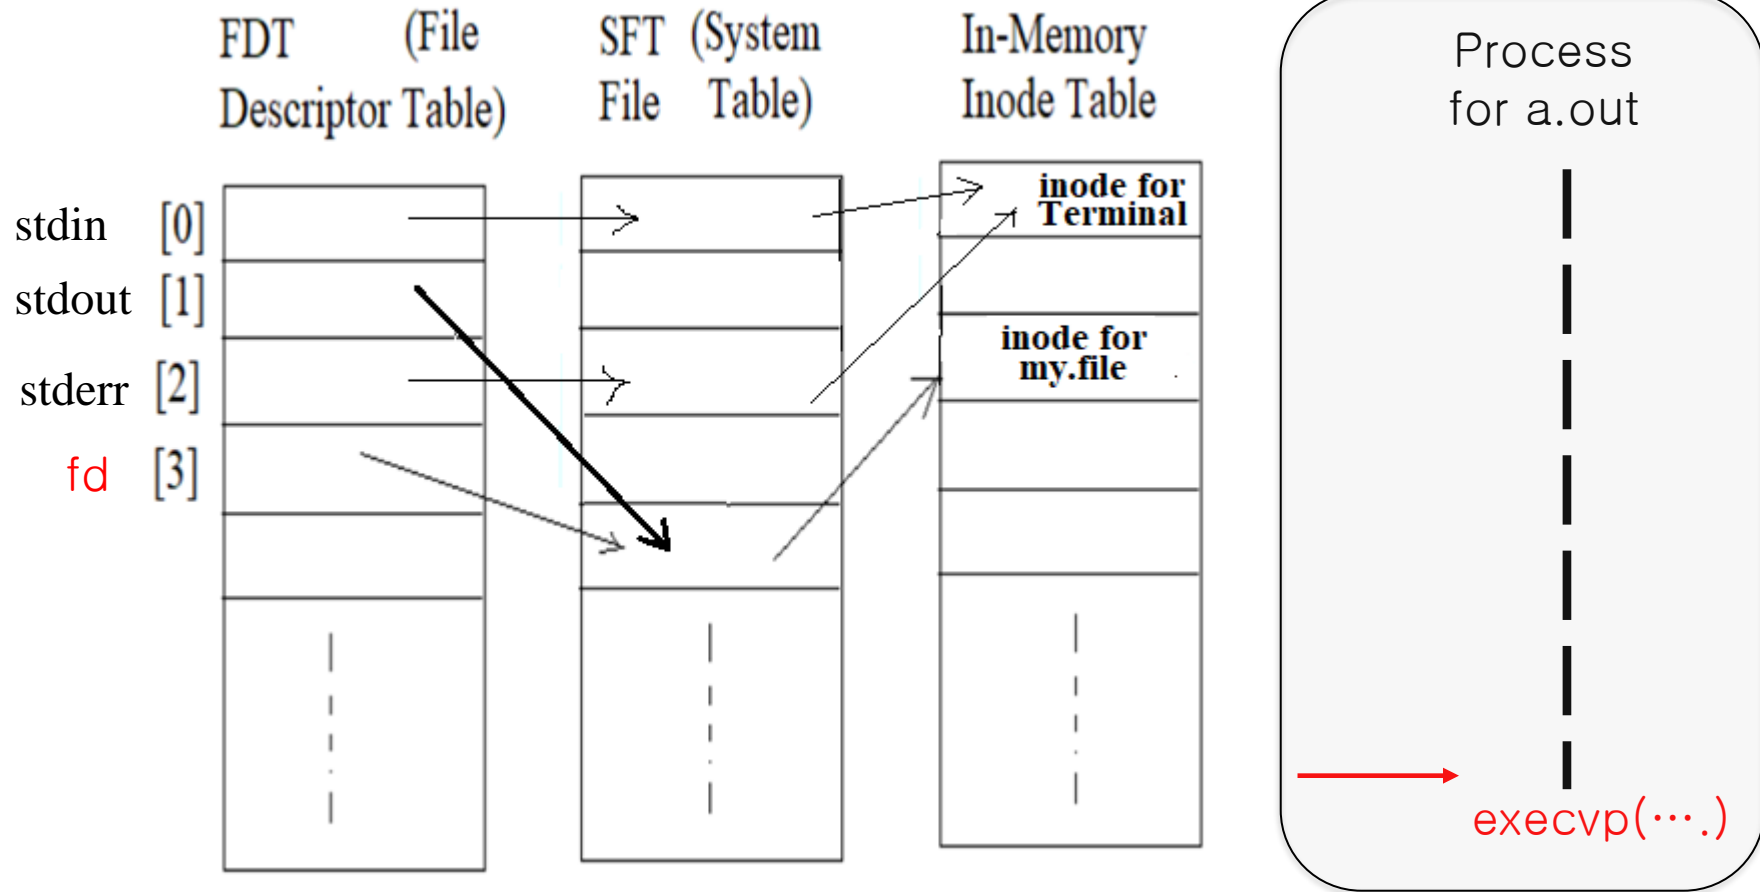
```
fd = open("my.file", O_CREAT, S_IRUSR|S_IWUSR); /* Create the file – my.file */  
→ close(1); /* Close 1 (Standard output) */  
dup(fd); /* Duplicate fd */  
execvp(cmd, argv); /* Execute the command ls -l */
```

I/O redirection (continue)



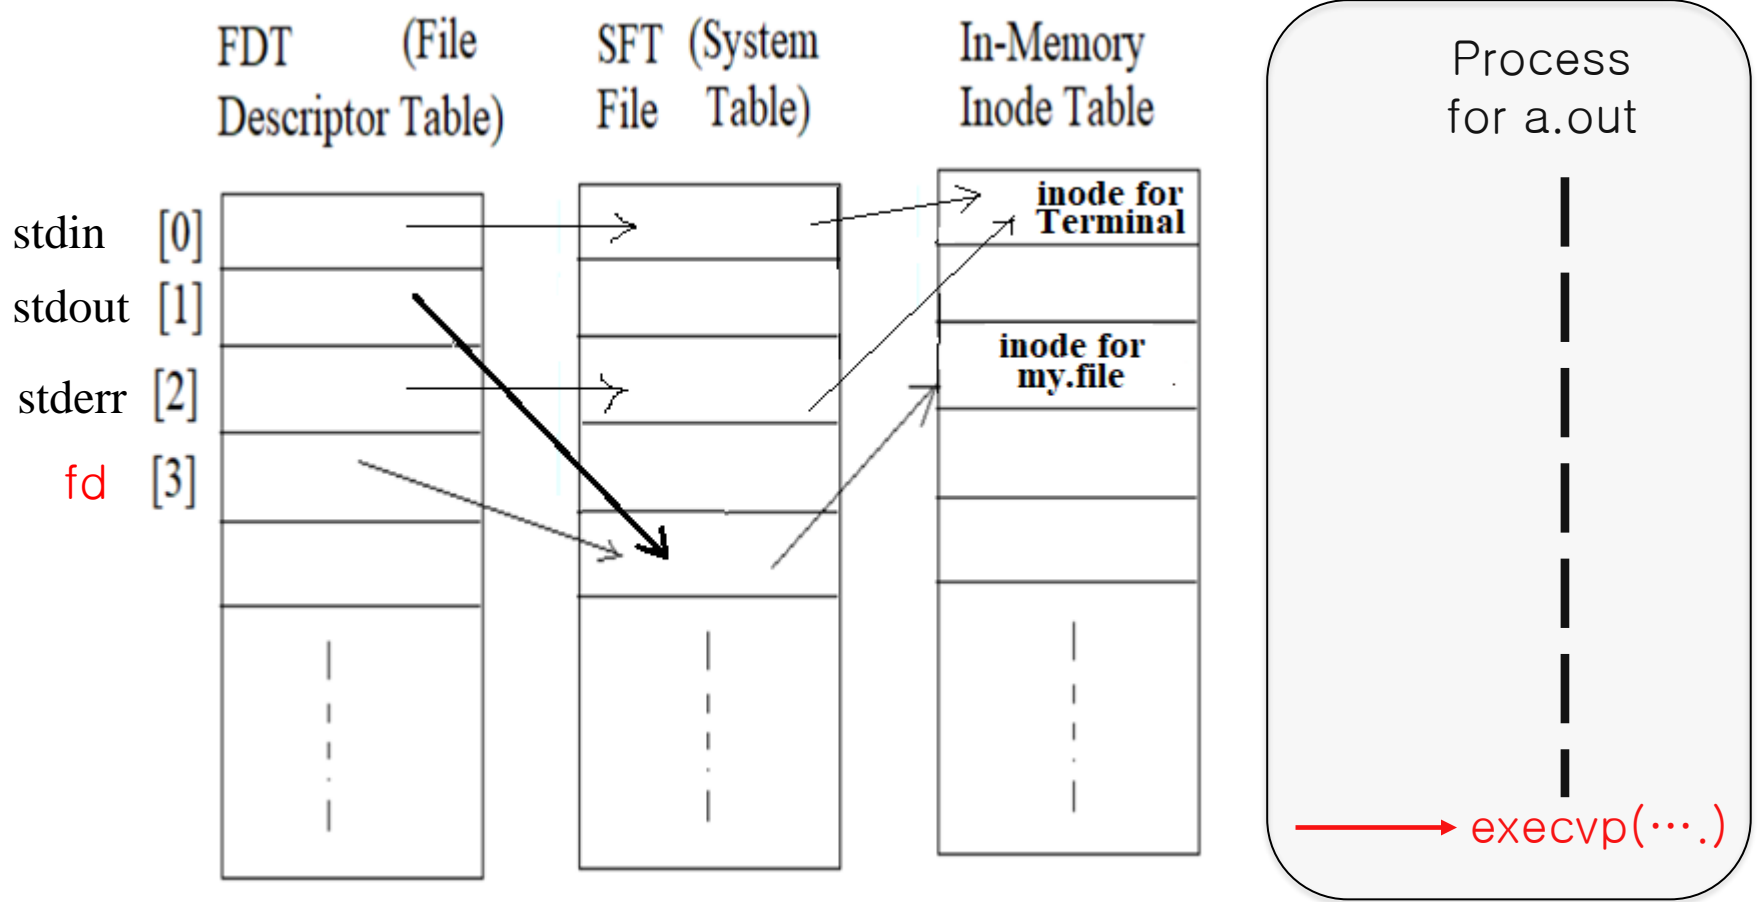
```
fd = open("my.file", O_CREAT, S_IRUSR|S_IWUSR); /* Create the file – my.file */
close(1); /* Close 1 (Standard output)*/
→ dup(fd); /* Duplicate fd */
execvp(cmd, argv); /* Execute the command ls -l */
```


I/O redirection (continue; **before execvp(...)**)

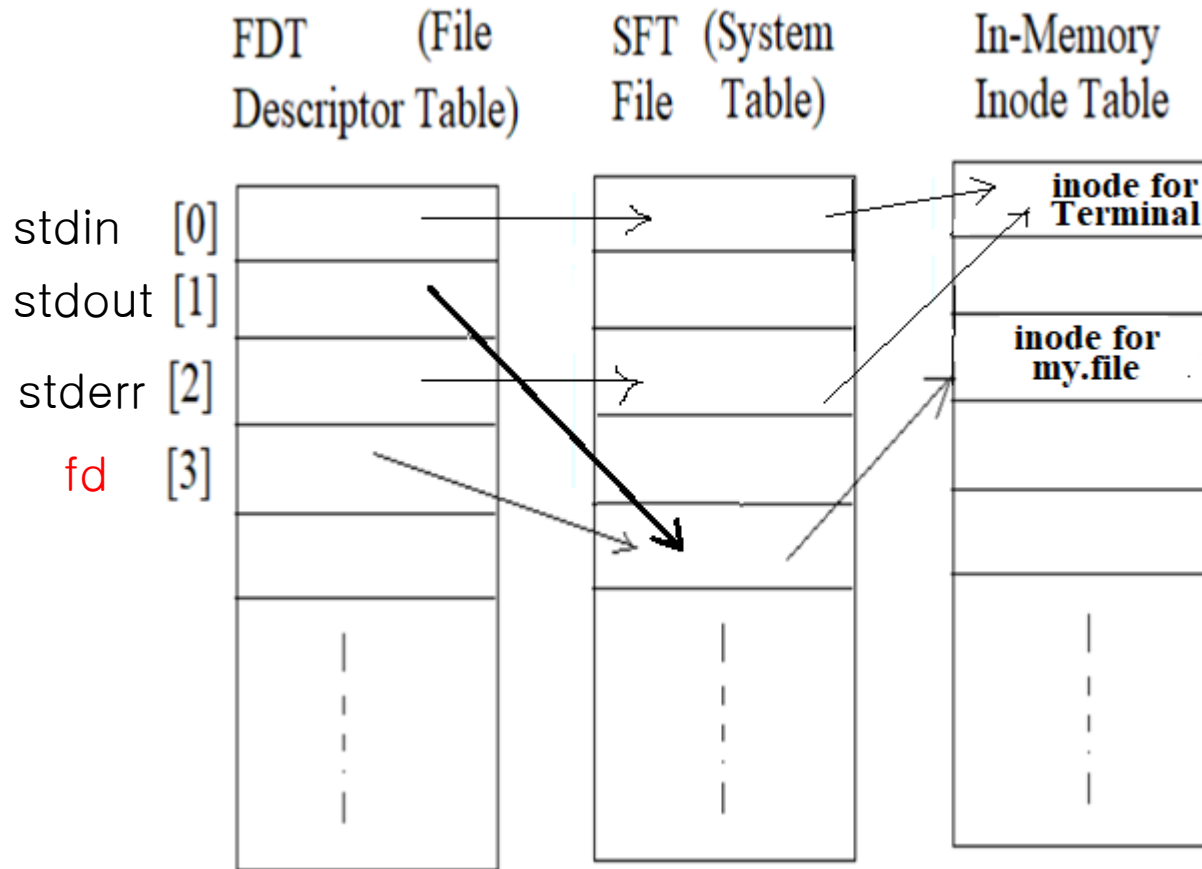


```
fd = open("my.file", O_CREAT, S_IRUSR|S_IWUSR); /* Create the file – my.file */
close(1); /* Close 1 (Standard output)*/
dup(fd); /* Duplicate fd */
→ execvp(cmd, argv); /* Execute the command ls -l */
```

I/O redirection (continue; **executing `execvp(...)`**)

[illegible]

I/O redirection (continue; after `execvp(...)`)



Process for “ls -l”
(totally replace the
body for Process a.out)


Note that
stdout of
Process
for “ls -l”
has been
redirected
to my.file

```
fd = open("my.file", O_CREAT, S_IRUSR|S_IWUSR); /* Create the file – my.file */
close(1); /* Close 1 (Standard output)*/
dup(fd); /* Duplicate fd */
→ execvp(cmd, argv); /* Execute the command ls -l */
```

Communication between parent/child processes via pipe

- System call `pipe()` returns two file descriptors by which we can access the input/output of a pipe (an I/O mechanism)

```
int fd[2];  
  
int pipe(int fd[2]);  
    return: 0 success; -1 error
```



Question: How to implement “*ls -l / wc -l*”?

Implement “ls -l | wc -l”

```
#include <sys/types.h>
#include <sys/wait.h>
#include <unistd.h>
#include <stdio.h>
#include <stdlib.h>
#include <sys/stat.h>
#include <fcntl.h>


main()
{
    int fd[2];
    int ret;
    char *cmd;
    char *argv[3];

    pipe(fd);
    if ( (ret=fork()) > 0 ){
        /* Parent process*/
        close(1);
        dup(fd[1]); close(fd[0]); close(fd[1]);

        cmd = "ls"; argv[0] = "ls"; argv[1] = "-l"; argv[2] = NULL;
        execvp(cmd, argv);
    } else if (ret == 0 ){
        /* Child process*/
        close(0);
        dup( fd[0]); close( fd[0]); close ( fd[1]);

        cmd = "wc"; argv[0] = "wc"; argv[1] = "-l"; argv[2] = NULL;
        execvp(cmd, argv);
    } else{
        /* Error in fork()*/
        printf("Error occurs when executing fork().\n");
        exit(-1);
    }
}
```

Parent Process
for a.out



Implement “ls -l | wc -l”

```
#include <sys/types.h>
#include <sys/wait.h>
#include <unistd.h>
#include <stdio.h>
#include <stdlib.h>
#include <sys/stat.h>
#include <fcntl.h>


main()
{
    int fd[2];
    int ret;
    char *cmd;
    char *argv[3];

    pipe(fd);
    if ( (ret=fork()) > 0 ){
        /* Parent process*/
        close(1);
        dup(fd[1]); close(fd[0]); close(fd[1]);

        cmd = "ls"; argv[0] = "ls"; argv[1] = "-l"; argv[2] = NULL;
        execvp(cmd, argv);
    } else if (ret == 0 ){
        /* Child process*/
        close(0);
        dup( fd[0]); close( fd[0]); close ( fd[1]);

        cmd = "wc"; argv[0] = "wc"; argv[1] = "-l"; argv[2] = NULL;
        execvp(cmd, argv);
    } else{
        /* Error in fork()*/
        printf("Error occurs when executing fork().\n");
        exit(-1);
    }
}
```

Parent Process
for a.out



Implement “ls -l | wc -l”


```
#include <sys/types.h>
#include <sys/wait.h>
#include <unistd.h>
#include <stdio.h>
#include <stdlib.h>
#include <sys/stat.h>
#include <fcntl.h>

main()
{
    int fd[2];
    int ret;
    char *cmd;
    char *argv[3];

    pipe(fd);
    if ( (ret=fork()) > 0 ){
        /* Parent process*/
        close(1);
        dup(fd[1]); close(fd[0]); close(fd[1]);

        cmd = "ls"; argv[0] = "ls"; argv[1] = "-l"; argv[2] = NULL;
        execvp(cmd, argv);
    } else if (ret == 0 ){
        /* Child process*/
        close(0);
        dup( fd[0]); close( fd[0]); close ( fd[1]);

        cmd = "wc"; argv[0] = "wc"; argv[1] = "-l"; argv[2] = NULL;
        execvp(cmd, argv);
    } else{
        /* Error in fork()*/
        printf("Error occurs when executing fork().\n");
        exit(-1);
    }
}
```



Implement “ls -l | wc -l”

```
#include <sys/types.h>
#include <sys/wait.h>
#include <unistd.h>
#include <stdio.h>
#include <stdlib.h>
#include <sys/stat.h>
#include <fcntl.h>

main()
{
    int fd[2];
    int ret;
    char *cmd;
    char *argv[3];

    pipe(fd);
    if ( (ret=fork()) > 0 ){
        /* Parent process*/
        close(1);
        dup(fd[1]); close(fd[0]); close(fd[1]);

        cmd = "ls"; argv[0] = "ls"; argv[1] = "-l"; argv[2] = NULL;
        execvp(cmd, argv);
    }else if (ret == 0 ){
        /* Child process*/
        close(0);
        dup( fd[0]); close( fd[0]); close ( fd[1]);

        cmd = "wc"; argv[0] = "wc"; argv[1] = "-l"; argv[2] = NULL;
        execvp(cmd, argv);
    } else{
        /* Error in fork()*/
        printf("Error occurs when executing fork().\n");
        exit(-1);
    }
}
```

Parent Process
for a.out

Child Process
for a.out

fork()

fork()

execvp("ls")

execvp("wc")

Implement “ls -l | wc -l”

```
#include <sys/types.h>
#include <sys/wait.h>
#include <unistd.h>
#include <stdio.h>
#include <stdlib.h>
#include <sys/stat.h>
#include <fcntl.h>

main()
{
    int fd[2];
    int ret;
    char *cmd;
    char *argv[3];


    pipe(fd);
    if ( (ret=fork()) > 0 ){
        /* Parent process*/
        close(1);
        dup(fd[1]); close(fd[0]); close(fd[1]);

        cmd = "ls"; argv[0] = "ls"; argv[1] = "-l"; argv[2] = NULL;
        execvp(cmd, argv);
    }else if (ret == 0 ){
        /* Child process*/
        close(0);
        dup( fd[0]); close( fd[0]); close ( fd[1]);

        cmd = "wc"; argv[0] = "wc"; argv[1] = "-l"; argv[2] = NULL;
        execvp(cmd, argv);
    } else{
        /* Error in fork()*/
        printf("Error occurs when executing fork().\n");
        exit(-1);
    }
}
```

Parent Process
for ls

Child Process
for wc



Implement “ls -l | wc -l”

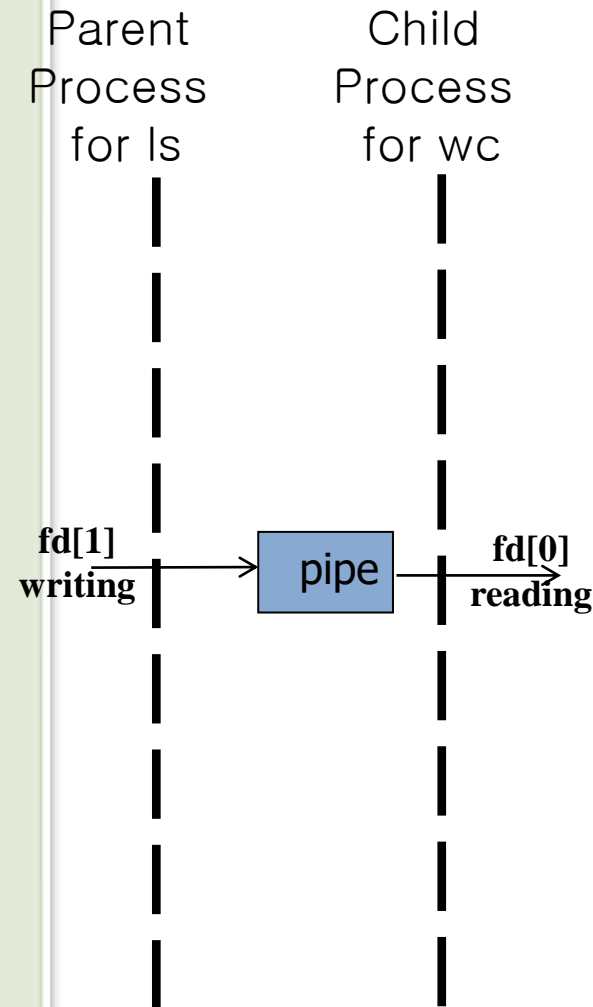
```
#include <sys/types.h>
#include <sys/wait.h>
#include <unistd.h>
#include <stdio.h>
#include <stdlib.h>
#include <sys/stat.h>
#include <fcntl.h>

main()
{
    int fd[2];
    int ret;
    char *cmd;
    char *argv[3];

    pipe(fd);
    if ( (ret=fork()) > 0 ){
        /* Parent process*/
        close(1);
        dup(fd[1]); close(fd[0]); close(fd[1]);

        cmd = "ls"; argv[0] = "ls"; argv[1] = "-l"; argv[2] = NULL;
        execvp(cmd, argv);
    }else if (ret == 0 ){
        /* Child process*/
        close(0);
        dup( fd[0]); close( fd[0]); close ( fd[1]);

        cmd = "wc"; argv[0] = "wc"; argv[1] = "-l"; argv[2] = NULL;
        execvp(cmd, argv);
    } else{
        /* Error in fork()*/
        printf("Error occurs when executing fork().\n");
        exit(-1);
    }
}
```



Implement “ls -l | wc -l”

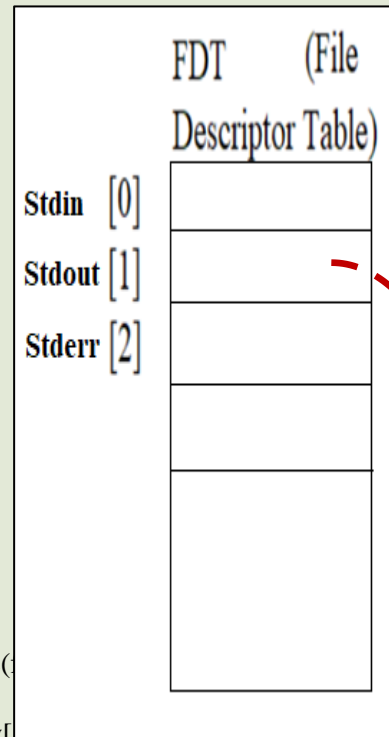
```
#include <sys/types.h>
#include <sys/wait.h>
#include <unistd.h>
#include <stdio.h>
#include <stdlib.h>
#include <sys/stat.h>
#include <fcntl.h>
```

```
main()
{
    int fd[2];
    int ret;
    char *cmd;
    char *argv[3];

    pipe(fd);
    if ( (ret=fork()) > 0 ){
        /* Parent process*/
        close(1);
        dup(fd[1]); close(fd[0]); close(fd[1]);

        cmd = "ls"; argv[0] = "ls"; argv[1] = "-l"; argv[2] = NULL;
        execvp(cmd, argv);
    } else if (ret == 0 ){
        /* Child process*/
        close(0);
        dup( fd[0]); close( fd[0]); close ( fd[1]);

        cmd = "wc"; argv[0] = "wc"; argv[1] = "-l"; argv[2] = NULL;
        execvp(cmd, argv);
    } else{
        /* Error in fork()*/
        printf("Error occurs when executing fork().\n");
        exit(-1);
    }
}
```



Parent Process
for ls

fd[1]
writing

Child Process
for wc

fd[0]
reading

pipe

Implement “ls -l | wc -l”

```
#include <sys/types.h>
#include <sys/wait.h>
#include <unistd.h>
#include <stdio.h>
#include <stdlib.h>
#include <sys/stat.h>
#include <fcntl.h>
```

```
main()
{
```

```
    int fd[2];
    int ret;
    char *cmd;
    char *argv[3];
```

```
    pipe(fd);
```

```
    if ( (ret=fork()) > 0 ){
```

```
        /* Parent process*/
```

```
        close(1);
```

```
        dup(fd[1]); close(fd[0]); close(
```

```
        cmd = "ls"; argv[0] = "ls"; argv[1] = "-l"; argv[2] = NULL;
```

```
        execvp(cmd, argv);
```

```
    }else if (ret == 0 ){
```

```
        /* Child process*/
```

```
        close(0);
```

```
        dup( fd[0]); close( fd[0]); close ( fd[1]);
```

```
        cmd = "wc"; argv[0] = "wc"; argv[1] = "-l"; argv[2] = NULL;
```

```
        execvp(cmd, argv);
```

```
    } else{
```

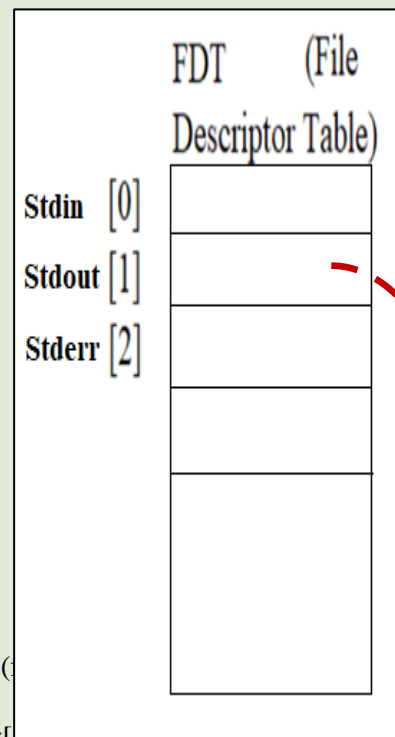
```
        /* Error in fork()*/
```

```
        printf("Error occurs when executing fork().\n");
```

```
        exit(-1);
```

```
    }
```

```
}
```



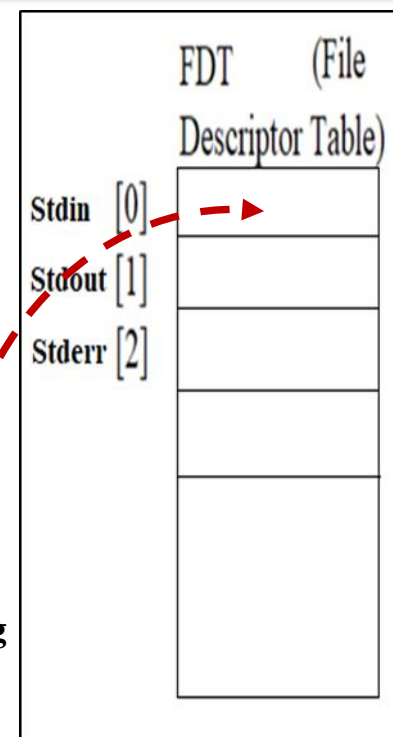
Parent Process
for ls

fd[1]
writing

Child Process
for wc

fd[0]
reading

pipe



More about Pipe (<https://man7.org/linux/man-pages/man7/pipe.7.html>)

- ❑ Pipes provide a unidirectional interprocess communication channel. A pipe has a read end and a write end. Data written to the write end of a pipe can be read from the read end of the pipe.
- ❑ If a process attempts to read from an empty pipe, then `read()` will block until data is available.
- ❑ The communication channel provided by a pipe is a byte stream: there is no concept of message boundaries.
- ❑ If all file descriptors referring to the write end of a pipe have been closed, then an attempt to `read()` from the pipe will see end-of-file (`read()` will return 0). If all file descriptors referring to the read end of a pipe have been closed, then a `write()` will cause a `SIGPIPE` signal to be generated for the calling process. If the calling process is ignoring this signal, then `write()` fails with the error `EPIPE`.

Writing Immediately with `fsync()`

- ❑ The file system will **buffer** writes in memory for some time (for performance reasons), e.g. 5 or 30 seconds.
- ❑ At that later point in time, the write(s) will **actually be issued** to the storage device.
 - ◆ Writes seem to complete quickly; but data can be lost (e.g., machine crashes).
 - ◆ However, some applications require persistence guarantee, e.g. DBMS requires force writes to disk from time to time (your bank transactions).

```
int fsync(int fd)
```

- force all dirty (i.e., not yet written) data written to disk
- `fsync()` returns once all of these writes are complete.

Writing Immediately with fsync() (Cont.)

```
#include <unistd.h>
```

```
int fsync(int fd)
```

On success, these system calls return zero. On error, -1 is returned, and errno is set appropriately.

■ An Example of fsync().

```
int fd = open("foo", O_CREAT | O_WRONLY | O_TRUNC);  
assert (fd > -1)  
int rc = write(fd, buffer, size);  
assert (rc == size);  
rc = fsync(fd);  
assert (rc == 0);
```

- ◆ In some cases, this code needs to fsync() the directory that contains the file foo.

Renaming Files

- ❑ `rename(char* old, char *new)`
 - ◆ Rename a file to different name.
 - ◆ It implemented as an **atomic call**.
 - e.g. Change from `foo` to `bar`:

```
prompt> mv foo bar           // mv uses the system call rename()
```

- How to update a file atomically:

```
int fd = open("foo.txt.tmp", O_WRONLY|O_CREAT|O_TRUNC, S_IRWXU);
char buffer[20] = "hello";
write(fd, buffer, sizeof(buffer)); // write out new version of file
fsync(fd);
close(fd);
rename("foo.txt.tmp", "foo.txt");
```


Getting Information About Files

▣ `stat()`, `fstat()` : Show the file metadata

◆ **Metadata** is information about each file.

- Size, Low-level name, Permission, ...

◆ `stat` structure is below:

```
struct stat {
    dev_t st_dev;           /* ID of device containing file */
    ino_t st_ino;           /* inode number */
    mode_t st_mode;         /* protection */
    nlink_t st_nlink;       /* number of hard links */
    uid_t st_uid;           /* user ID of owner */
    gid_t st_gid;           /* group ID of owner */
    dev_t st_rdev;          /* device ID (if special file) */
    off_t st_size;          /* total size, in bytes */
    blksize_t st_blksize;   /* blocksize for filesystem I/O */
    blkcnt_t st_blocks;     /* number of blocks allocated */
    time_t st_atime;        /* time of last access */
    time_t st_mtime;        /* time of last modification */
    time_t st_ctime;        /* time of last status change */
};
```

Getting Information About Files (Cont.)

- To see stat information, you can use the command line tool `stat`.

```
prompt> echo hello > file
prompt> stat file

File: `file'
Size: 6 Blocks: 8 IO Block: 4096 regular file
Device: 811h/2065d Inode: 67158084 Links: 1
Access: (0640/-rw-r-----) Uid: (30686/ root) Gid: (30686/ remzi)
Access: 2011-05-03 15:50:20.157594748 -0500
Modify: 2011-05-03 15:50:20.157594748 -0500
Change: 2011-05-03 15:50:20.157594748 -0500
```

- ◆ File system keeps this type of information in an `inode` structure.

Removing Files

- ❑ `rm` is Linux command to remove a file
 - ◆ `rm` call `unlink()` to remove a file.

```
prompt> strace rm foo
...
unlink("foo")          = 0          // return 0 upon success
...
prompt>
```

Why it calls `unlink()` not "remove or delete"?
We can get the answer later.

Making Directories

▣ `mkdir()`: Make a directory

```
prompt> strace mkdir foo
...
mkdir("foo", 0777)           = 0
prompt>
```

- ◆ When a directory is created, it is **empty**.
- ◆ Empty directory have two entries: `.` (itself), `..` (parent)

```
prompt> ls -a
./      ../
prompt> ls -al
total 8
drwxr-x---  2 remzi remzi    6 Apr 30 16:17 ./
drwxr-x--- 26 remzi remzi 4096 Apr 30 16:17 ../
```

Reading Directories

- A sample code to read directory entries (like `ls`).

```
#include <sys/stat.h>
#include <assert.h>
#include <stdio.h>
#include <dirent.h>

int main(int argc, char *argv[]) {
    DIR *dp = opendir(".");           // open current directory
    assert(dp != NULL);
    struct dirent *d;
    while ((d = readdir(dp)) != NULL) // read one directory entry
    {
        // print out the name and inode number of each file
        printf("%d %s\n", (int) d->d_ino, d->d_name);
    }
    closedir(dp);                    // close current directory
    return 0;
}
```

- ◆ The information available within struct dirent

```
struct dirent {
    char          d_name[256];        /* filename */
    ino_t         d_ino;              /* inode number */
    off_t         d_off;              /* offset to the next direct */
    unsigned short d_reclen;          /* length of this record */
    unsigned char  d_type;            /* type of file */
}
```

Deleting Directories

- `rmdir()` : Delete a directory.
 - ◆ Require that the directory be **empty**
 - i.e. Only has “.” and “..” entries.
 - ◆ If you call `rmdir()` to a non-empty directory, it will fail.

Hard Links

- ▣ `link(old pathname, new one)`
 - ◆ **Link** a new file name to an old one
 - ◆ Create another way to refer to *the same file*
 - ◆ The command-line link program : `ln`

```
prompt> echo hello > file
prompt> cat file
hello
prompt> ln file file2 // create a hard link, link file to file2
prompt> cat file2
hello
```

Hard Links (Cont.)

- The way `link` works:
 - ◆ **Create** another name in the directory.
 - ◆ **Refer** it to the same inode number of the original file.
 - The file is not copied in any way.
 - ◆ Then, we now just have two human names (`file` and `file2`) that both refer to the same file.

Hard Links (Cont.)

□ The result of `link()`

```
prompt> ls -li file file2
67158084 file /* inode value is 67158084 */
67158084 file2 /* inode value is 67158084 */
prompt>
```

- ◆ Two files have **same inode** number, but two file name (file, file2).
- ◆ There is **no difference** between file and file2.
 - Both just link to the underlying metadata about the file.

Hard Links (Cont.)

- Thus, to remove a file, we call `unlink()`.

```
prompt> rm file
removed 'file'
prompt> cat file2           // Still access the file
hello
```

- ◆ *reference count*

- Track how many different file names have been linked to this inode.
- When `unlink()` is called, the reference count decrements.
- If the reference count reaches zero, the files system free the inode and related data blocks.
→ truly “delete” the file

Hard Links (Cont.)

▣ The result of `unlink()`

- ◆ `stat()` shows the reference count of a file.

```
prompt> echo hello > file          /* create file*/
prompt> stat file
... Inode: 67158084 Links: 1 ...    /* Link count is 1 */
prompt> ln file file2              /* hard link file2 */
prompt> stat file
... Inode: 67158084 Links: 2 ...    /* Link count is 2 */
prompt> stat file2
... Inode: 67158084 Links: 2 ...    /* Link count is 2 */
prompt> ln file2 file3             /* hard link file3 */
prompt> stat file
... Inode: 67158084 Links: 3 ...    /* Link count is 3 */
prompt> rm file                    /* remove file */
prompt> stat file2
... Inode: 67158084 Links: 2 ...    /* Link count is 2 */
prompt> rm file2                   /* remove file2 */
prompt> stat file3
... Inode: 67158084 Links: 1 ...    /* Link count is 1 */
prompt> rm file3
```

Symbolic Links (Soft Link)

- ❑ Symbolic link is more **useful** than Hard link.
 - ◆ Hard Link cannot be created for a directory.
 - ◆ Hard Link cannot create to a file to other partition.
 - Because inode numbers are only unique within a file system.

- ❑ Create a symbolic link: `ln -s`

```
prompt> echo hello > file
prompt> ln -s file file2 /* option -s : create a symbolic link, */
prompt> cat file2
hello
```

Symbolic Links (Cont.)

- ❑ What is different between *Symbolic link* and *Hard Link*?
 - ◆ Symbolic links are a **third type** the file system knows about.

```
prompt> stat file
... regular file ...
prompt> stat file2
... symbolic link ...           // Actually a file itself of a different type
```

- ◆ The size of symbolic link (file2) is **4 bytes**.

```
prompt> ls -al
drwxr-x---  2 remzi remzi   29 May 3 19:10 ./
drwxr-x--- 27 remzi remzi 4096 May 3 15:14 ../           // directory
-rw-r----- 1 remzi remzi    6 May 3 19:10 file         // regular file
lrwxrwxrwx  1 remzi remzi    4 May 3 19:10 file2 -> file // symbolic link
```

- A symbolic link holds the pathname of the linked-to file as the data of the link file.

Symbolic Links (Cont.)

- If we link to a longer pathname, our link file would be bigger.

```
prompt> echo hello > alongerfilename
prompt> ln -s alongerfilename file3
prompt> ls -al alongerfilename file3
-rw-r----- 1 remzi remzi  6 May 3 19:17 alongerfilename
lrwxrwxrwx 1 remzi remzi 15 May 3 19:17 file3 -> alongerfilename
```

Symbolic Links (Cont.)

❑ Dangling reference

- ◆ When remove an original file, symbolic link points nothing.

```
prompt> echo hello > file
prompt> ln -s file file2
prompt> cat file2
hello
prompt> rm file           // remove the original file
prompt> cat file2
cat: file2: No such file or directory
```

Summary

- ▣ System calls related to file & directory
 - ◆ `open()`, `read()`, `write()`, `lseek()`
 - ◆ `dup()`, `pipe()`, `fsync()`, `rename()`, `stat()`
 - ◆ `mkdir()`, `rmdir()`, `opendir()`, `readdir()`
 - ◆ `unlink()`, hard/symbolic link
- ▣ Next: File System Implementation
 - ◆ [Chapter 40](#)