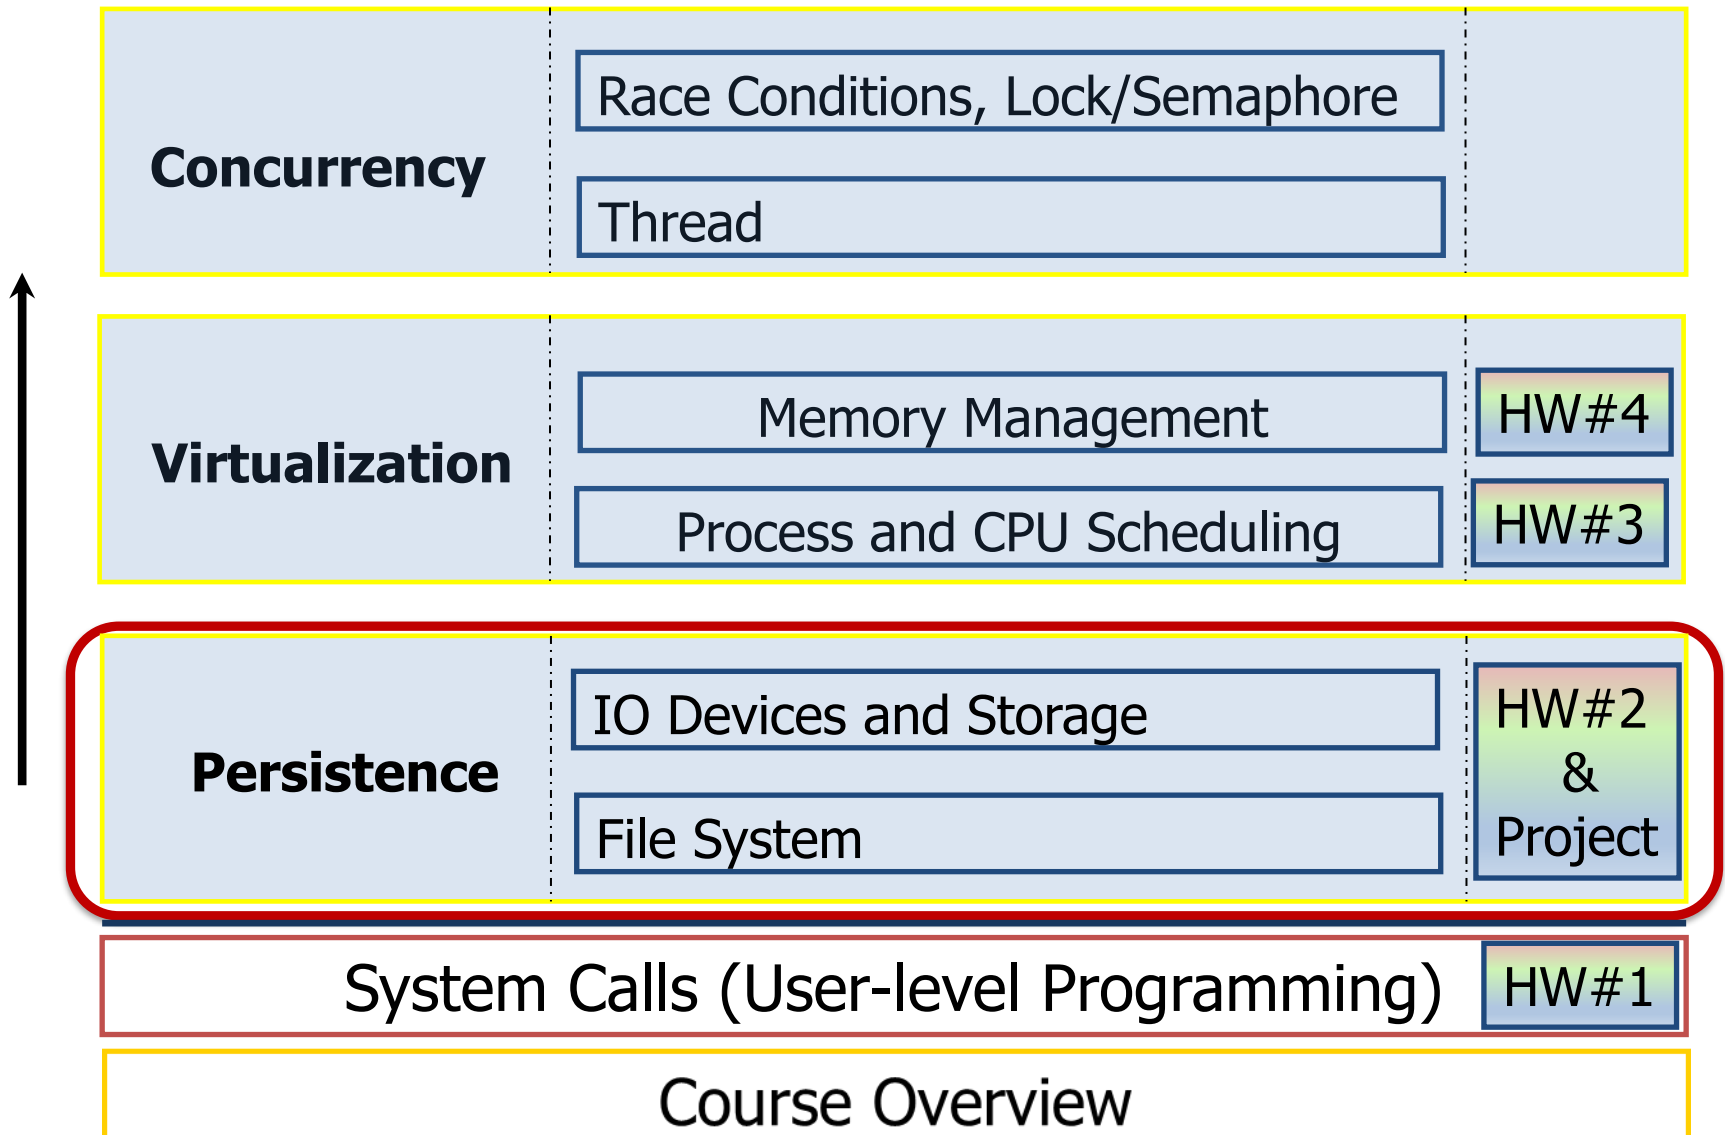



Lecture 6: IO Devices

The Course Organization (Bottom-up)



- ▣ I/O is **critical** for computer systems to **interact with IO systems**.
- ▣ Issue :
 - ◆ How should I/O be integrated into systems?
 - ◆ What are the general mechanisms?
 - ◆ How can we make them work efficiently?

Structure of input/output (I/O) device



Prototypical System Architecture

CPU is attached to the main memory of the system via some kind of memory bus.

Some devices are connected to the system via a general I/O bus.

□ Buses

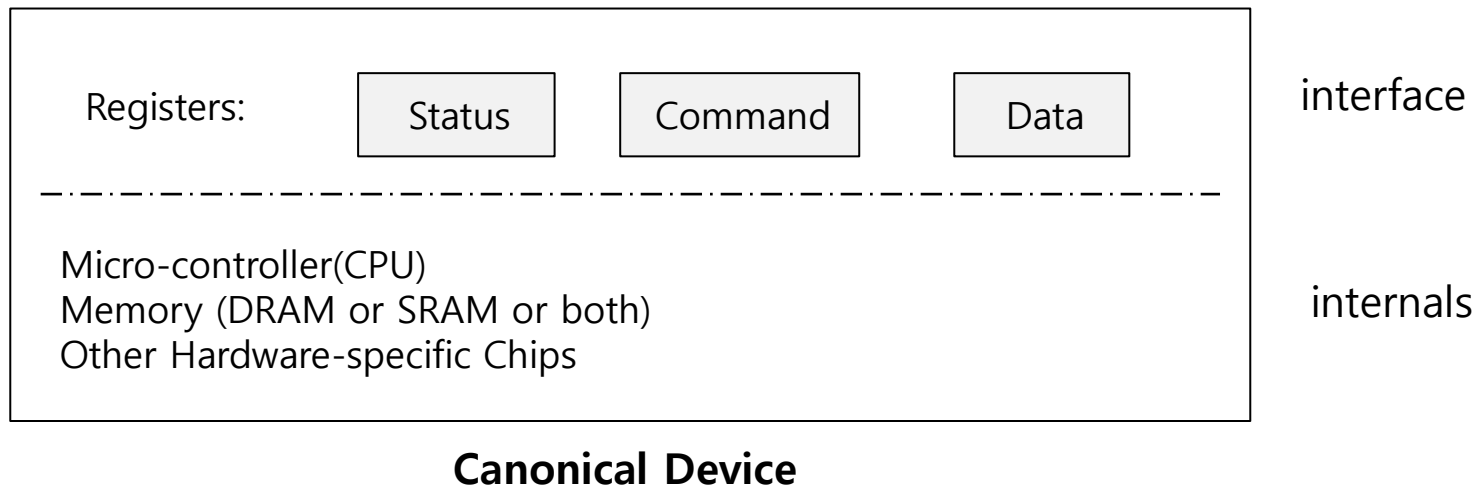
- ◆ Data paths that are provided to communicate information among CPU(s), RAM, and I/O devices.

□ I/O bus

- ◆ Data path that connects CPU to I/O devices.
- ◆ I/O device is connected to I/O bus by three hardware components: I/O ports, interfaces and device controllers.

Canonical Device

- Canonical Devices has two important components.
 - ◆ **Hardware interface** allows the system software to control its operation.
 - ◆ **Internals** which are implementation specific.



Hardware interface of Canonical Device

- ❑ **status register**

- ◆ See the current status of the device

- ❑ **command register**

- ◆ Tell the device to perform a certain task

- ❑ **data register**

- ◆ Pass data to the device, or get data from the device

By reading and writing above **three registers,
the operating system can **control device behavior**.**

Hardware interface of Canonical Device (Cont.)

▣ Typical interaction example

```
while ( STATUS == BUSY)
    ; //wait until device is not busy
write data to data register
write command to command register
    Doing so starts the device and executes the command
while ( STATUS == BUSY)
    ; //wait until device is done with your request
```


Polling

- ❑ Operating system waits until the device is ready by **repeatedly** reading the status register.
 - ◆ Positive aspect: Simple and it works.
 - ◆ **However, it wastes CPU time just waiting for the device.**
 - Switching to another ready process may better utilize the CPU.

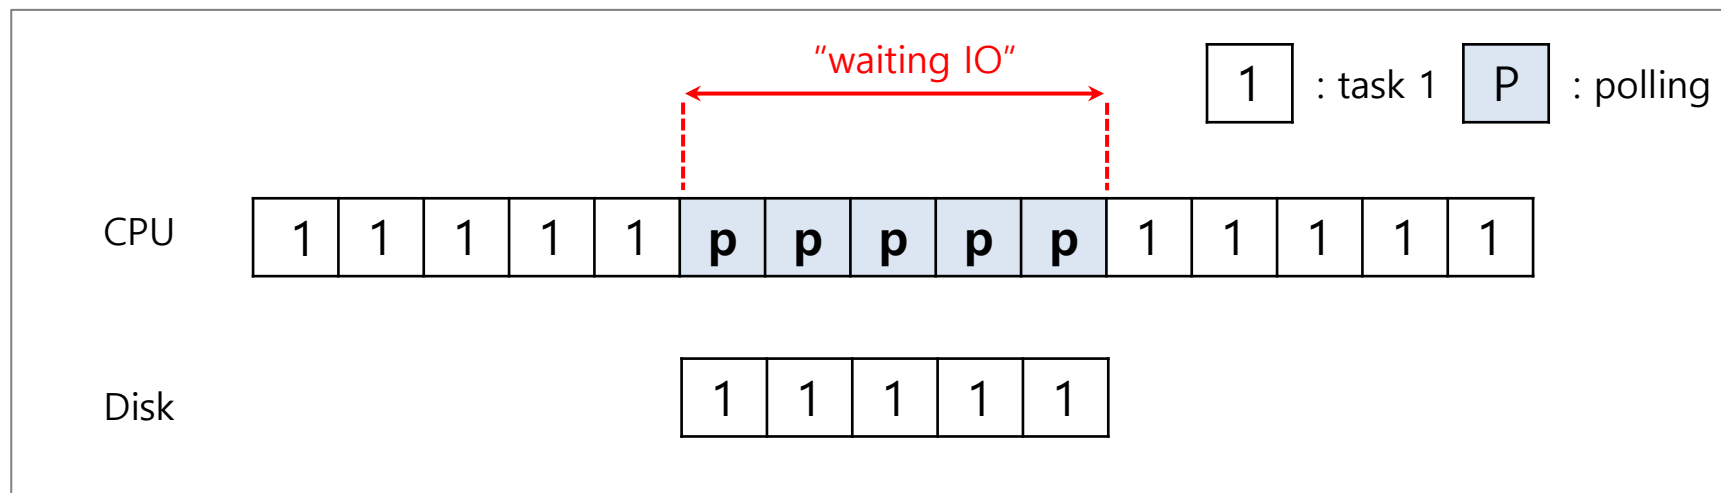


Diagram of CPU utilization by polling

Interrupt

- ❑ Put the I/O request process to sleep and context switch to another.
- ❑ When the device is finished, wake the process waiting for the I/O by **interrupt**.
 - ◆ Positive aspect is to allow **CPU and the disk are properly utilized**.




Diagram of CPU utilization by interrupt

Polling vs interrupts

- *However, “interrupts is not always the best solution”*
 - ◆ If device performs very quickly, interrupt will “slow down” the system.
 - ◆ Because **context switch is expensive** (switching to another process)

If a device is fast → **poll** is best.
If it is slow → **interrupts** is better.

CPU is once again over-burdened

- ❑ CPU wastes a lot of time to copy *a large chunk of data* from memory to the device.

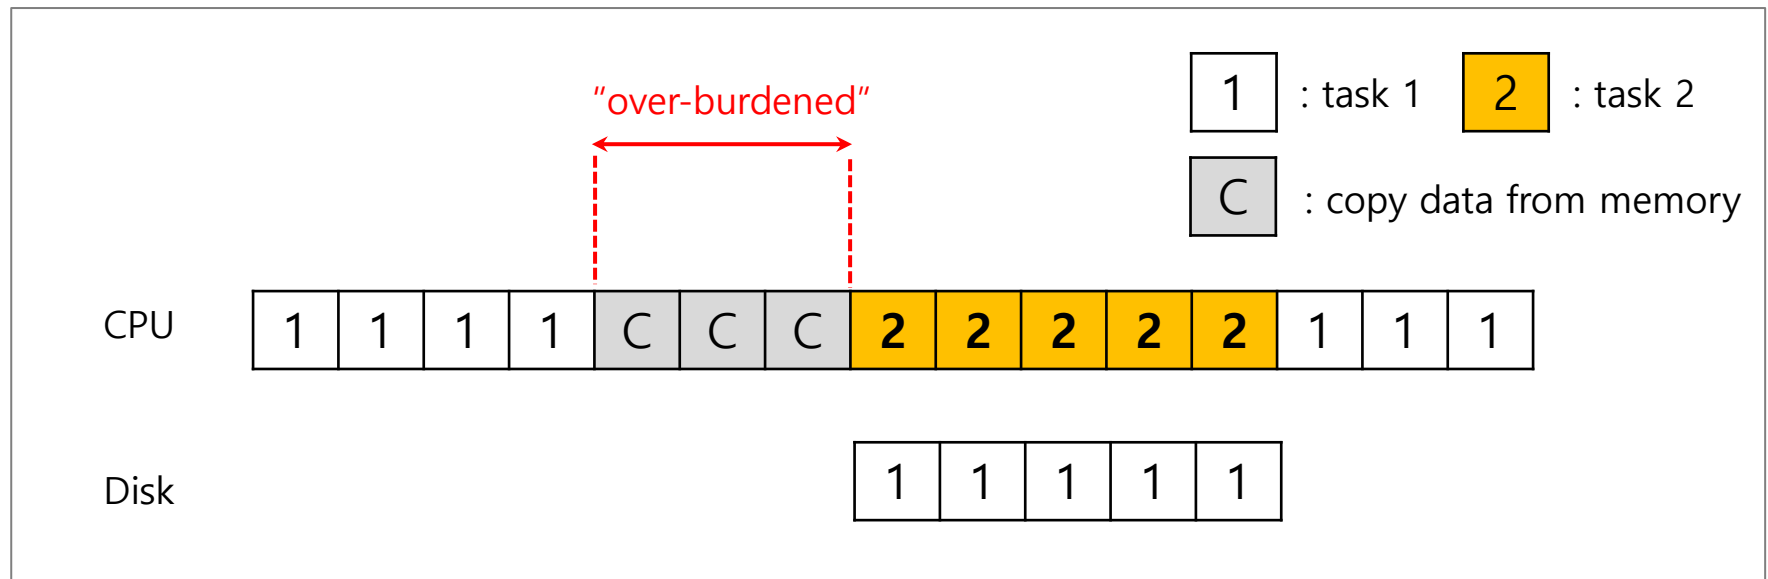


Diagram of CPU utilization

DMA (Direct Memory Access)

- ❑ **Copy data** in memory by knowing “where the data lives in memory, how much data to copy”
- ❑ When completed, DMA raises an interrupt, I/O begins on Disk.




Diagram of CPU utilization by DMA

Device interaction

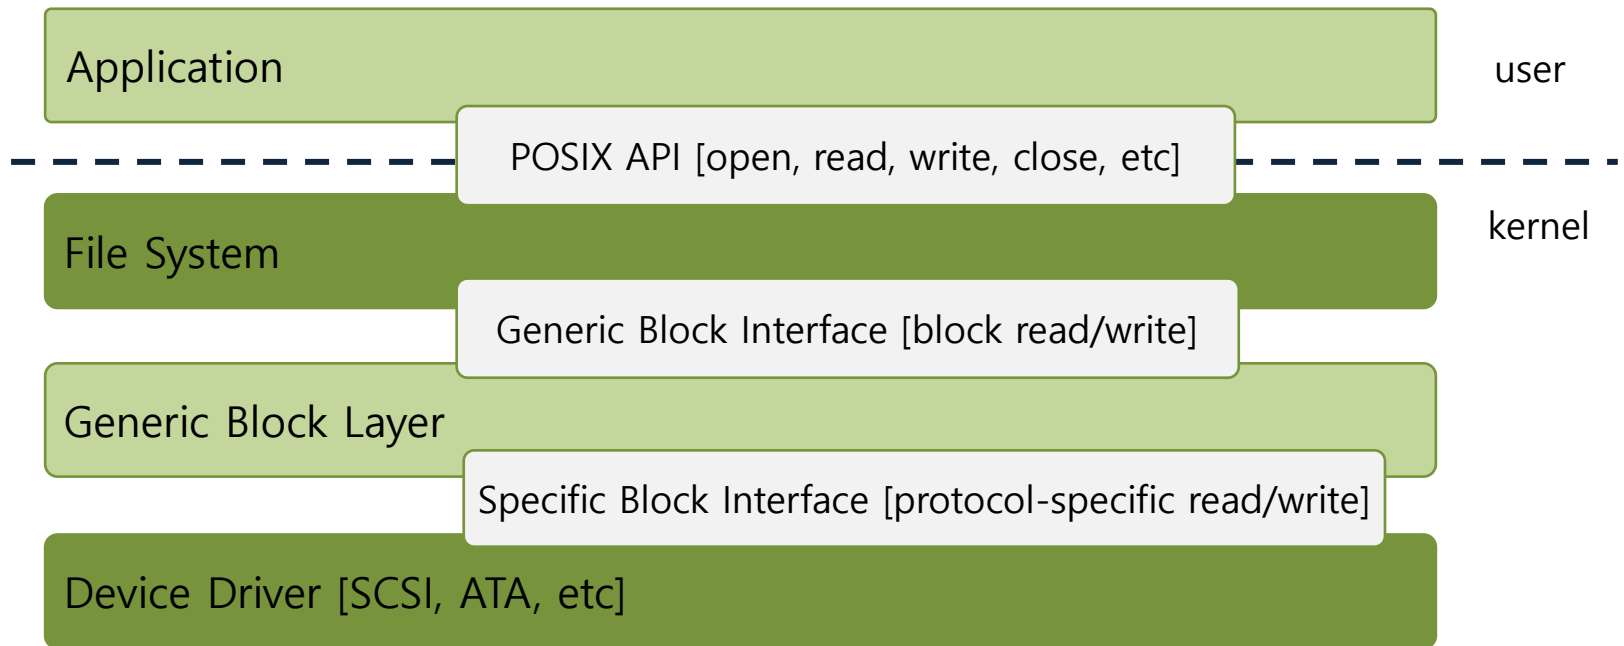
- ❑ How the OS communicates with the **device**?
- ❑ Solutions
 - ◆ **I/O instructions**: a way for the OS to send data to specific device registers.
 - Ex) `in` and `out` instructions on x86
 - ◆ **memory-mapped I/O**
 - Device registers available as if they were memory locations.
 - The OS `load` (to read) or `store` (to write) to the device instead of main memory.

Device interaction (Cont.)

- How does the OS interact with **different specific interfaces**?
 - ◆ Ex) We'd like to build a file system that worked on top of SCSI disks, IDE disks, USB keychain drivers, and so on.
- Solutions: **Abstraction**
 - ◆ Abstraction encapsulate **any specifics of device interaction**.

File system Abstraction

- File system **specifics** of which disk class it is using.
 - ◆ Ex) It issues **block read** and **write** request to the generic block layer.



The File System Stack

Problem of File system Abstraction

- ❑ If there is a device having many special capabilities, these capabilities **will go unused** in the generic interface layer.
- ❑ **Over 70% of OS** code is found in device drivers.
 - ◆ Any device drivers are needed because you might plug it to your system.
 - ◆ They are primary contributor to **kernel crashes**, making **more bugs**.

A Simple IDE Disk Driver

- Four types of register
 - ◆ Control, command block, status and error
 - ◆ `in` and `out` I/O instruction

A Simple IDE Disk Driver

- Control Register:

Address 0x3F6 = 0x80 (0000 1RE0): R=reset, E=0 means "enable interrupt"

- Command Block Registers:

Address 0x1F0 = Data Port

Address 0x1F1 = Error

Address 0x1F2 = Sector Count

Address 0x1F3 = LBA low byte

Address 0x1F4 = LBA mid byte

Address 0x1F5 = LBA hi byte

Address 0x1F6 = 1B1D TOP4LBA: B=LBA, D=drive

Address 0x1F7 = Command/status

A Simple IDE Disk Driver

- ▣ Status Register (Address 0x1F7):

| | | | | | | | |
|------|-------|-------|------|-----|------|-------|-------|
| 7 | 6 | 5 | 4 | 3 | 2 | 1 | 0 |
| BUSY | READY | FAULT | SEEK | DRQ | CORR | IDDEX | ERROR |

- ▣ Error Register (Address 0x1F1): (check when Status ERROR==1)

| | | | | | | | |
|-----|-----|----|------|-----|------|------|------|
| 7 | 6 | 5 | 4 | 3 | 2 | 1 | 0 |
| BBK | UNC | MC | IDNF | MCR | ABRT | T0NF | AMNF |

- ◆ BBK = Bad Block
- ◆ UNC = Uncorrectable data error
- ◆ MC = Media Changed
- ◆ IDNF = ID mark Not Found
- ◆ MCR = Media Change Requested
- ◆ ABRT = Command aborted
- ◆ T0NF = Track 0 Not Found
- ◆ AMNF = Address Mark Not Found

A Simple IDE Disk Driver

- ❑ **Wait for drive to be ready.** Read Status Register (0x1F7) until drive is not busy and READY.
- ❑ **Write parameters to command registers.** Write the sector count, logical block address (LBA) of the sectors to be accessed, and drive number (master=0x00 or slave=0x10, as IDE permits just two drives) to command registers (0x1F2-0x1F6).
- ❑ **Start the I/O** by issuing read/write to command register. Write READ—WRITE command to command register (0x1F7).
- ❑ **Data transfer (for writes):** Wait until drive status is READY and DRQ (drive request for data); write data to data port.
- ❑ **Handle interrupts.** In the simplest case, handle an interrupt for each sector transferred; more complex approaches allow batching and thus one final interrupt when the entire transfer is complete.
- ❑ **Error handling.** After each operation, read the status register. If the ERROR bit is on, read the error register for details.

A Simple IDE Disk Driver

```
static int ide_wait_ready() {  
    while (((int r = inb(0x1f7)) & IDE_BSY) ||  
           !(r & IDE_DRDY))  
        ; // loop until drive isn't busy  
}
```

A Simple IDE Disk Driver

```
static void ide_start_request(struct buf *b) {  
    ide_wait_ready();  
    outb(0x3f6, 0); // generate interrupt  
    outb(0x1f2, 1); // how many sectors?  
    outb(0x1f3, b->sector & 0xff); // LBA goes here ...  
    outb(0x1f4, (b->sector >> 8) & 0xff); // ... and here  
    outb(0x1f5, (b->sector >> 16) & 0xff); // ... and here!  
    outb(0x1f6, 0xe0 | ((b->dev&1)<<4) | ((b->sector>>24)&0x0f));  
    if(b->flags & B_DIRTY){  
        outb(0x1f7, IDE_CMD_WRITE); // this is a WRITE  
        outsl(0x1f0, b->data, 512/4); // transfer data too!  
    } else {  
        outb(0x1f7, IDE_CMD_READ); // this is a READ (no data)  
    }  
}
```

A Simple IDE Disk Driver

```
void ide_rw(struct buf *b) {  
    acquire(&ide_lock);  
  
    for (struct buf **pp = &ide_queue; *pp; pp=&(*pp)->qnext)  
        ; // walk queue  
  
    *pp = b; // add request to end  
  
    if (ide_queue == b) // if q is empty  
        ide_start_request(b); // send req to disk  
  
    while ((b->flags & (B_VALID|B_DIRTY)) != B_VALID)  
        sleep(b, &ide_lock); // wait for completion  
  
    release(&ide_lock);  
}
```


A Simple IDE Disk Driver

```
void ide_intr() {  
    struct buf *b;  
  
    acquire(&ide_lock);  
  
    if (!(b->flags & B_DIRTY) && ide_wait_ready(1) >= 0)  
        insl(0x1f0, b->data, 512/4); // if READ: get data  
  
    b->flags |= B_VALID;  
  
    b->flags &= B_DIRTY;  
  
    wakeup(b); // wake waiting process  
  
    if ((ide_queue = b->qnext) != 0) // start next request  
        ide_start_request(ide_queue); // (if one exists)  
  
    release(&ide_lock);  
  
}
```

Summary

- IO Devices
 - ◆ IO system architecture
 - ◆ General mechanisms
 - ◆ Device interactions and software development
- Next: Virtualizing CPU – Process
 - ◆ [Chapter 4](#)
 - ◆ [Chapter 6](#)