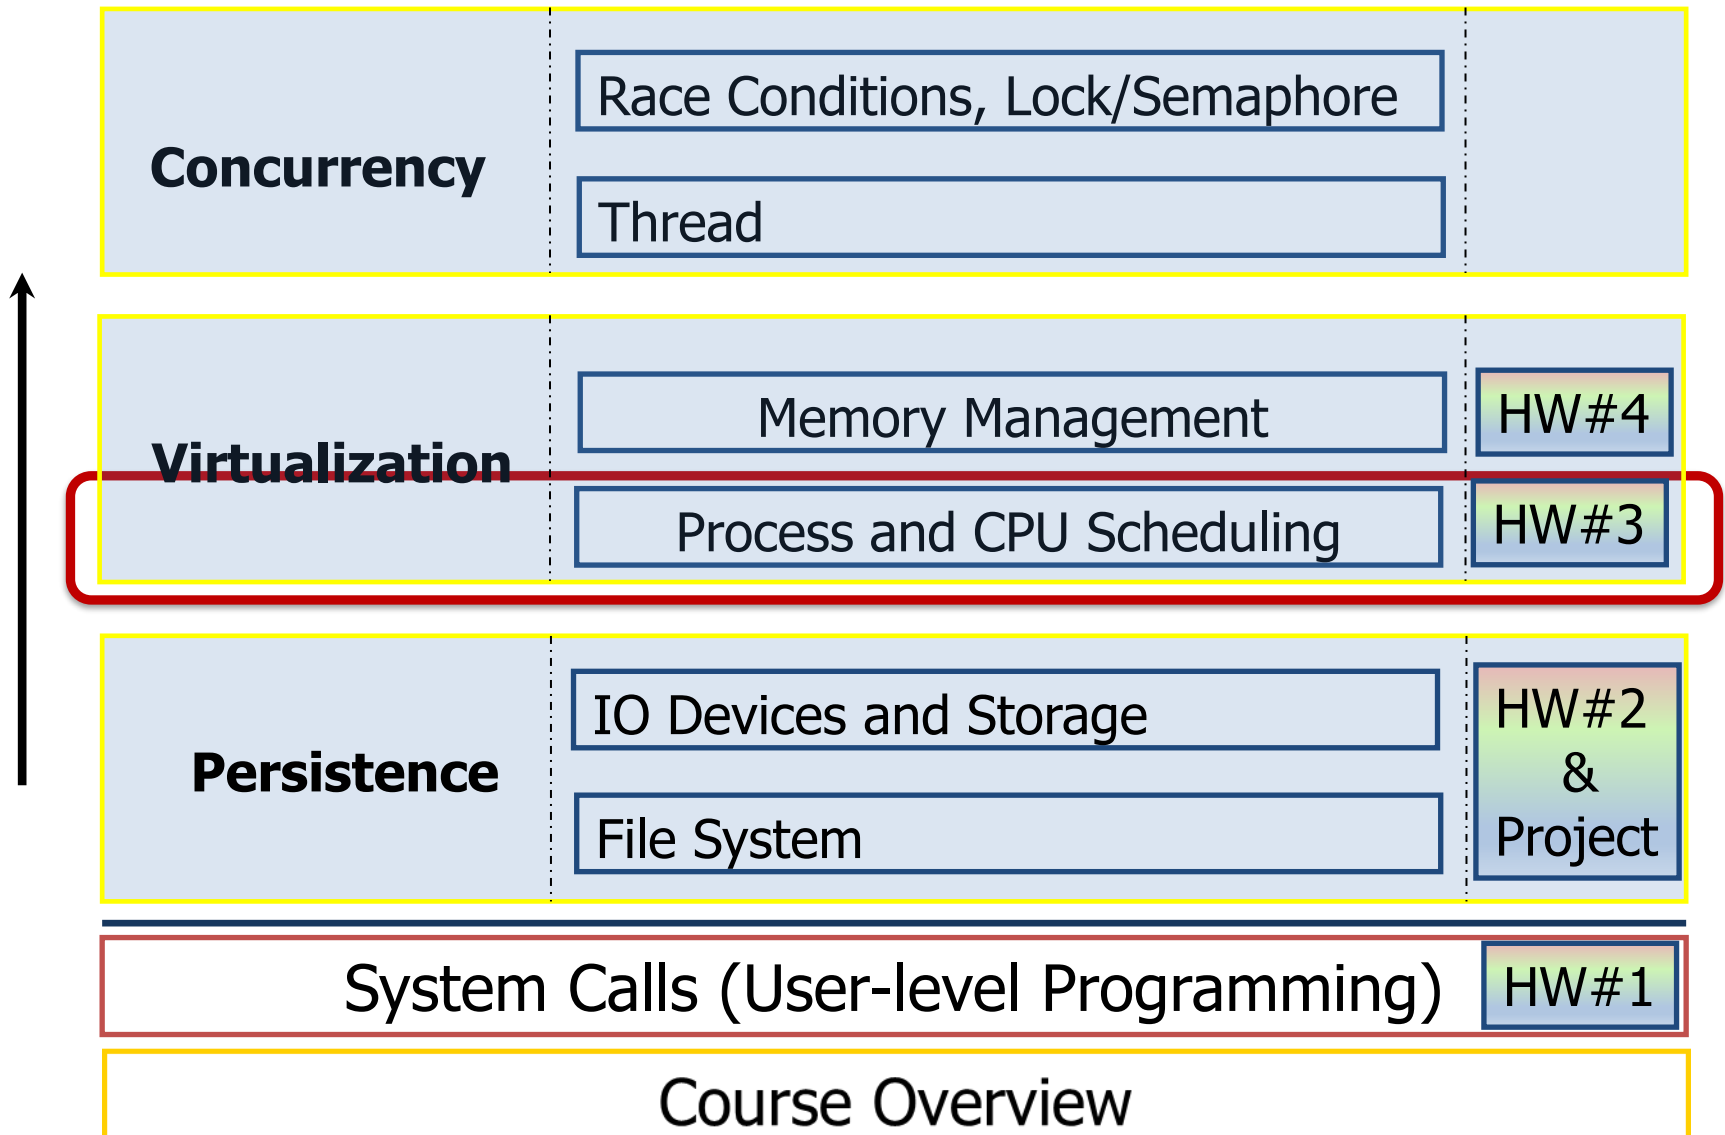


# **Lecture 7: Virtualizing CPU - Process**

---

# The Course Organization (Bottom-up)




# OS – Resource management via virtualization

OS provides services via **System Call** (typically a few hundred) to run **process**, access memory/devices/files, etc.

The OS **manages resources** such as *CPU*, *memory* and *disk* via **virtualization**.

- many programs to run (processes) → Sharing the CPU
- many processes to *concurrently* access their own instructions and data → Sharing memory
- many processes to access devices → Sharing disks



**The Design Of The Unix Operating System (Maurice Bach, 1986)**

# **Part I. The Abstraction: The Process**

---

# Virtualizing CPUs

- ❑ The OS can provide the illusion that many virtual CPUs exist.
- ❑ **Time sharing:** Running one process, then stopping it and running another
  - ◆ The potential cost is **performance**.

**A process is a **running program**.**

- ▣ Comprising of a process:
  - ◆ Memory (address space)
    - Instructions
    - Data section
  - ◆ Registers
    - Program counter
    - Stack pointer

# Process API

- ▣ These APIs are available on any modern OS.
  - ◆ **Create**
    - Create a new process to run a program
  - ◆ **Destroy**
    - Halt a runaway process
  - ◆ **Wait**
    - Wait for a process to stop running
  - ◆ **Miscellaneous Control**
    - Some kind of method to suspend a process and then resume it
  - ◆ **Status**
    - Get some status info about a process

# Process Creation

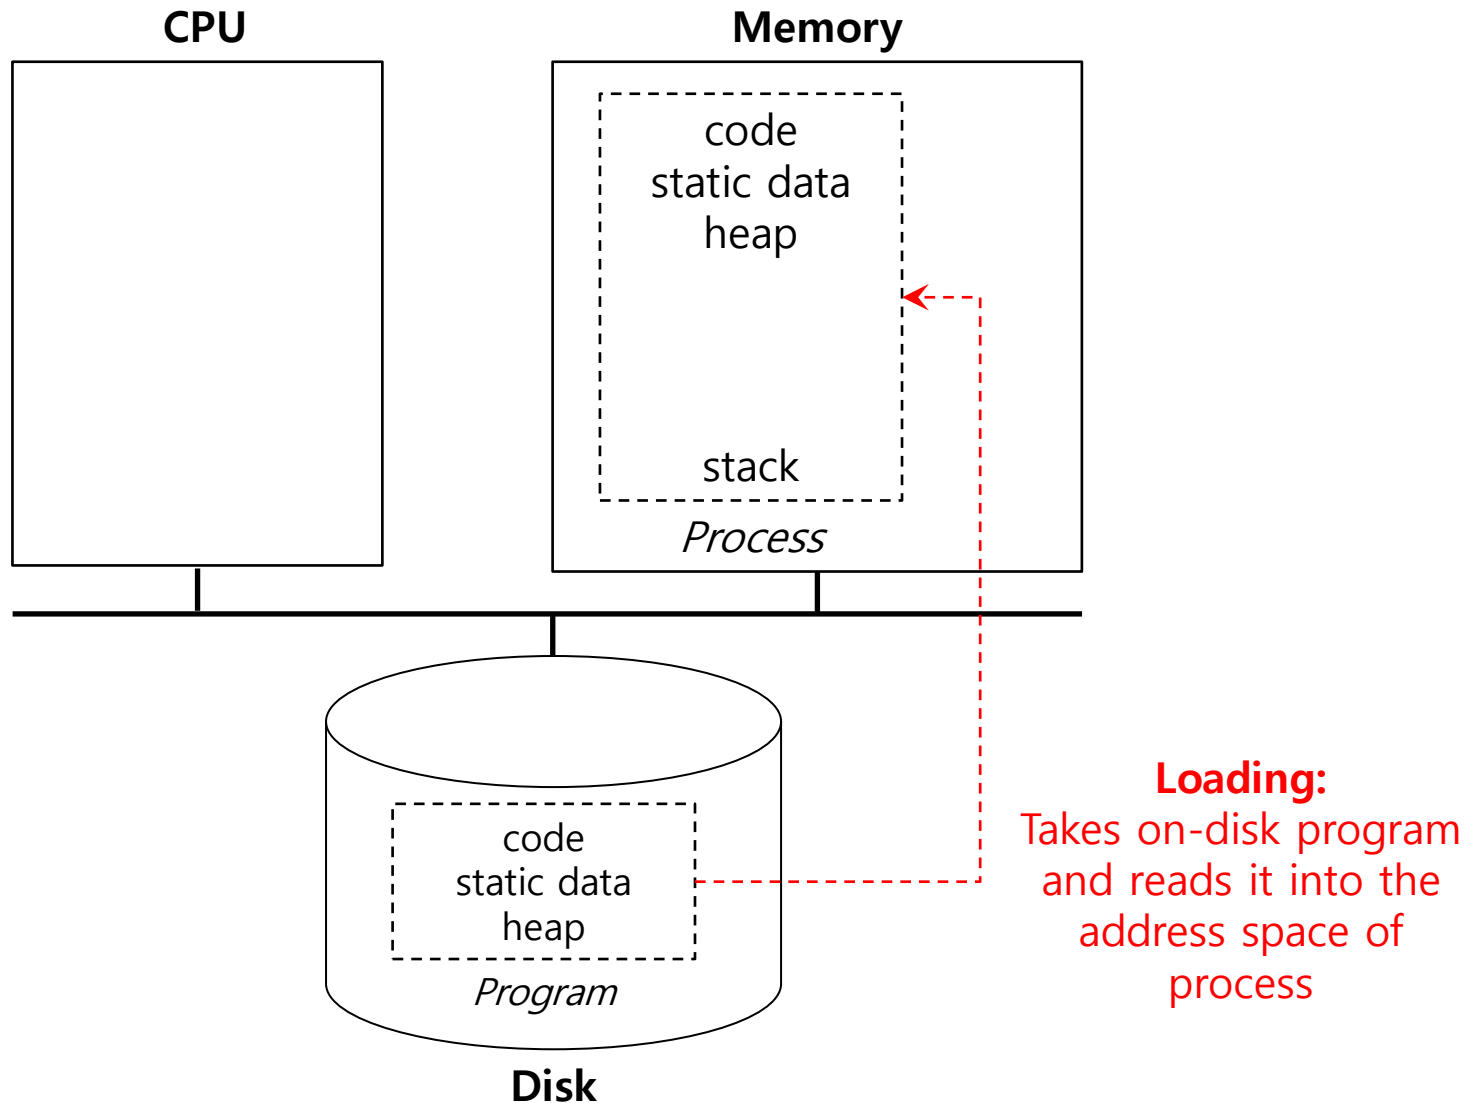
1. **Load** a program code into memory, into the address space of the process.
  - ◆ Programs initially reside on disk in *executable format*.
  - ◆ OS perform the loading process **lazily**.
    - Loading pieces of code or data only as they are needed during program execution.
2. The program's run-time **stack** is allocated.
  - ◆ Use the stack for *local variables*, *function parameters*, and *return address*.
  - ◆ Initialize the stack with arguments → `argc` and the `argv` array of `main()` function



## Process Creation (Cont.)

3. The program's **heap** is created.
  - ◆ Used for explicitly requested dynamically allocated data.
  - ◆ Program request such space by calling `malloc()` and free it by calling `free()`.
4. The OS do some other initialization tasks.
  - ◆ Input/output (I/O) setup
    - Each process by default has three open file descriptors.
    - Standard input, output and error
5. **Start the program** running at the entry point, namely `main()`.
  - ◆ The OS *transfers control* of the CPU to the newly-created process.

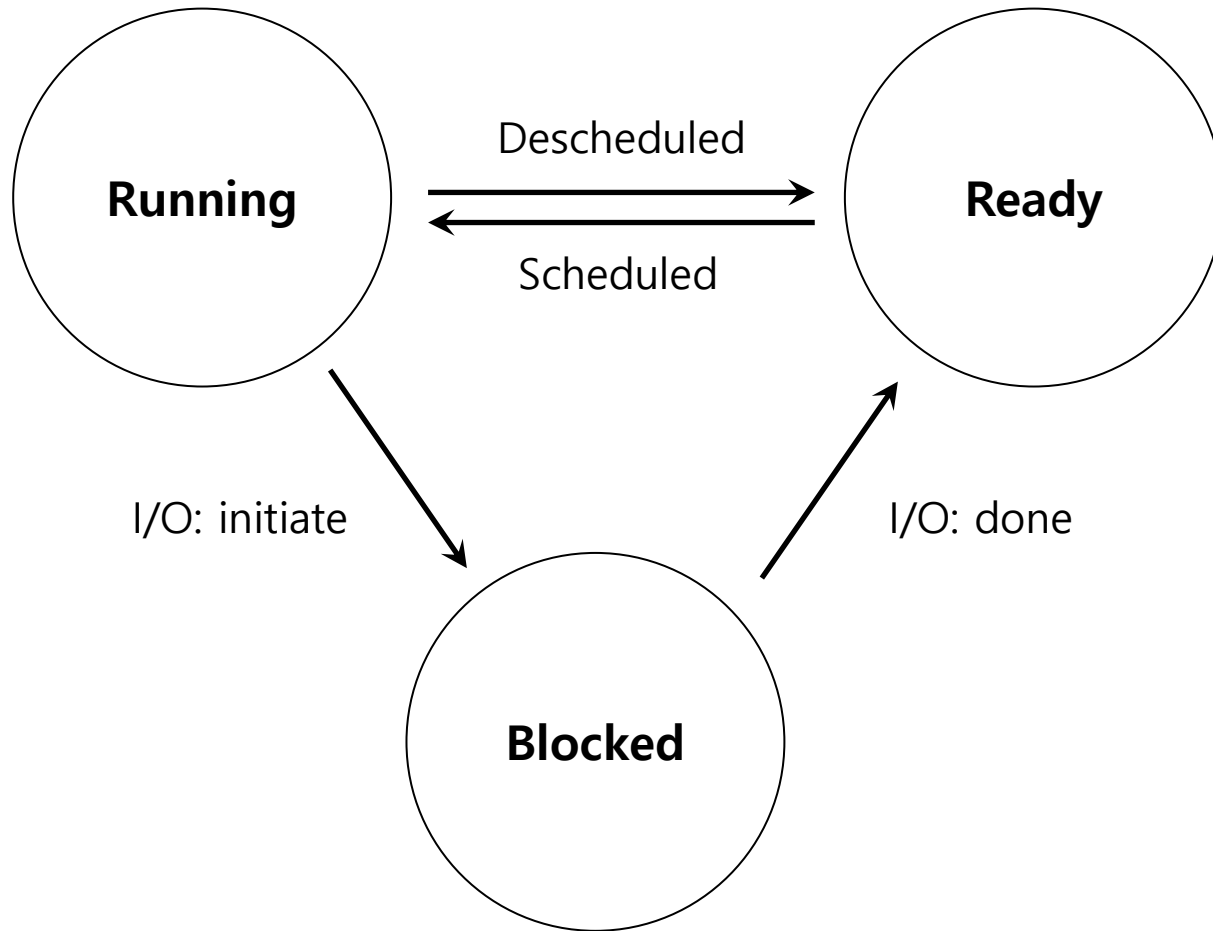
# Loading: From Program To Process



# Process States

- A process can be one of three states.
  - ◆ **Running**
    - A process is running on a processor.
  - ◆ **Ready**
    - A process is ready to run but for some reason the OS has chosen not to run it at this given moment.
  - ◆ **Blocked**
    - A process has performed some kind of operation.
    - When a process initiates an I/O request to a disk, it becomes blocked and thus some other process can use the processor.

# Process State Transition



# Data structures

- ❑ The OS has **some key data structures** that track various relevant pieces of information.
  - ◆ **Process list**
    - Ready processes
    - Blocked processes
    - Current running process
  - ◆ **Register context**
- ❑ PCB (Process Control Block)
  - ◆ An in-memory data structure that contains information **about each process**.

## Example: The xv6 kernel Proc Structure (proc.h)

```
// proc.h in xv6
// the registers xv6 will save and restore
// to stop and subsequently restart a process
struct context {
    uint edi;
    uint esi;
    uint ebx;
    uint ebp;
    uint eip;
};

// the different states a process can be in
enum proc_state { UNUSED, EMBRYO, SLEEPING,
                  RUNNABLE, RUNNING, ZOMBIE };
```

<https://github.com/mit-pdos/xv6-public>

**xv6** is a re-implementation of **Dennis Ritchie's** and **Ken Thompson's Unix Version 6** (v6). xv6 loosely follows the structure and style of v6, but is implemented for a modern x86-based multiprocessor using ANSI C.

## Example: The xv6 kernel Proc Structure (Cont.)

```
// the information xv6 tracks about each process
// including its register context and state
struct proc {
    char *mem;           // Start of process memory
    uint sz;             // Size of process memory
    char *kstack;        // Bottom of kernel stack
                        // for this process
    enum proc_state state; // Process state
    int pid;             // Process ID
    struct proc *parent; // Parent process
    void *chan;          // If non-zero, sleeping on chan
    int killed;          // If non-zero, have been killed
    struct file *ofile[NOFILE]; // Open files
    struct inode *cwd;    // Current directory
    struct context context; // Switch here to run process
    struct trapframe *tf; // Trap frame for the
                        // current interrupt
};
```

## **Part II: Limited Direct Execution Mechanism**

---



# How to efficiently virtualize the CPU with control?

- ▣ The OS needs to share the physical CPU by **time sharing**.
- ▣ Issue
  - ◆ **Performance:** How can we implement virtualization without adding excessive overhead to the system?
  - ◆ **Control:** How can we run processes efficiently while retaining control over the CPU?

# Direct Execution

- Just run the program directly on the CPU.

OS	Program
<ol style="list-style-type: none"><li>1. Create entry for process list</li><li>2. Allocate memory for program</li><li>3. Load program into memory</li><li>4. Set up stack with <code>argc / argv</code></li><li>5. Clear registers</li><li>6. Execute call <code>main()</code></li></ol> <ol style="list-style-type: none"><li>9. Free memory of process</li><li>10. Remove from process list</li></ol>	<ol style="list-style-type: none"><li>7. Run <code>main()</code></li><li>8. Execute <code>return from main()</code></li></ol>

Without *limits* on running programs,  
the OS wouldn't be in control of anything and  
thus would be "just a library"

# Problem 1: Restricted Operation

- ❑ What if a process wishes to perform some kind of restricted operation such as ...
  - ◆ Issuing an I/O request to a disk
  - ◆ Gaining access to more system resources such as CPU or memory
  
- ❑ **Solution:** Using protected control transfer
  - ◆ **User mode:** Applications do not have full access to hardware resources.
  - ◆ **Kernel mode:** The OS has access to the full resources of the machine

# System Call

- ▣ Allow the kernel to **carefully expose** certain key pieces of functionality to user program, such as ...
  - ◆ Accessing the file system
  - ◆ Creating and destroying processes
  - ◆ Communicating with other processes
  - ◆ Allocating more memory

# System Call (Cont.)

## ❑ **Trap** instruction

- ◆ Jump into the kernel
- ◆ Raise the privilege level to kernel mode

## ❑ **Return-from-trap** instruction

- ◆ Return into the calling user program
- ◆ Reduce the privilege level back to user mode

# Example: Use “strace” to trace system calls and signals

```
Terminal - csci3150@csci3150-VirtualBox: /
File Edit View Terminal Tabs Help

csci3150@csci3150-VirtualBox:/$
csci3150@csci3150-VirtualBox:/$
csci3150@csci3150-VirtualBox:/$ strace ls
execve("/bin/ls", ["ls"], 0xbfbfd53f0 /* 48 vars */) = 0
brk(NULL) = 0x1806000
access("/etc/ld.so.nohwcap", F_OK) = -1 ENOENT (No such file or directory)
mmap2(NULL, 8192, PROT_READ|PROT_WRITE, MAP_PRIVATE|MAP_ANONYMOUS, -1, 0) = 0xb7ef6000
access("/etc/ld.so.preload", R_OK) = -1 ENOENT (No such file or directory)
openat(AT_FDCWD, "/etc/ld.so.cache", O_RDONLY|O_CLOEXEC) = 3
fstat64(3, {st_mode=S_IFREG|0644, st_size=79757, ...}) = 0
mmap2(NULL, 79757, PROT_READ, MAP_PRIVATE, 3, 0) = 0xb7ee2000
close(3) = 0
access("/etc/ld.so.nohwcap", F_OK) = -1 ENOENT (No such file or directory)
openat(AT_FDCWD, "/lib/i386-linux-gnu/libselinux.so.1", O_RDONLY|O_CLOEXEC) = 3
read(3, "\177ELF\1\1\1\0\0\0\0\0\0\0\0\3\0\3\0\1\0\0\0\0L\0\0004\0\0\0"... , 512) = 512
fstat64(3, {st_mode=S_IFREG|0644, st_size=169960, ...}) = 0
mmap2(NULL, 179612, PROT_READ|PROT_EXEC, MAP_PRIVATE|MAP_DENYWRITE, 3, 0) = 0xb7eb6000
mmap2(0xb7edf000, 8192, PROT_READ|PROT_WRITE, MAP_PRIVATE|MAP_FIXED|MAP_DENYWRITE, 3, 0x28000) = 0xb7edf000
mmap2(0xb7ee1000, 3484, PROT_READ|PROT_WRITE, MAP_PRIVATE|MAP_FIXED|MAP_ANONYMOUS, -1, 0) = 0xb7ee1000
close(3) = 0
access("/etc/ld.so.nohwcap", F_OK) = -1 ENOENT (No such file or directory)
openat(AT_FDCWD, "/lib/i386-linux-gnu/libc.so.6", O_RDONLY|O_CLOEXEC) = 3
read(3, "\177ELF\1\1\1\3\0\0\0\0\0\0\0\3\0\3\0\1\0\0\0\020\220\1\0004\0\0\0"... , 512) = 512
fstat64(3, {st_mode=S_IFREG|0755, st_size=1942840, ...}) = 0
mmap2(NULL, 1948188, PROT_READ|PROT_EXEC, MAP_PRIVATE|MAP_DENYWRITE, 3, 0) = 0xb7cda000
mprotect(0xb7eaf000, 4096, PROT_NONE) = 0
mmap2(0xb7eb0000, 12288, PROT_READ|PROT_WRITE, MAP_PRIVATE|MAP_FIXED|MAP_DENYWRITE, 3, 0x1d5000) = 0xb7eb0000
mmap2(0xb7eb3000, 10780, PROT_READ|PROT_WRITE, MAP_PRIVATE|MAP_FIXED|MAP_ANONYMOUS, -1, 0) = 0xb7eb3000
```

## Example: Use command “ausyscall” to see the system call list

```
Terminal - csci3150@csci3150-VirtualBox: /
File Edit View Terminal Tabs Help
csci3150@csci3150-VirtualBox:/$ ausyscall --dump
Using i386 syscall table:
0      restart_syscall
1      exit
2      fork
3      read
4      write
5      open
6      close
7      waitpid
8      creat
9      link
10     unlink
11     execve
12     chdir
13     time
14     mknod
15     chmod
16     lchown
17     break
18     oldstat
19     lseek
20     getpid
21     mount
22     umount
23     setuid
24     getuid
25     stime
```

# Limited Direction Execution Protocol

OS @ boot  
(kernel mode)

Hardware

initialize trap table

remember address of ...  
syscall handler

OS @ run  
(kernel mode)

Hardware

Program  
(user mode)

Create entry for process list  
Allocate memory for program  
Load program into memory  
Setup user stack with argv  
Fill kernel stack with reg/PC  
**return-from -trap**

restore regs from kernel stack  
move to user mode  
jump to main

Run main()  
...  
Call system  
**trap** into OS



# Limited Direction Execution Protocol (Cont.)

OS @ run  
(kernel mode)

Hardware

Program  
(user mode)

---

*(Cont.)*

save regs to kernel stack  
move to kernel mode  
jump to trap handler

Handle trap  
Do work of syscall  
**return-from-trap**

restore regs from kernel stack  
move to user mode  
jump to PC after trap

Free memory of process  
Remove from process list

...  
return from main  
trap (via `exit()`)

## Problem 2: Switching Between Processes

- How can the OS **regain control** of the CPU so that it can switch between *processes*?
  - ◆ A cooperative Approach: **Wait for system calls**
  - ◆ A Non-Cooperative Approach: **The OS takes control**

# A cooperative Approach: Wait for system calls

- ❑ Processes **periodically give up the CPU** by making **system calls** such as `yield`.
  - ◆ The OS decides to run some other task.
  - ◆ Application also transfer control to the OS when they do something illegal.
    - Divide by zero
    - Try to access memory that it shouldn't be able to access
  - ◆ e.g Early versions of the Macintosh OS, The old Xerox Alto system

**A process gets stuck in an infinite loop.  
→ Reboot the machine**

# A Non-Cooperative Approach: OS Takes Control

## □ A timer interrupt

- ◆ During the boot sequence, the OS start the timer.
- ◆ The timer raises an interrupt every so many milliseconds.
- ◆ When the interrupt is raised :
  - The currently running process is halted.
  - Save enough of the state of the program
  - A pre-configured interrupt handler in the OS runs.

**A timer interrupt gives OS the ability to run again on a CPU.**

# Saving and Restoring Context

- **Scheduler** makes a decision:
  - ◆ Whether to continue running the **current process**, or switch to a **different one**.
  - ◆ If the decision is made to switch, the OS executes context switch.

# Context Switch

- A low-level piece of assembly code
  - ◆ **Save the values of necessary registers** for the current process onto its kernel stack
    - General purpose registers
    - PC
    - kernel stack pointer
  - ◆ **Restore the register values** for the soon-to-be-executing process from its kernel stack
  - ◆ **Switch to the kernel stack** for the soon-to-be-executing process

# Limited Direction Execution Protocol (Timer interrupt)

OS @ boot  
(kernel mode)

Hardware

**initialize trap table**

remember address of ...  
syscall handler  
timer handler

**start interrupt timer**

start timer  
interrupt CPU in X ms

OS @ run  
(kernel mode)

Hardware

Program  
(user mode)

Process A

...

**timer interrupt**

save regs(A) to k-stack(A)  
move to kernel mode  
jump to trap handler

# Limited Direction Execution Protocol (Timer interrupt)

OS @ run  
(kernel mode)

Hardware

Program  
(user mode)

*(Cont.)*

Handle the trap

Call switch() routine

    save regs(A) to proc-struct(A)

    restore regs(B) from proc-struct(B)

    switch to k-stack(B)

**return-from-trap (into B)**

restore regs(B) from k-stack(B)

move to user mode

jump to B's PC

Process B

...



## The xv6 Context Switch Code (swtch.S)

[illegible]

# Concurrency Problems?

- ❑ What happens if, during interrupt or trap handling, another interrupt occurs?
- ❑ OS handles these situations:
  - ◆ **Disable interrupts** during interrupt processing
  - ◆ Use a number of sophisticated **locking** schemes to protect concurrent access to internal data structures.

# Summary

## ❑ Virtualize CPU

- ◆ The abstraction of process – Process in OS kernel
  - Process creation process, Process state, Process data structure in OS kernel
- ◆ Limited Direct Execution
  - User/kernel mode, System call (the interface between user/kernel), System call working process, and Process switch

## ❑ Next: CPU Scheduling

- ◆ [Chapter 7 \(Scheduling\)](#), [Chapter 8 \(Multi-level Feedback Queue\)](#), [Chapter 9 \(Proportional Share\)](#), [Chapter 10 \(Multi-CPU Scheduling\)](#)