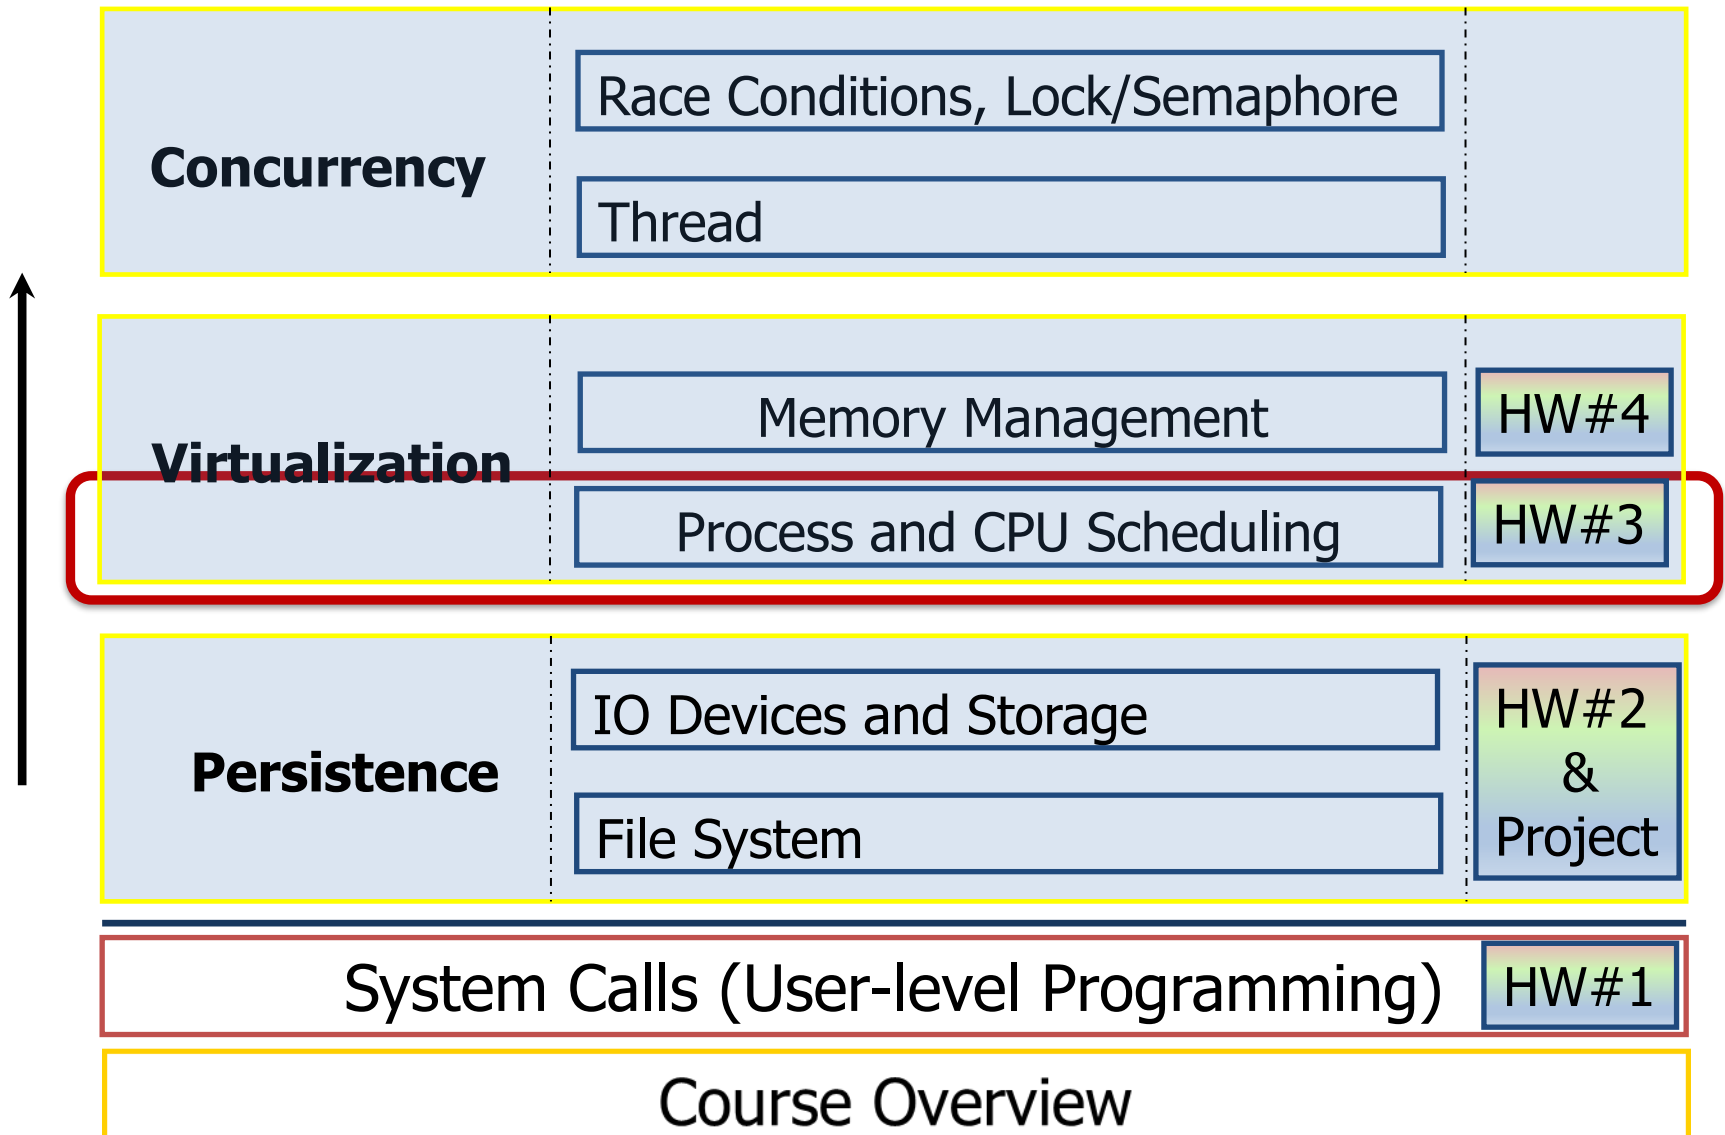


Lecture 8: Virtualizing CPU - Scheduling

The Course Organization (Bottom-up)



Part I. Scheduling: Introduction

Scheduling: Introduction

- Workload assumptions:

1. Each job runs for the **same amount of time**.
2. All jobs **arrive** at the same time.
3. All jobs only use the **CPU** (i.e., they perform no I/O).
4. The **run-time** of each job is known.

Scheduling Metrics

- Performance metric: **Turnaround time**

- ◆ The time at which **the job completes** minus the time at which **the job arrived** in the system.

$$T_{turnaround} = T_{completion} - T_{arrival}$$

- Another metric is **fairness**.

- ◆ Performance and fairness are often at odds in scheduling.


First In, First Out (FIFO)

- First Come, First Served (FCFS)

- ◆ Very simple and easy to implement

- Example:


- ◆ A arrived just before B which arrived just before C.
- ◆ Each job runs for 10 seconds.



$$\text{Average turnaround time} = \frac{10 + 20 + 30}{3} = 20 \text{ sec}$$

Why FIFO is not that great? – Convoy effect


- Let's relax assumption 1: Each job **no longer** runs for the same amount of time.
- Example:
 - ◆ A arrived just before B which arrived just before C.
 - ◆ A runs for 100 seconds, B and C run for 10 each.



$$\text{Average turnaround time} = \frac{100 + 110 + 120}{3} = \mathbf{110 \text{ sec}}$$

Shortest Job First (SJF)


- Run the shortest job first, then the next shortest, and so on
 - ◆ Non-preemptive scheduler
- Example:
 - ◆ A arrived just before B which arrived just before C.
 - ◆ A runs for 100 seconds, B and C run for 10 each.



$$\text{Average turnaround time} = \frac{10 + 20 + 120}{3} = 50 \text{ sec}$$

SJF with Late Arrivals from B and C

- Let's relax assumption 2: Jobs can arrive at any time.
- Example:
 - ◆ A arrives at $t=0$ and needs to run for 100 seconds.
 - ◆ B and C arrive at $t=10$ and each need to run for 10 seconds



$$\text{Average turnaround time} = \frac{100 + (110 - 10) + (120 - 10)}{3} = 103.33 \text{ sec}$$


Shortest Time-to-Completion First (STCF)

- Add **preemption** to SJF
 - ◆ Also known as Preemptive Shortest Job First (PSJF)
- A new job enters the system:
 - ◆ Determine of the remaining jobs and new job
 - ◆ Schedule the job which has the least time left

Shortest Time-to-Completion First (STCF)

□ Example:

- ♦ A arrives at $t=0$ and needs to run for 100 seconds.
- ♦ B and C arrive at $t=10$ and each need to run for 10 seconds



$$\text{Average turnaround time} = \frac{(120 - 0) + (20 - 10) + (30 - 10)}{3} = 50 \text{ sec}$$

New scheduling metric: Response time

- The time from **when the job arrives** to the **first time it is scheduled**.

$$T_{response} = T_{firstrun} - T_{arrival}$$

- ♦ STCF and related disciplines are not particularly good for response time.

**How can we build a scheduler that is
sensitive to response time?**

Round Robin (RR) Scheduling


□ Time slicing Scheduling

- ◆ Run a job for a **time slice** and then switch to the next job in the **run queue** until the jobs are finished.
 - Time slice is sometimes called a scheduling quantum.
- ◆ It repeatedly does so until the jobs are finished.
- ◆ The length of a time slice must be *a multiple of* the timer-interrupt period.

**RR is fair, but performs poorly on metrics
such as turnaround time**


RR Scheduling Example

- A, B and C arrive at the same time.
- They each wish to run for 5 seconds.



SJF (Bad for Response Time)

$$T_{average\ response} = \frac{0 + 5 + 10}{3} = 5sec$$



RR with a time-slice of 1sec (Good for Response Time)

$$T_{average\ response} = \frac{0 + 1 + 2}{3} = 1sec$$

The length of the time slice is critical.


- ❑ The shorter time slice
 - ◆ Better response time
 - ◆ The cost of context switching will dominate overall performance.
- ❑ The longer time slice
 - ◆ Amortize the cost of switching
 - ◆ Worse response time

**Deciding on the length of the time slice presents
a **trade-off** to a system designer**


Incorporating I/O

- ▣ Let's relax assumption 3: All programs perform I/O
- ▣ Example:
 - ◆ A and B need 50ms of CPU time each.
 - ◆ A runs for 10ms and then issues an I/O request
 - I/Os each take 10ms
 - ◆ B simply uses the CPU for 50ms and performs no I/O
 - ◆ The scheduler runs A first, then B after

Incorporating I/O (Cont.)



Poor Use of Resources



Overlap Allows Better Use of Resources

**Maximize the
CPU utilization**

Incorporating I/O (Cont.)

- When a job initiates an I/O request.
 - ◆ The job is blocked waiting for I/O completion.
 - ◆ The scheduler should schedule another job on the CPU.

- When the I/O completes
 - ◆ An interrupt is raised.
 - ◆ The OS moves the process from blocked back to the ready state.

Part II: Scheduling: The Multi-Level Feedback Queue

Multi-Level Feedback Queue (MLFQ)

- A Scheduler that learns from the past to predict the future.
- Objective:
 - ◆ Optimize **turnaround time** → Run shorter jobs first
 - ◆ Minimize **response time** without *a priori knowledge of job length*.

MLFQ: Basic Rules

- MLFQ has a number of distinct **queues**.
 - ◆ Each queue is assigned a different priority level.

- A job that is ready to run is on a single queue.
 - ◆ A job **on a higher queue** is chosen to run.
 - ◆ Use round-robin scheduling among jobs in the same queue


Rule 1: If $\text{Priority}(A) > \text{Priority}(B)$, A runs (B doesn't).

Rule 2: If $\text{Priority}(A) = \text{Priority}(B)$, A & B run in RR.

MLFQ: Basic Rules (Cont.)

- ▣ MLFQ varies the priority of a job based on its observed behavior.
- ▣ Example:
 - ◆ A job repeatedly relinquishes the CPU while waiting IOs → Keep its priority high
 - ◆ A job uses the CPU intensively for long periods of time → Reduce its priority.

MLFQ Example




MLFQ: How to Change Priority

- MLFQ priority adjustment algorithm:
 - ◆ **Rule 3:** When a job enters the system, it is placed at the highest priority
 - ◆ **Rule 4a:** If a job uses up an entire time slice while running, its priority is reduced (i.e., it moves down on queue).
 - ◆ **Rule 4b:** If a job gives up the CPU before the time slice is up, it stays at the same priority level

In this manner, MLFQ approximates SJF

Example 1: A Single Long-Running Job


- ▣ A three-queue scheduler with time slice 10ms



Example 2: Along Came a Short Job

□ Assumption:


- ◆ **Job A:** A long-running CPU-intensive job
- ◆ **Job B:** A short-running interactive job (20ms runtime)
- ◆ A has been running for some time, and then B arrives at time $T=100$.



Example 3: What About I/O?

□ Assumption:

- ◆ **Job A:** A long-running CPU-intensive job
- ◆ **Job B:** An interactive job that need the CPU only for 1ms before performing an I/O



A Mixed I/O-intensive and CPU-intensive Workload (msec)

The MLFQ approach keeps an interactive job at the highest priority

Problems with the Basic MLFQ

❑ Starvation

- ◆ If there are “too many” interactive jobs in the system.
- ◆ Long-running jobs will never receive any CPU time.

❑ Game the scheduler

- ◆ After running 99% of a time slice, issue an I/O operation.
- ◆ The job gains a higher percentage of CPU time.

❑ A program may change its behavior over time.



- ◆ CPU bound process → I/O bound process

The Priority Boost




- ❑ **Rule 5:** After some time period S , move all the jobs in the system to the topmost queue.

- ◆ Example:

- A long-running job(A) with two short-running interactive job(B, C)





Without(Left) and With(Right) Priority Boost

A:  B:  C: 

Better Accounting

- How to prevent gaming of our scheduler?
- Solution:
 - ◆ **Rule 4** (Rewrite Rules 4a and 4b): Once a job **uses up its time allotment** at a given level (regardless of how many times it has given up the CPU), **its priority is reduced**(i.e., it moves down on queue).




Without(Left) and With(Right) Gaming Tolerance

Tuning MLFQ And Other Issues

Lower Priority, Longer Quanta

- ◆ The high-priority queues → Short time slices
 - E.g., 10 or fewer milliseconds
- ◆ The Low-priority queue → Longer time slices
 - E.g., 100 milliseconds



Example) 10ms for the highest queue, 20ms for the middle, 40ms for the lowest

The Solaris MLFQ implementation

- ▣ For the Time-Sharing scheduling class (TS)
 - ◆ 60 Queues
 - ◆ Slowly increasing time-slice length
 - The highest priority: 20msec
 - The lowest priority: A few hundred milliseconds
 - ◆ Priorities boosted around every 1 second or so.

MLFQ: Summary

- The refined set of MLFQ rules:
 - ◆ **Rule 1:** If $\text{Priority}(A) > \text{Priority}(B)$, A runs (B doesn't).
 - ◆ **Rule 2:** If $\text{Priority}(A) = \text{Priority}(B)$, A & B run in RR.
 - ◆ **Rule 3:** When a job enters the system, it is placed at the highest priority.
 - ◆ **Rule 4:** Once a job uses up its time allotment at a given level (regardless of how many times it has given up the CPU), its priority is reduced(i.e., it moves down on queue).
 - ◆ **Rule 5:** After some time period S, move all the jobs in the system to the topmost queue.

Part III: Scheduling: Proportional Share

Proportional Share Scheduler

- Fair-share scheduler
 - ◆ Guarantee that each job obtain *a certain percentage* of CPU time.
 - ◆ Not optimized for turnaround or response time

Basic Concept

□ Tickets

- ◆ Represent the share of a resource that a process should receive
- ◆ The percent of tickets represents its share of the system resource in question.

□ Example

- ◆ There are two processes, A and B.
 - Process A has 75 tickets → receive 75% of the CPU
 - Process B has 25 tickets → receive 25% of the CPU

Lottery scheduling

- ❑ The scheduler picks a winning ticket.
 - ◆ Load the state of that *winning process* and runs it.

- ❑ Example

- ◆ There are 100 tickets
 - Process A has 75 tickets: 0 ~ 74
 - Process B has 25 tickets: 75 ~ 99

Scheduler's winning tickets: 63 85 70 39 76 17 29 41 36 39 10 99 68 83 63

Resulting scheduler: A B A A B A A A A A A B A B A

**The longer these two jobs compete,
The more likely they are to achieve the desired percentages.**

Ticket Mechanisms

□ Ticket currency

- ◆ A user allocates tickets among their own jobs in whatever currency they would like.
- ◆ The system converts the currency into the correct global value.
- ◆ Example
 - There are 200 tickets (Global currency)
 - Process A has 100 tickets
 - Process B has 100 tickets

User A $\rightarrow 500$ (A's currency) to A1 $\rightarrow 50$ (global currency)
 $\rightarrow 500$ (A's currency) to A2 $\rightarrow 50$ (global currency)

User B $\rightarrow 10$ (B's currency) to B1 $\rightarrow 100$ (global currency)

Ticket Mechanisms (Cont.)

□ Ticket transfer


- ◆ A process can temporarily hand off *its tickets* to another process.

□ Ticket inflation

- ◆ A process can temporarily raise or lower the number of tickets it owns.
- ◆ If any one process needs *more CPU time*, it can boost its tickets.

Implementation

- ❑ Example: There are three processes, A, B, and C.
 - ◆ Keep the processes in a list:




```
1      // counter: used to track if we've found the winner yet
2      int counter = 0;
3
4      // winner: use some call to a random number generator to
5      // get a value, between 0 and the total # of tickets
6      int winner = getRandom(0, totaltickets);
7
8      // current: use this to walk through the list of jobs
9      node_t *current = head;
10
11     // loop until the sum of ticket values is > the winner
12     while (current) {
13         counter = counter + current->tickets;
14         if (counter > winner)
15             break; // found the winner
16         current = current->next;
17     }
18     // 'current' is the winner: schedule it...
```


Implementation (Cont.)

- U: unfairness metric
 - ◆ The time the first job completes divided by the time that the second job completes.
- Example:
 - ◆ There are two jobs, each jobs has runtime 10.
 - First job finishes at time 10
 - Second job finishes at time 20
 - ◆ $U = \frac{10}{20} = 0.5$
 - ◆ U will be close to 1 when both jobs finish at nearly the same time.

Lottery Fairness Study

- There are two jobs.
 - ◆ Each jobs has the same number of tickets (100).



**When the job length is not very long,
average unfairness can be quite severe.**

Stride Scheduling

▣ Stride of each process

- ◆ (A large number) / (the number of tickets of the process)
- ◆ Example: A large number = 10,000
 - Process A has 100 tickets → stride of A is 100
 - Process B has 50 tickets → stride of B is 200

▣ A process runs, increment a counter(=pass value) for it by its stride.

- ◆ Pick the process to run that has the lowest pass value

```
current = remove_min(queue);           // pick client with minimum pass
schedule(current);                     // use resource for quantum
current->pass += current->stride;        // compute next pass using stride
insert(queue, current);                // put back into the queue
```

A pseudo code implementation

Stride Scheduling Example

Pass(A) (stride=100)	Pass(B) (stride=200)	Pass(C) (stride=40)	Who Runs?
0	0	0	A
100	0	0	B
100	200	0	C
100	200	40	C
100	200	80	C
100	200	120	A
200	200	120	C
200	200	160	C
200	200	200	...

If new job enters with pass value 0,
It will **monopolize** the CPU!

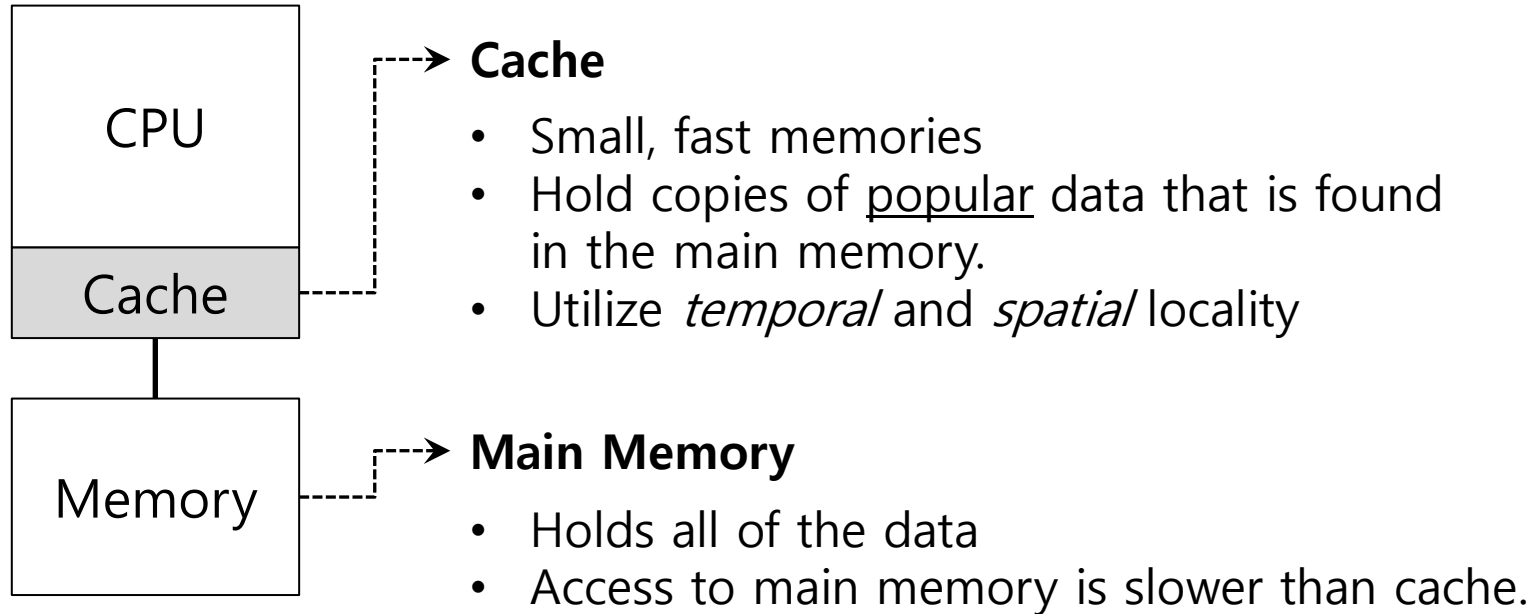
Part IV. Multiprocessor Scheduling

Multiprocessor Scheduling

- ❑ The rise of the **multicore processor** is the source of multiprocessor-scheduling proliferation.
 - ◆ **Multicore:** Multiple CPU cores are packed onto a single chip.
- ❑ Adding more CPUs does not make that single application run faster. → You'll have to rewrite application to run in parallel, using **threads**.

How to schedule jobs on **Multiple CPUs?**

Single CPU with cache




By keeping data in cache, the system can make slow memory
appear to be a fast one

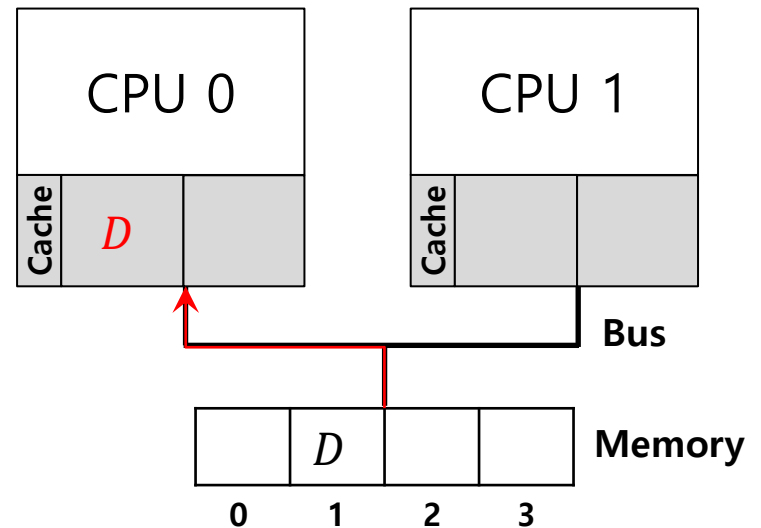
Cache coherence

- Consistency of shared resource data stored in multiple caches.

0. Two CPUs with caches sharing memory

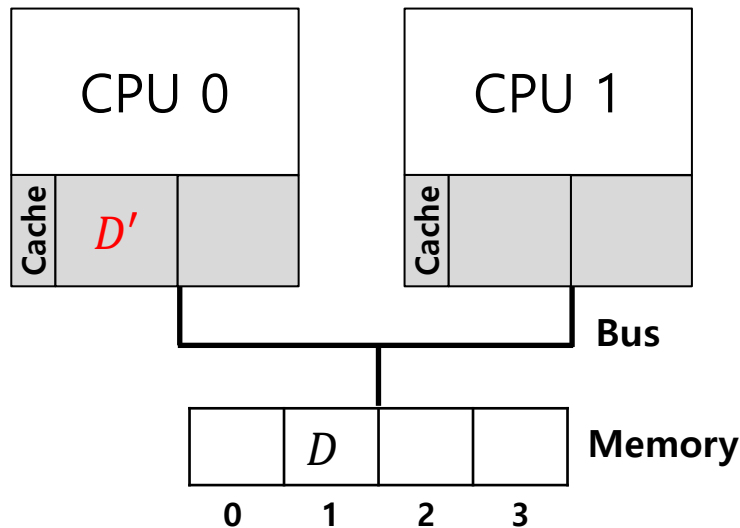


1. CPU0 reads a data at address 1.

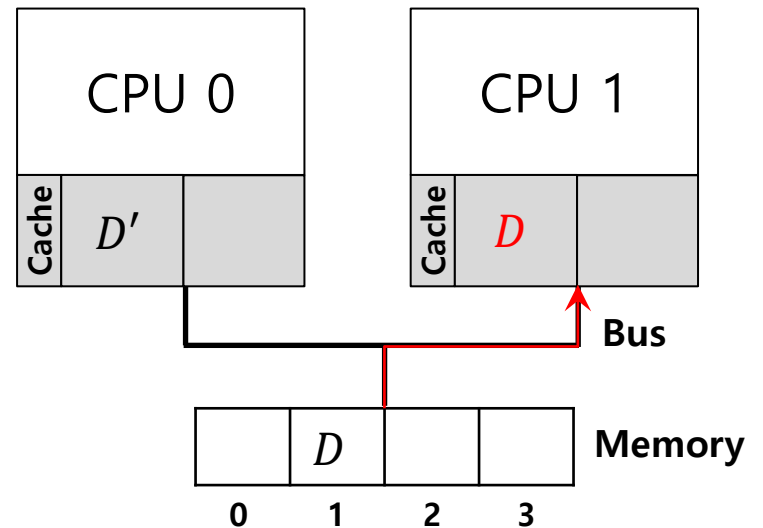


Cache coherence (Cont.)

2. D is updated and CPU1 is scheduled.



3. CPU1 re-reads the value at address A



CPU1 gets the D value instead of the correct value D' .

Cache coherence solution

□ Bus snooping

- ◆ Each cache pays attention to memory updates by **observing the bus**.
- ◆ When a CPU sees an update for a data item it holds in its cache, it will notice the change and either invalidate its copy or update it.

Don't forget synchronization

- When accessing shared data across CPUs, **mutual exclusion** primitives should likely be used to guarantee correctness.

```
1      typedef struct __Node_t {
2          int value;
3          struct __Node_t *next;
4      } Node_t;
5
6      int List_Pop() {
7          Node_t *tmp = head;           // remember old head ...
8          int value = head->value;       // ... and its value
9          head = head->next;             // advance head to next pointer
10         free(tmp);                    // free old head
11         return value;                 // return value at head
12     }
```

Simple List Delete Code

Don't forget synchronization (Cont.)

▣ Solution

```
1      pthread_mutex_t m;  
2      typedef struct __Node_t {  
3          int value;  
4          struct __Node_t *next;  
5      } Node_t;  
6  
7      int List_Pop() {  
8          lock(&m)  
9          Node_t *tmp = head;           // remember old head ...  
10         int value = head->value;       // ... and its value  
11         head = head->next;             // advance head to next pointer  
12         free(tmp);                     // free old head  
13         unlock(&m)  
14         return value;                  // return value at head  
15     }
```

Simple List Delete Code with lock

Cache Affinity


- ❑ Keep a process on **the same CPU** if at all possible
 - ◆ A process builds up a fair bit of state in the cache of a CPU.
 - ◆ The next time the process run, it will run faster if some of its state is *already present* in the cache on that CPU.

A multiprocessor scheduler should consider **cache affinity when making its scheduling decision.**

Single queue Multiprocessor Scheduling (SQMS)

- Put all jobs that need to be scheduled into a single queue.
 - ◆ Each CPU simply picks the next job from the globally shared queue.
 - ◆ Cons:



- Some form of **locking** have to be inserted → Lack of scalability
- Cache affinity
- Example:



- Possible job scheduler across CPUs:

CPU0	A	E	D	C	B	... (repeat) ...
CPU1	B	A	E	D	C	... (repeat) ...
CPU2	C	B	A	E	D	... (repeat) ...
CPU3	D	C	B	A	E	... (repeat) ...

Scheduling Example with Cache affinity




- ◆ Preserving affinity for most
 - Jobs A through D are not moved across processors.
 - Only job e Migrating from CPU to CPU.
- ◆ Implementing such a scheme can be **complex**.

Multi-queue Multiprocessor Scheduling (MQMS)

- MQMS consists of **multiple scheduling queues**.
 - ◆ Each queue will follow a particular scheduling discipline.
 - ◆ When a job enters the system, it is placed on **exactly one** scheduling queue.
 - ◆ Avoid the problems of information sharing and synchronization.

MQMS Example

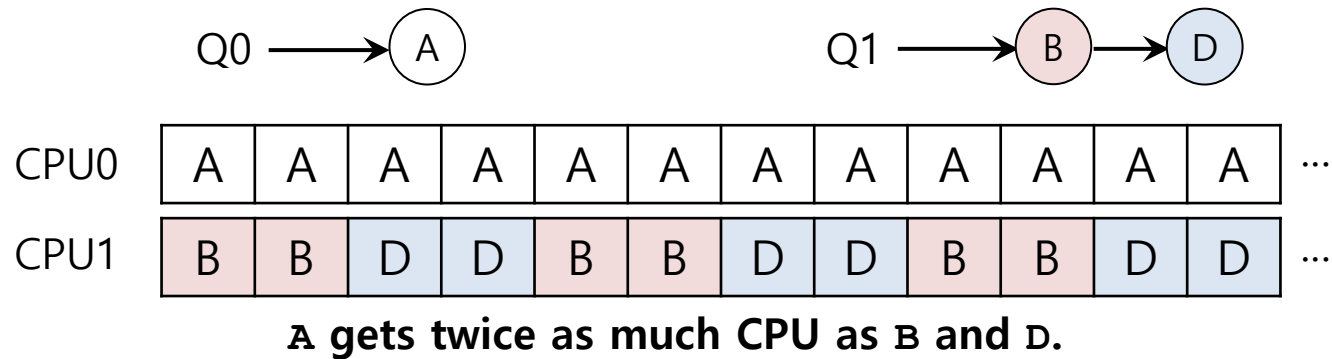
- With **round robin**, the system might produce a schedule that looks like this:




MQMS provides more **scalability** and **cache affinity**.

Load Imbalance issue of MQMS

- After job C in Q0 finishes:





- After job A in Q0 finishes:




How to deal with load imbalance?

- The answer is to move jobs (**Migration**).

- ◆ Example:



The OS moves one of B or D to CPU 0



Or




How to deal with load imbalance? (Cont.)

- A more tricky case:



- A possible migration pattern:

- ◆ Keep switching jobs



Work Stealing

□ Move jobs between queues

◆ Implementation:

- A source queue that is low on jobs is picked.
- The source queue occasionally peeks at another target queue.
- If the target queue is more full than the source queue, the source will “**steal**” one or more jobs from the target queue.

◆ Cons:

- *High overhead and trouble scaling*

Linux Multiprocessor Schedulers

▣ O(1)

- ◆ A Priority-based scheduler
- ◆ Use Multiple queues
- ◆ Change a process's priority over time
- ◆ Schedule those with highest priority
- ◆ Interactivity is a particular focus

▣ Completely Fair Scheduler (CFS)

- ◆ Deterministic proportional-share approach
- ◆ Multiple queues

Linux Multiprocessor Schedulers (Cont.)

- BF Scheduler (BFS)
 - ◆ A single queue approach
 - ◆ Proportional-share
 - ◆ Based on Earliest Eligible Virtual Deadline First (EEVDF)

Summary

□ CPU Scheduling

- ◆ Scheduling metrics: Turnaround time, Response time; Fairness
- ◆ Scheduling Strategy: FIFO, SJF, PSJF, RR, MLFQ
 - Simple: FIFO, SJF, PSJF, RR
 - MLFQ
 - Proportional Share
- ◆ Multiprocessor scheduling
 - Cache coherence
 - Cache affinity
 - SQMS, MQMS, Load balance

□ Next: Memory ([Chapters 13](#), [15](#), [16](#), [17](#), [18](#), [19](#), [20](#), [21](#), [22](#), [23](#))