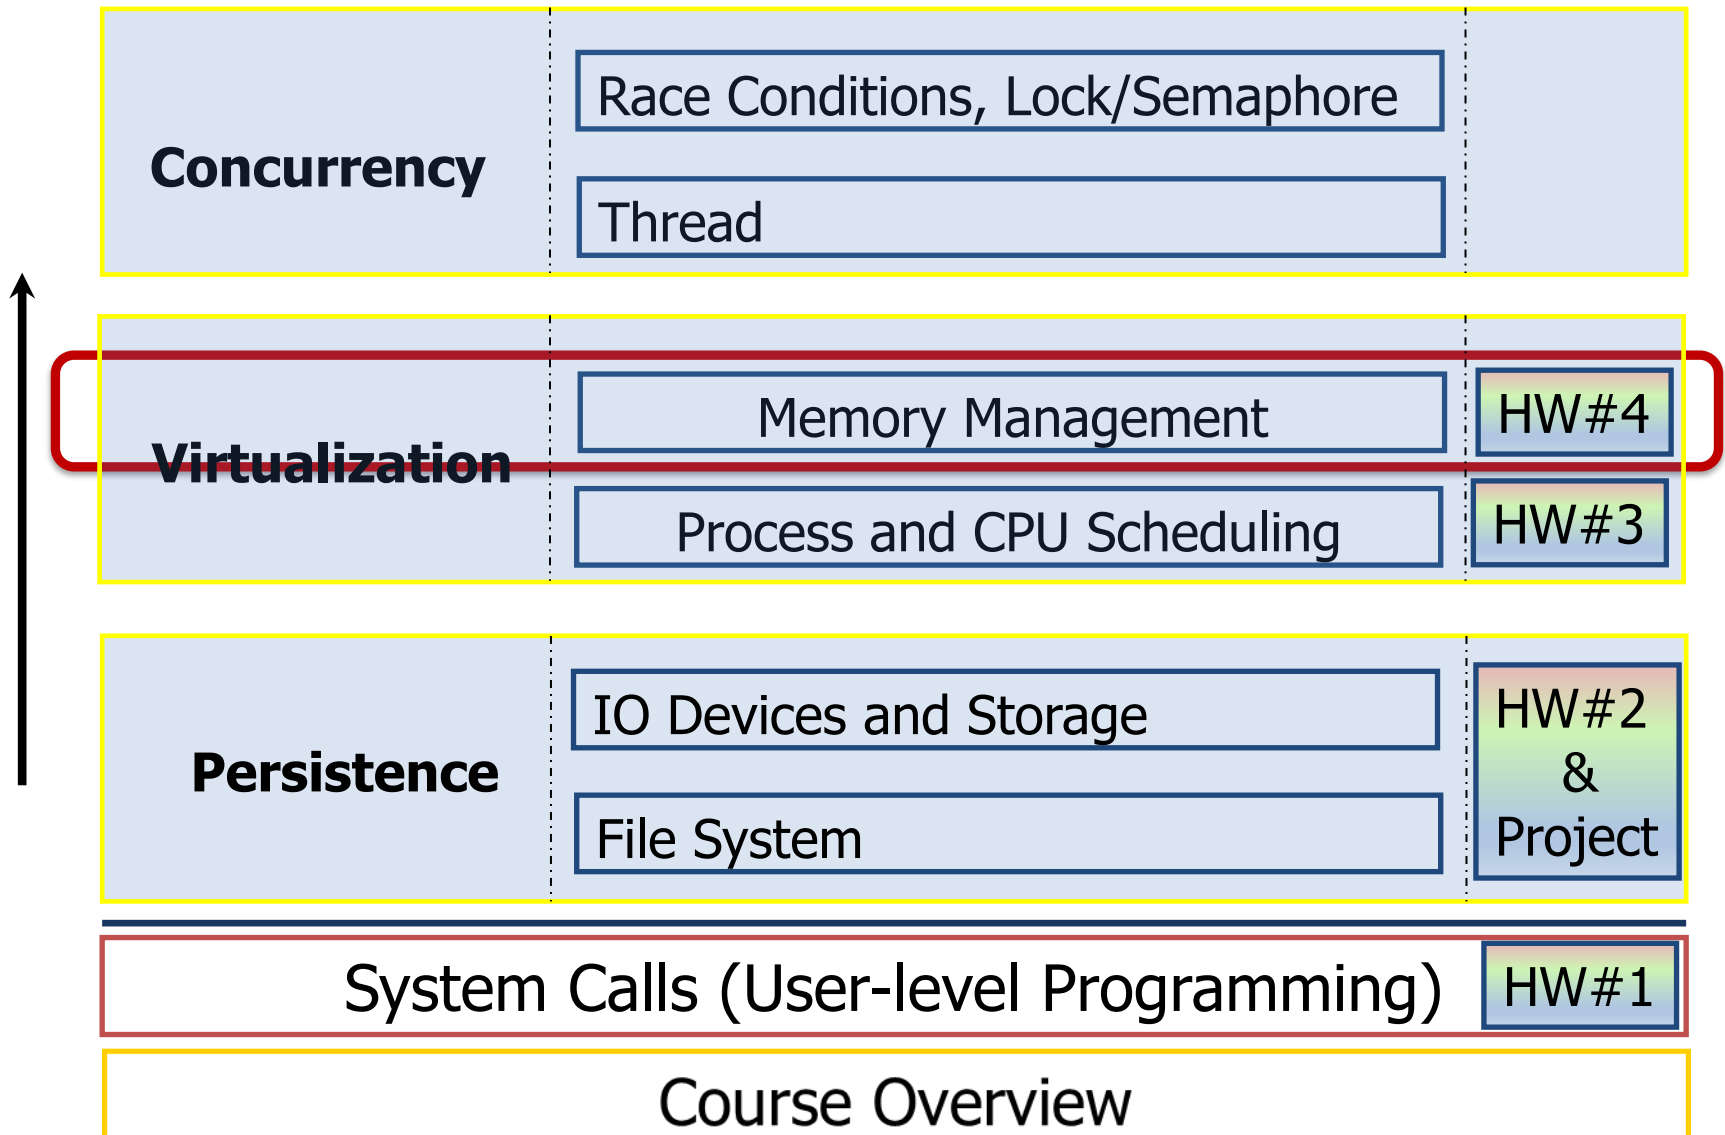


Lecture 9: Virtualizing Memory – Address Space and Address Translation

The Course Organization (Bottom-up)

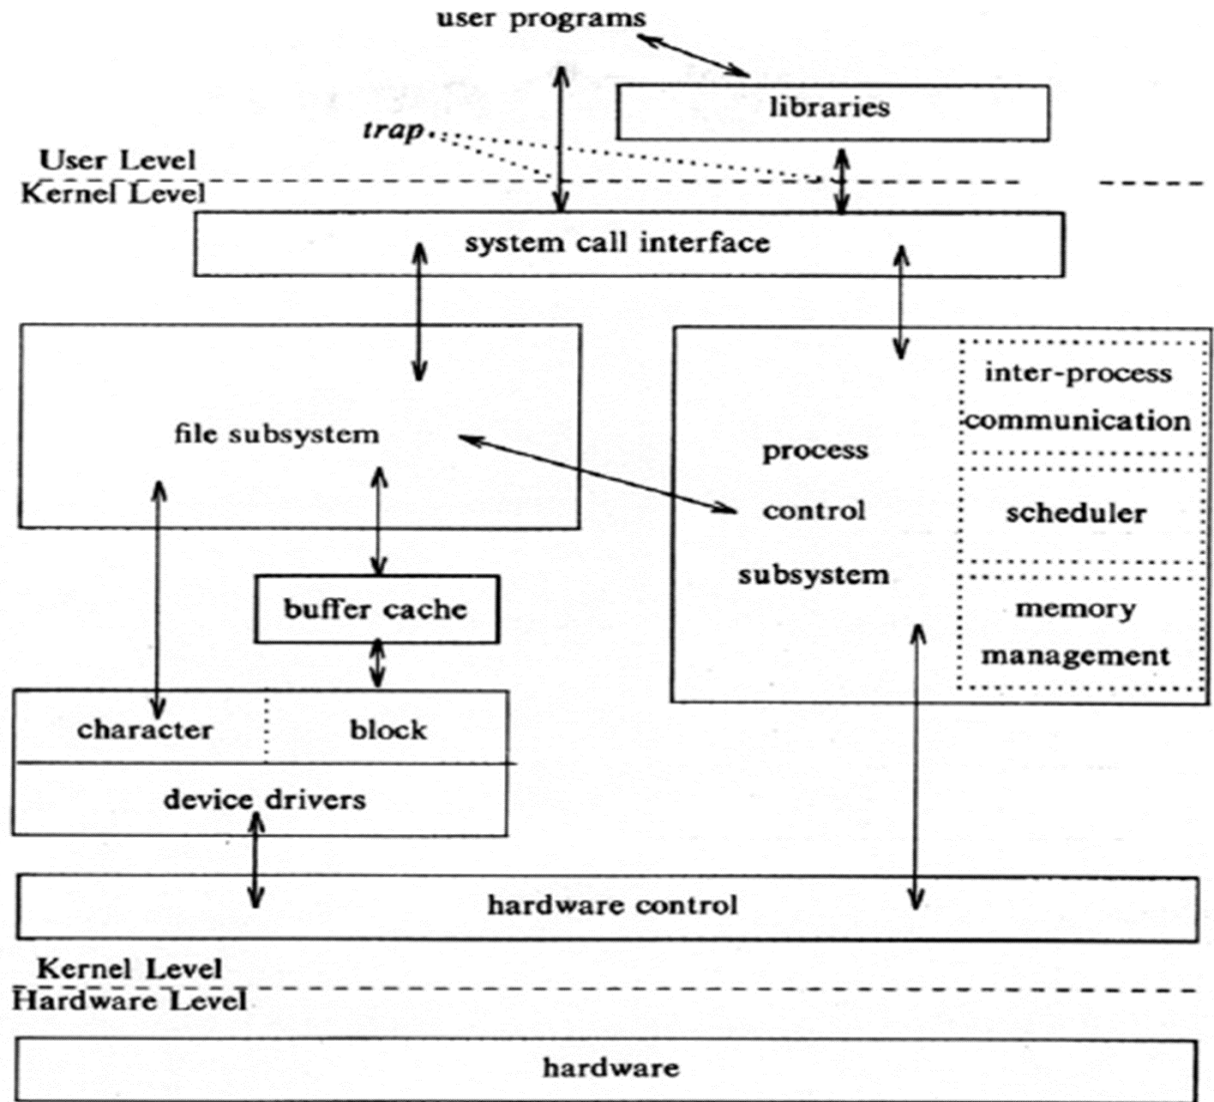


OS – Resource management via virtualization

OS provides services via **System Call** (typically a few hundred) to run **process**, access memory/devices/files, etc.

The OS **manages resources** such as *CPU*, *memory* and *disk* via **virtualization**.

- many programs to run (processes) → Sharing the CPU
- many processes to *concurrently* access their own instructions and data → Sharing memory
- many processes to access devices → Sharing disks



The Design Of The Unix Operating System (Maurice Bach, 1986)

Part I: Address Space

Memory Virtualization

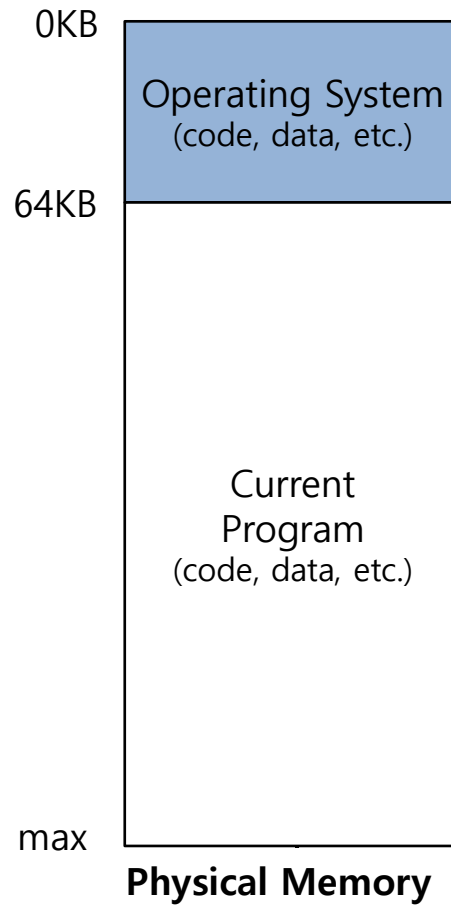
- What is **memory virtualization**?
 - ◆ OS virtualizes its physical memory.
 - ◆ OS provides an **illusion memory space** per each process.
 - ◆ It seems to be seen like **each process uses the whole memory** .

Benefit of Memory Virtualization

- ❑ Ease of use in programming
- ❑ Memory efficiency in terms of time and space
- ❑ The guarantee of isolation for processes as well as OS
 - ◆ Protection from errant accesses of other processes

OS in The Early System

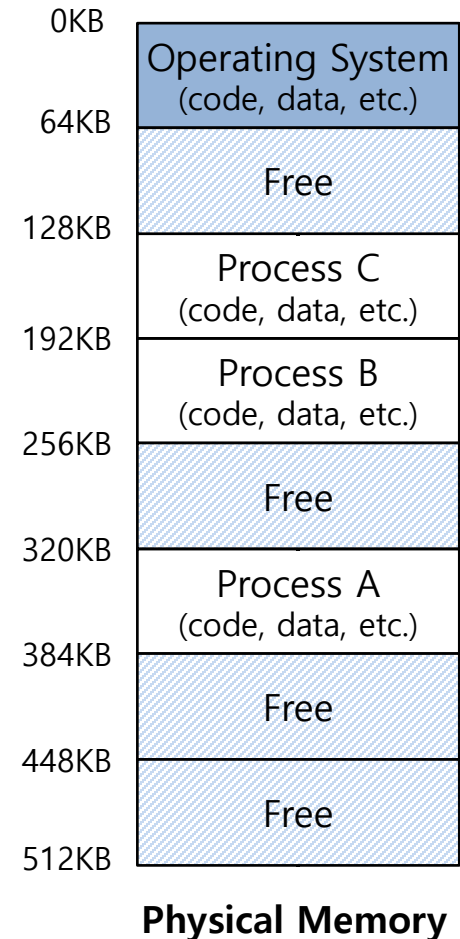
- ▣ Load only one process in memory.
 - ◆ Poor utilization and efficiency



Multiprogramming and Time Sharing


- ❑ **Load multiple processes** in memory.
 - ◆ Execute one for a short while.
 - ◆ Switch processes between them in memory.
 - ◆ Increase utilization and efficiency.

- ❑ Cause an important **protection issue**.
 - ◆ Errant memory accesses from other processes




Address Space

- ❑ OS creates an **abstraction** of physical memory.
 - ◆ The address space contains all about a running process.
 - ◆ That is consist of program code, data, heap, and stack





Address Space(Cont.)

- Code
 - ◆ Where instructions live
- Data
 - ◆ Global or static local variables
- Heap
 - ◆ Dynamically allocate memory.
 - `malloc` in C language
 - `new` in object-oriented language
- Stack
 - ◆ Store return addresses or values.
 - ◆ Contain local variables arguments to routines.



Address Space

Example from Linux



Virtual Address

- ❑ **Every address** in a running program is virtual.
 - ◆ OS translates the virtual address to physical address

```
#include <stdio.h>
#include <stdlib.h>

int main(int argc, char *argv[]){

    printf("location of code   : %p\n", (void *) main);
    printf("location of heap   : %p\n", (void *) malloc(1));
    int x = 3;
    printf("location of stack  : %p\n", (void *) &x);

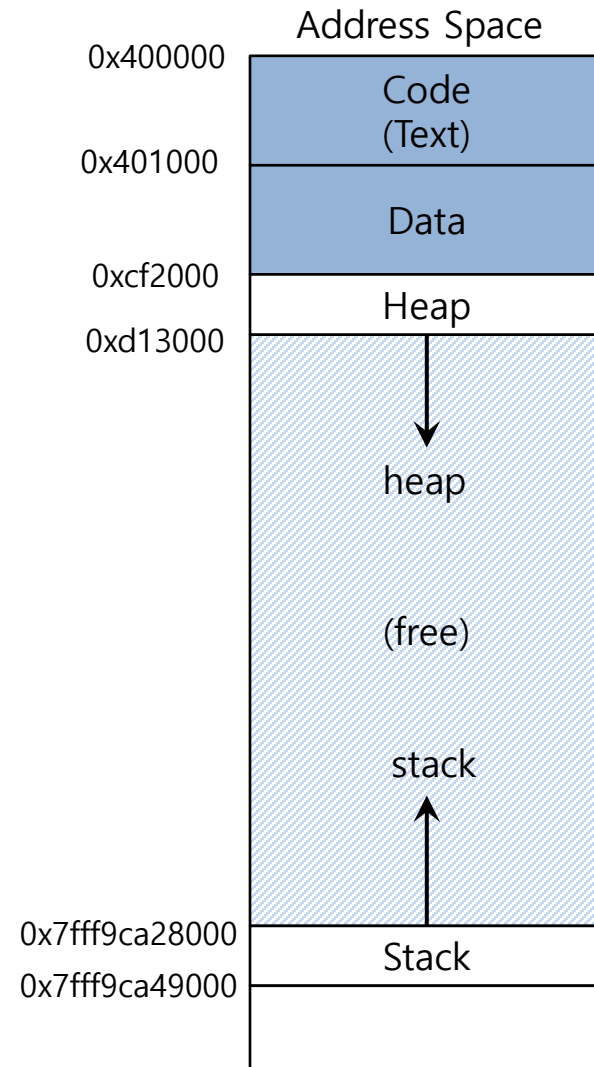
    return x;
}
```

A simple program that prints out addresses

Virtual Address(Cont.)

▣ The output in 64-bit Linux machine

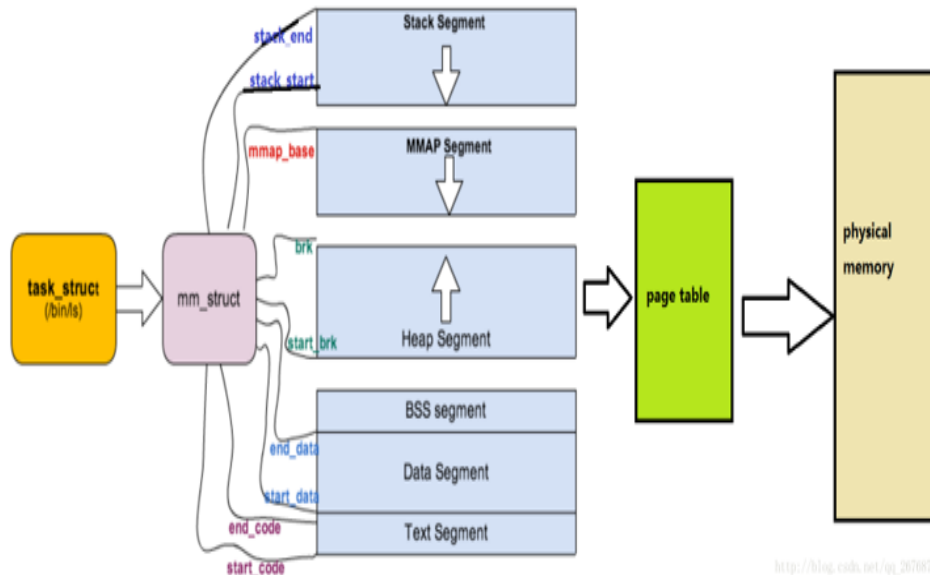
```
location of code   : 0x40057d
location of heap   : 0xcf2010
location of stack  : 0x7fff9ca45fcc
```



Example: pmap and cat /proc/<pid>/maps

cat /proc/<pid_number>/maps

```
csc3150@csc3150-VirtualBox:~/process-memory-test$ more /proc/10534/maps
004cd000-004ce000 r-xp 00000000 08:01 136307 /home/csc3150/process-memory-test/test
004ce000-004cf000 r--p 00000000 08:01 136307 /home/csc3150/process-memory-test/test
004cf000-004d0000 rw-p 00001000 08:01 136307 /home/csc3150/process-memory-test/test
b7d93000-b7f68000 r-xp 00000000 08:01 526179 /lib/i386-linux-gnu/libc-2.27.so
b7f68000-b7f69000 ---p 001d5000 08:01 526179 /lib/i386-linux-gnu/libc-2.27.so
b7f69000-b7f6b000 r--p 001d5000 08:01 526179 /lib/i386-linux-gnu/libc-2.27.so
b7f6b000-b7f6c000 rw-p 001d7000 08:01 526179 /lib/i386-linux-gnu/libc-2.27.so
b7f6c000-b7f6f000 rw-p 00000000 00:00 0
b7f83000-b7f85000 rw-p 00000000 00:00 0
b7f85000-b7f88000 r--p 00000000 00:00 0 [vvar]
b7f88000-b7f8a000 r-xp 00000000 00:00 0 [vdso]
b7f8a000-b7fb0000 r-xp 00000000 08:01 526151 /lib/i386-linux-gnu/ld-2.27.so
b7fb0000-b7fb1000 r--p 00025000 08:01 526151 /lib/i386-linux-gnu/ld-2.27.so
b7fb1000-b7fb2000 rw-p 00026000 08:01 526151 /lib/i386-linux-gnu/ld-2.27.so
bf829000-bf84a000 rw-p 00000000 00:00 0 [stack]
```



pmap -x <pid_number>

```
csc3150@csc3150-VirtualBox:~/process-memory-test$ pmap -x 10534
10534: ./test
Address      Kbytes      RSS      Dirty Mode  Mapping
004cd000      4           4         0 r-x-- test
004cd000      0           0         0 r-x-- test
004ce000      4           4         4 r--- test
004ce000      0           0         0 r--- test
004cf000      4           4         4 rw--- test
004cf000      0           0         0 rw--- test
b7d93000    1876        688         0 r-x-- libc-2.27.so
b7d93000      0           0         0 r-x-- libc-2.27.so
b7f68000      4           0         0 ----- libc-2.27.so
b7f68000      0           0         0 ----- libc-2.27.so
b7f69000      8           8         8 r--- libc-2.27.so
b7f69000      0           0         0 r--- libc-2.27.so
b7f6b000      4           4         4 rw--- libc-2.27.so
b7f6b000      0           0         0 rw--- libc-2.27.so
b7f6c000     12           8         8 rw--- [ anon ]
b7f6c000      0           0         0 rw--- [ anon ]
b7f83000      8           8         8 rw--- [ anon ]
b7f83000      0           0         0 rw--- [ anon ]
b7f85000     12           0         0 r--- [ anon ]
b7f85000      0           0         0 r--- [ anon ]
b7f88000      8           4         0 r-x-- [ anon ]
b7f88000      0           0         0 r-x-- [ anon ]
b7f8a000     152        152         0 r-x-- ld-2.27.so
b7f8a000      0           0         0 r-x-- ld-2.27.so
b7fb0000      4           4         4 r--- ld-2.27.so
b7fb0000      0           0         0 r--- ld-2.27.so
b7fb1000      4           4         4 rw--- ld-2.27.so
b7fb1000      0           0         0 rw--- ld-2.27.so
bf829000     132        12        12 rw--- [ stack ]
bf829000      0           0         0 rw--- [ stack ]
-----
total kB      2236        904        56
```

Part II: Address Translation

Memory Virtualizing with Efficiency and Control

- ❑ Memory virtualizing takes a similar strategy known as **limited direct execution(LDE)** for efficiency and control.
- ❑ In memory virtualizing, efficiency and control are attained by **hardware support**.
 - ◆ e.g., registers, TLB(Translation Look-aside Buffer)s, page-table

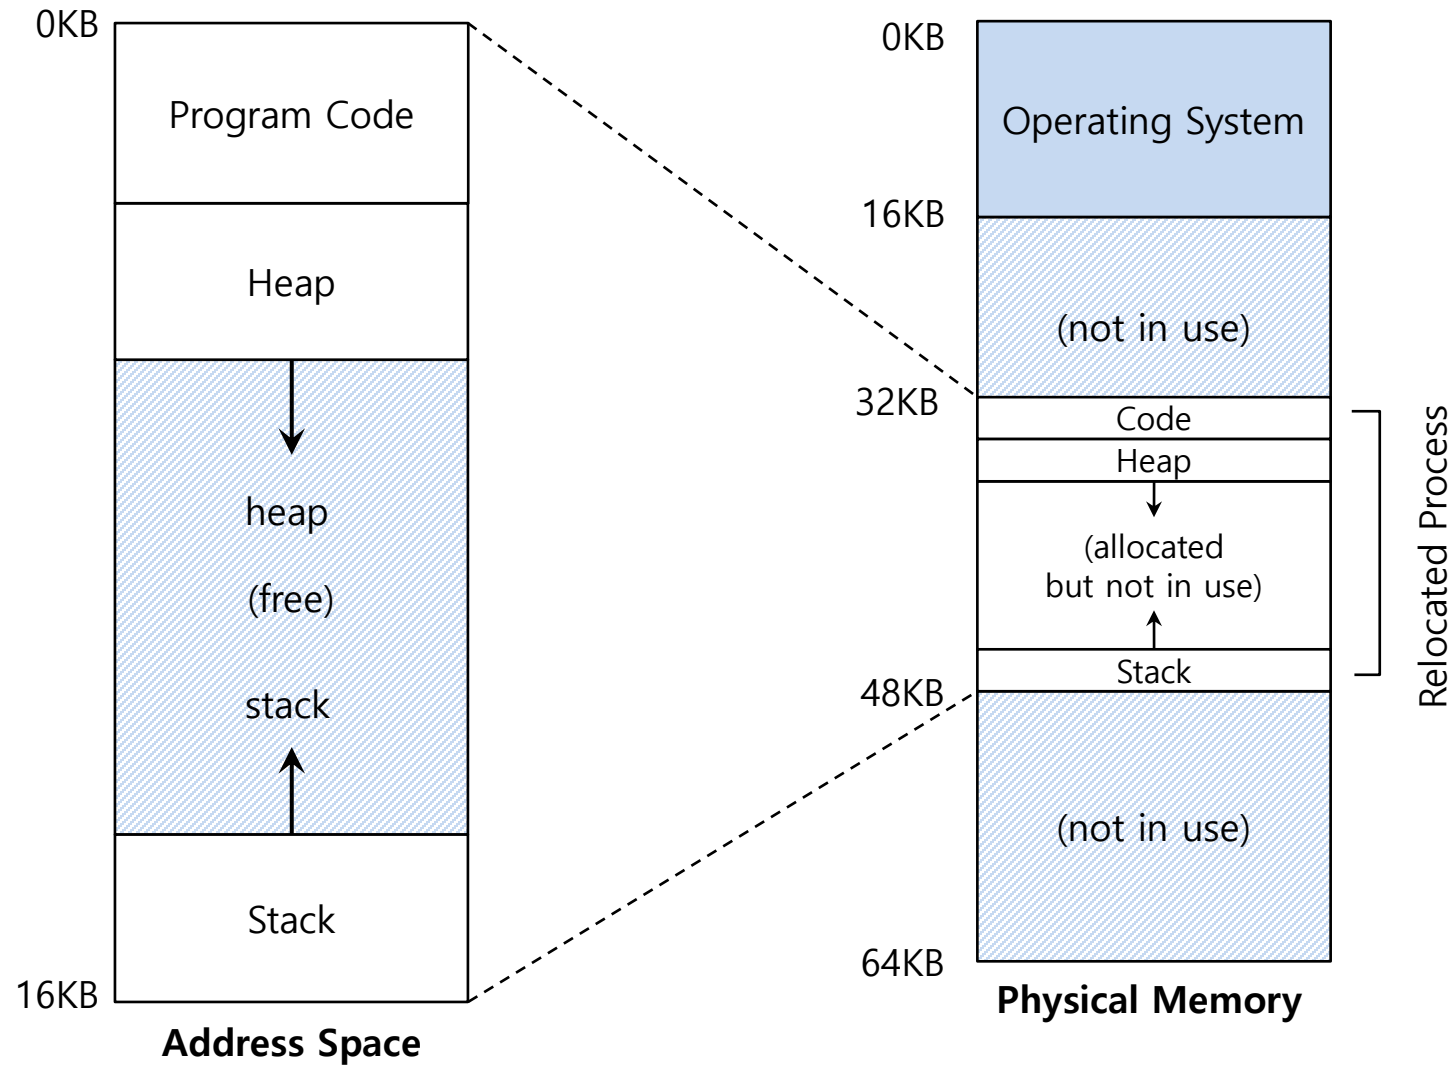
Address Translation

- ❑ Hardware transforms a **virtual address** to a **physical address**.
 - ◆ The desired information is actually stored in a physical address.
- ❑ The OS must get involved at key points to set up the hardware.
 - ◆ The OS must manage memory to judiciously intervene.

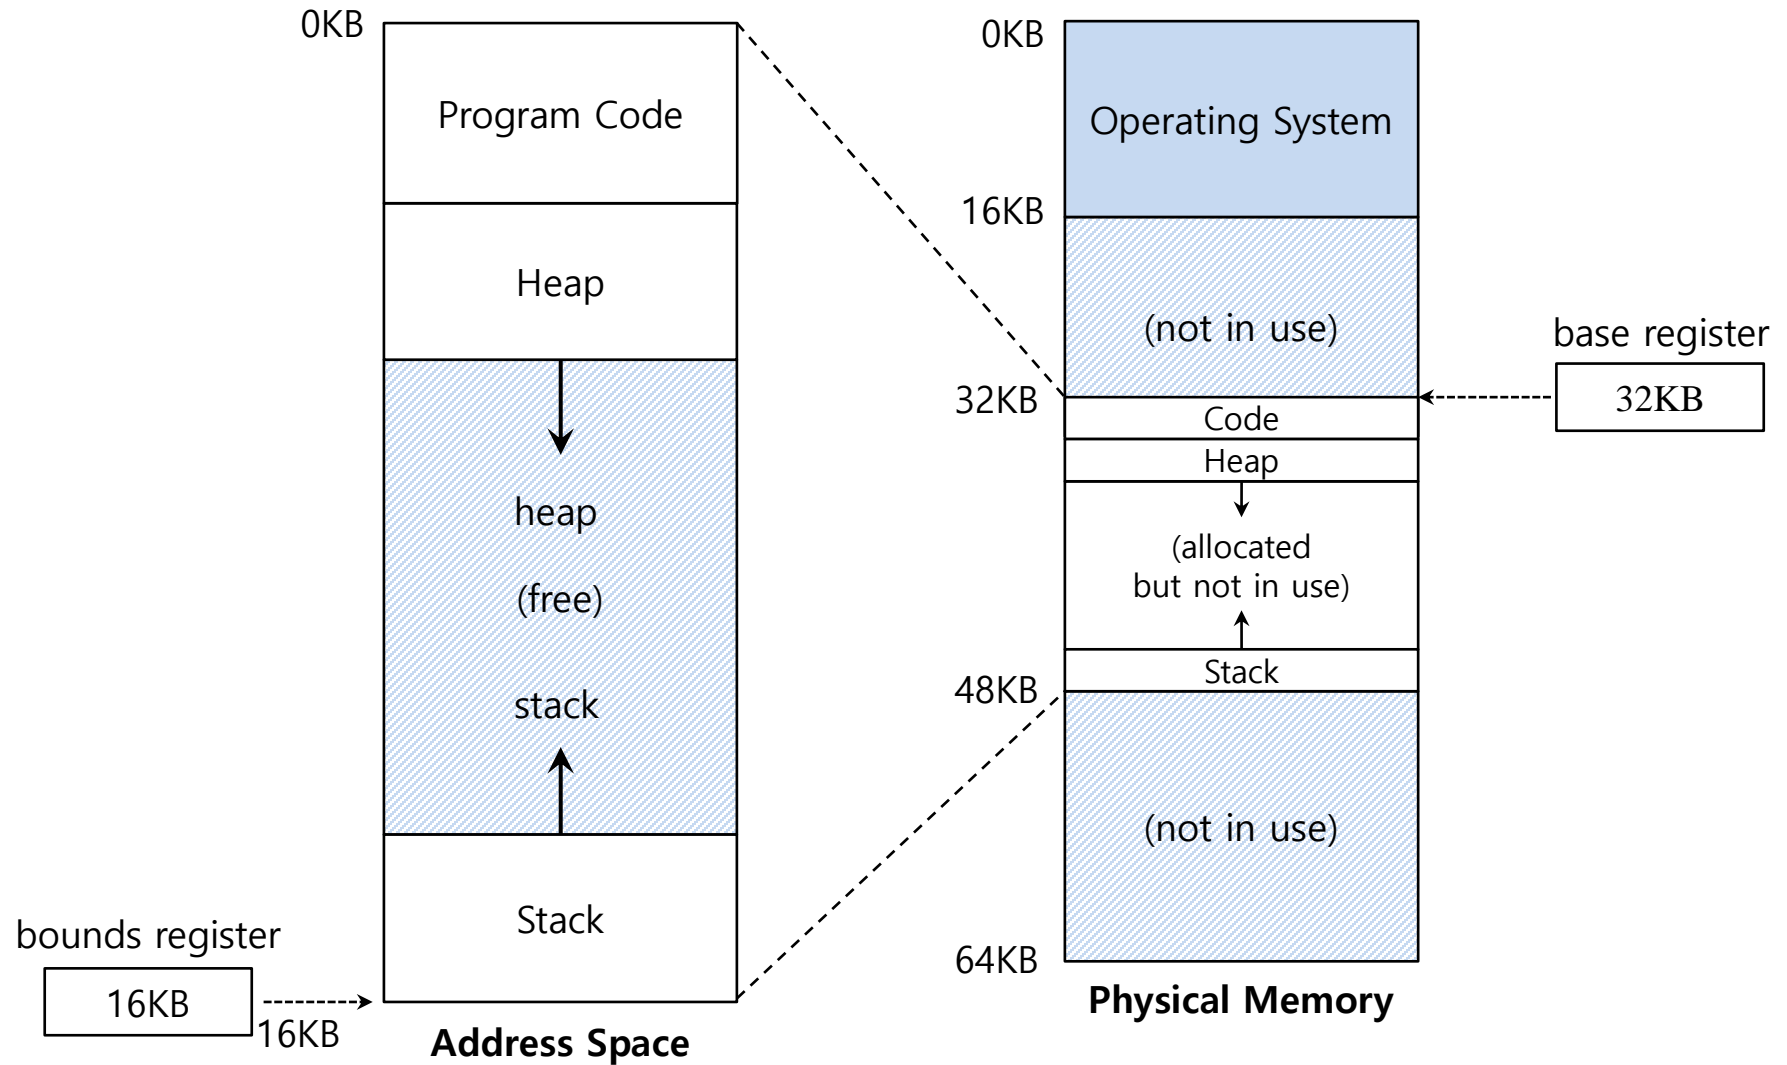
Relocation Address Space

- ❑ The address space start at address 0.
- ❑ The OS wants to place the process **somewhere else** in physical memory, not at address 0.

A Single Relocated Process



Base and Bounds Register



Dynamic(Hardware base) Relocation

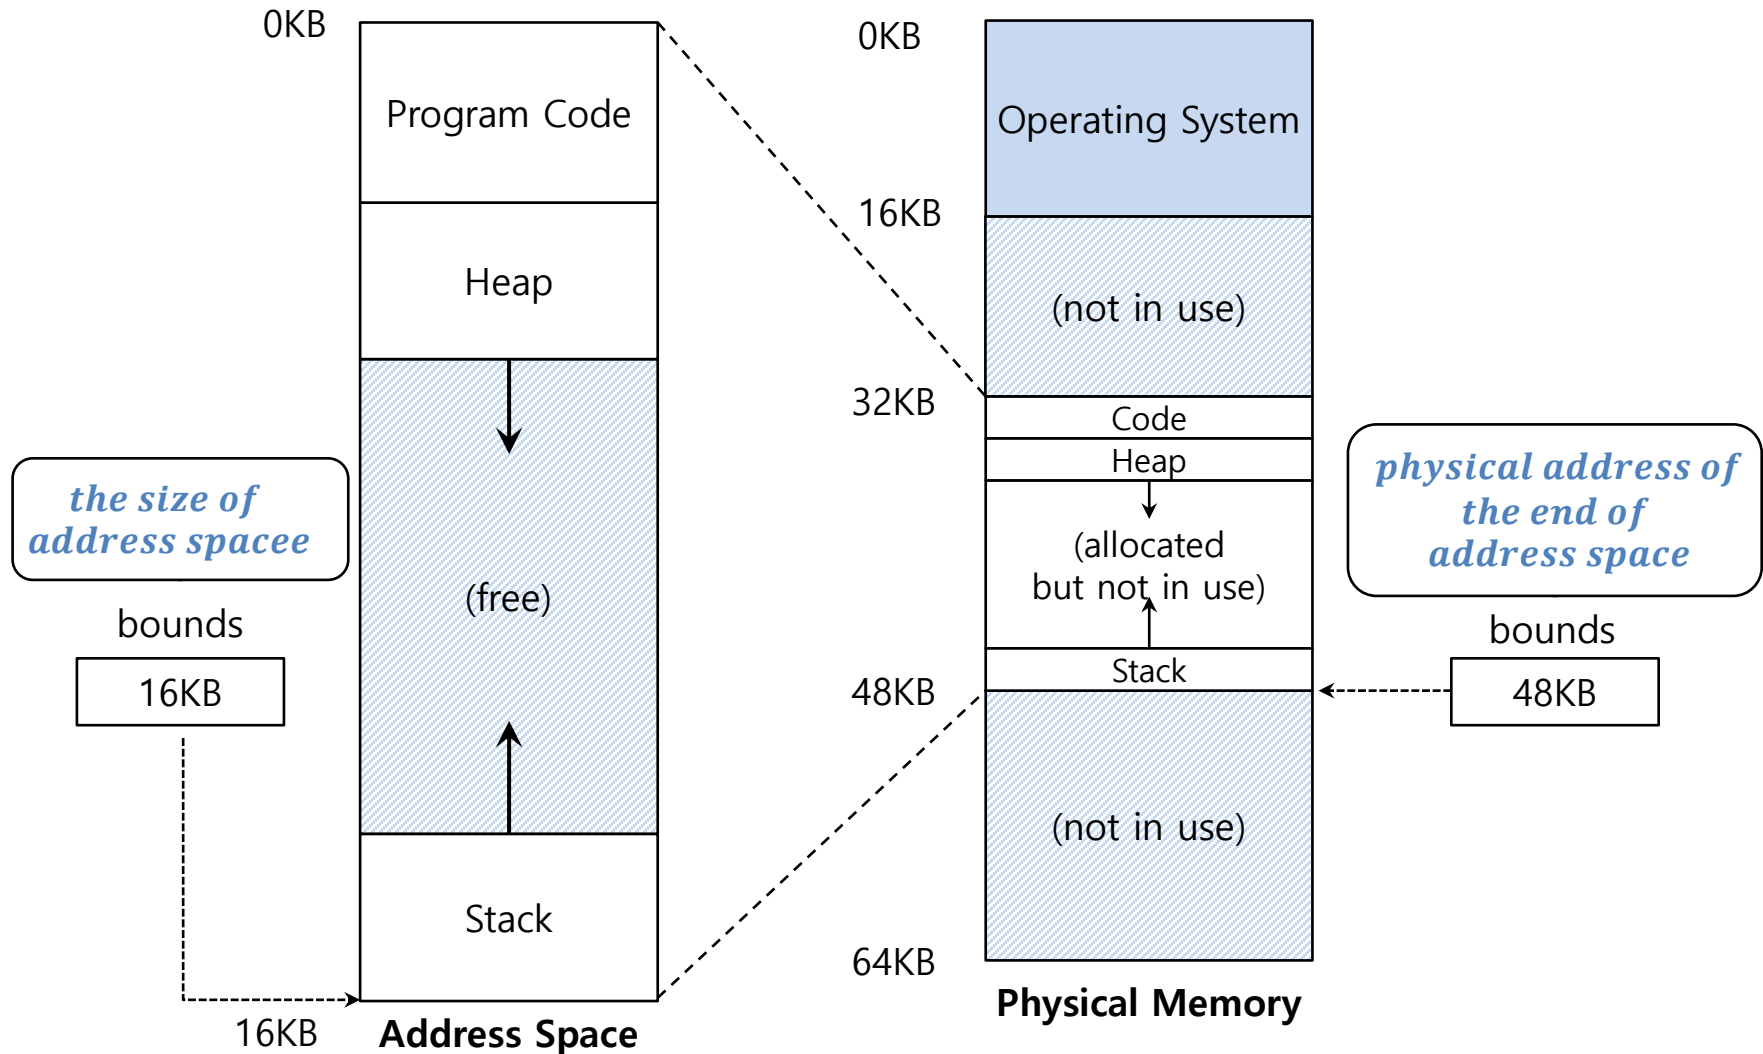
- When a program starts running, the OS decides **where** in physical memory a process should be **loaded**.
 - ◆ Set the **base** register a value.

$$\text{physical address} = \text{virtual address} + \text{base}$$


- ◆ Every virtual address must **not be greater than bound** and **negative**.

$$0 \leq \text{virtual address} < \text{bounds}$$

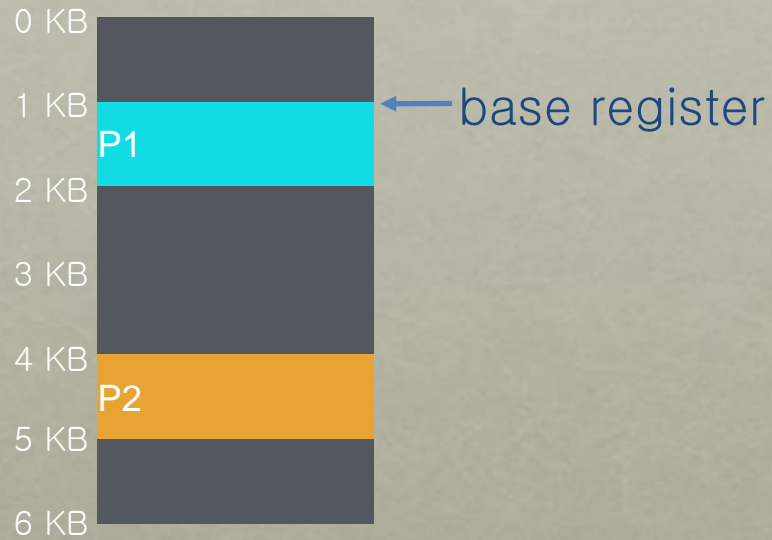
Two ways of Bounds Register



Base Register



Base Register



P1 is running

Base Register




Base Register



(Decimal notation)

Virtual	Physical
P1: load 100, R1	

Base Register



Virtual	Physical	
P1: load 100, R1	load 1124, R1	(1024 + 100)

Base Register



Virtual	Physical
P1: load 100, R1	load 1124, R1
P2: load 100, R1	

Base Register



Virtual	Physical	
P1: load 100, R1	load 1124, R1	
P2: load 100, R1	load 4196, R1	(4096 + 100)

Base Register



Virtual	Physical
P1: load 100, R1	load 1124, R1
P2: load 100, R1	load 4196, R1
P2: load 1000, R1	

Base Register



Virtual	Physical
P1: load 100, R1	load 1124, R1
P2: load 100, R1	load 4196, R1
P2: load 1000, R1	load 5096, R1

Base Register



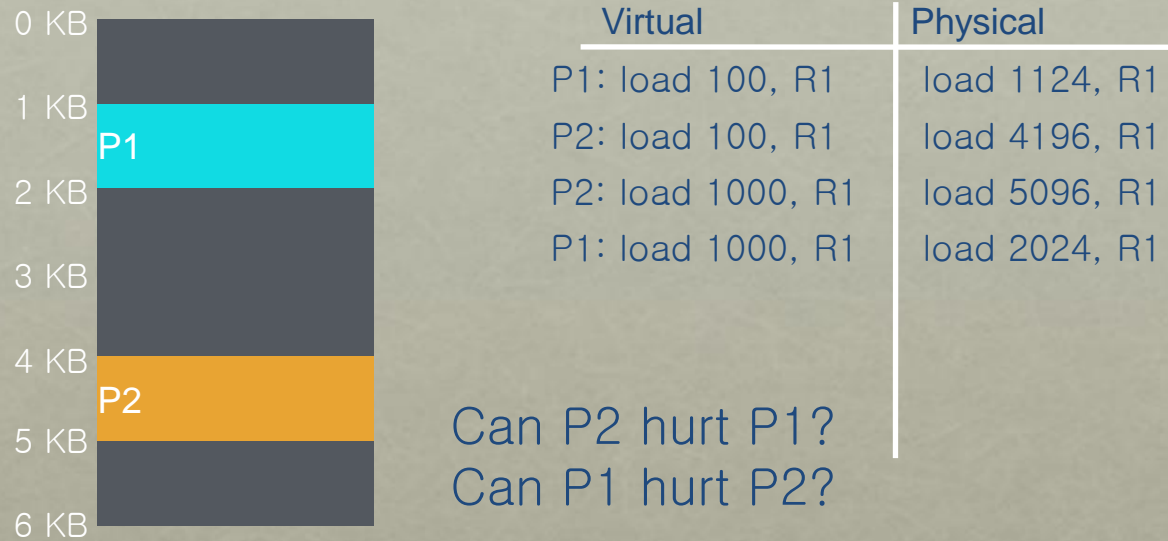
Virtual	Physical
P1: load 100, R1	load 1124, R1
P2: load 100, R1	load 4196, R1
P2: load 1000, R1	load 5096, R1
P1: load 1000, R1	

Base Register




Virtual	Physical
P1: load 100, R1	load 1124, R1
P2: load 100, R1	load 4196, R1
P2: load 1000, R1	load 5096, R1
P1: load 1000, R1	load 2024, R1

Base Register



How to protect process with base register?

Base Register



How to protect process with base register?

Base + Bounds



Base + Bounds



Base + Bounds



Virtual	Physical
P1: load 100, R1	load 1124, R1
P2: load 100, R1	load 4196, R1
P2: load 1000, R1	load 5096, R1
P1: load 1000, R1	load 2024, R1
P1: store 3072, R1	

Can P1 hurt P2?

Base + Bounds




Virtual	Physical
P1: load 100, R1	load 1124, R1
P2: load 100, R1	load 4196, R1
P2: load 1000, R1	load 5096, R1
P1: load 1000, R1	load 2024, R1
P1: store 3072, R1	Exception

3072 > 1024

Can P1 hurt P2?

Base + Bounds



Virtual	Physical
P1: load 100, R1	load 1124, R1
P2: load 100, R1	load 4196, R1
P2: load 1000, R1	load 5096, R1
P1: load 1000, R1	load 2024, R1
P1: store 3072, R1	Exception

3072 > 1024

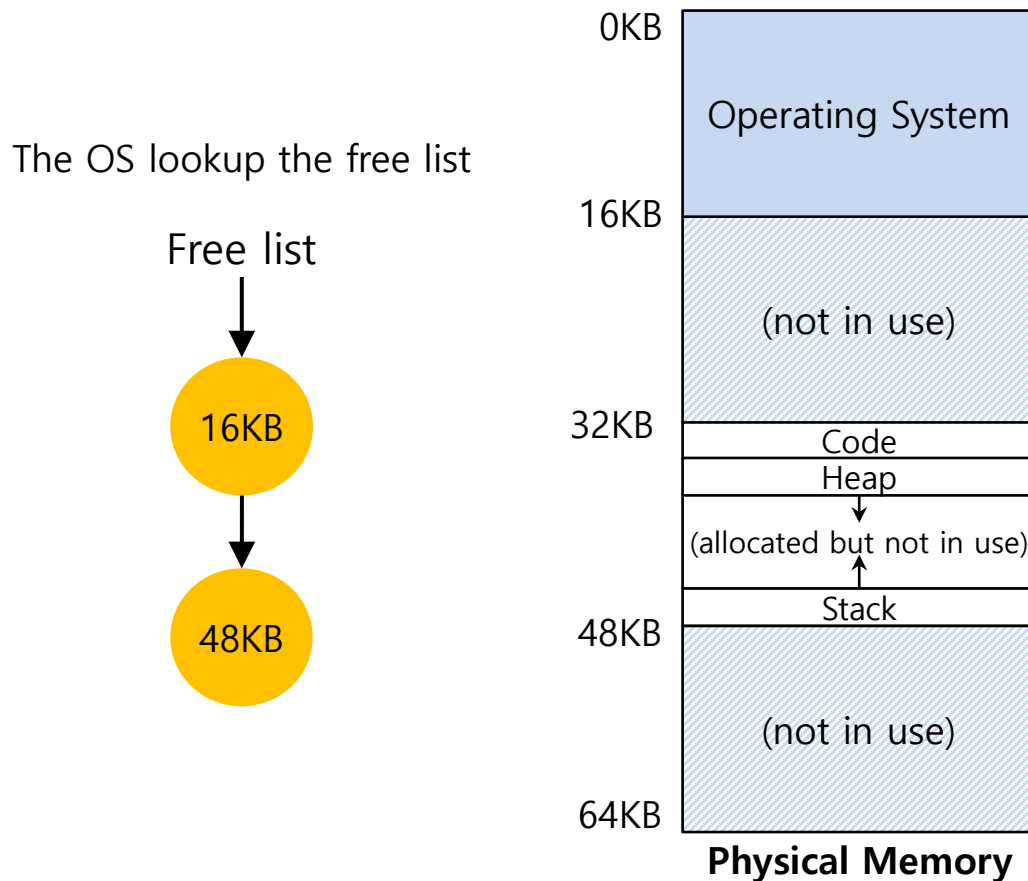
Can P1 hurt P2?

OS Issues for Memory Virtualizing

- ❑ The OS must **take action** to implement **base-and-bounds** approach.
- ❑ Three critical junctures:
 - ◆ When a process **starts running**:
 - Finding space for address space in physical memory
 - ◆ When a process is **terminated**:
 - Reclaiming the memory for use
 - ◆ When context **switch occurs**:
 - Saving and storing the base-and-bounds pair



OS Issues: When a Process Starts Running

- ❑ The OS must **find a room** for a new address space.
 - ◆ free list : A list of the range of the physical memory which are not in use.



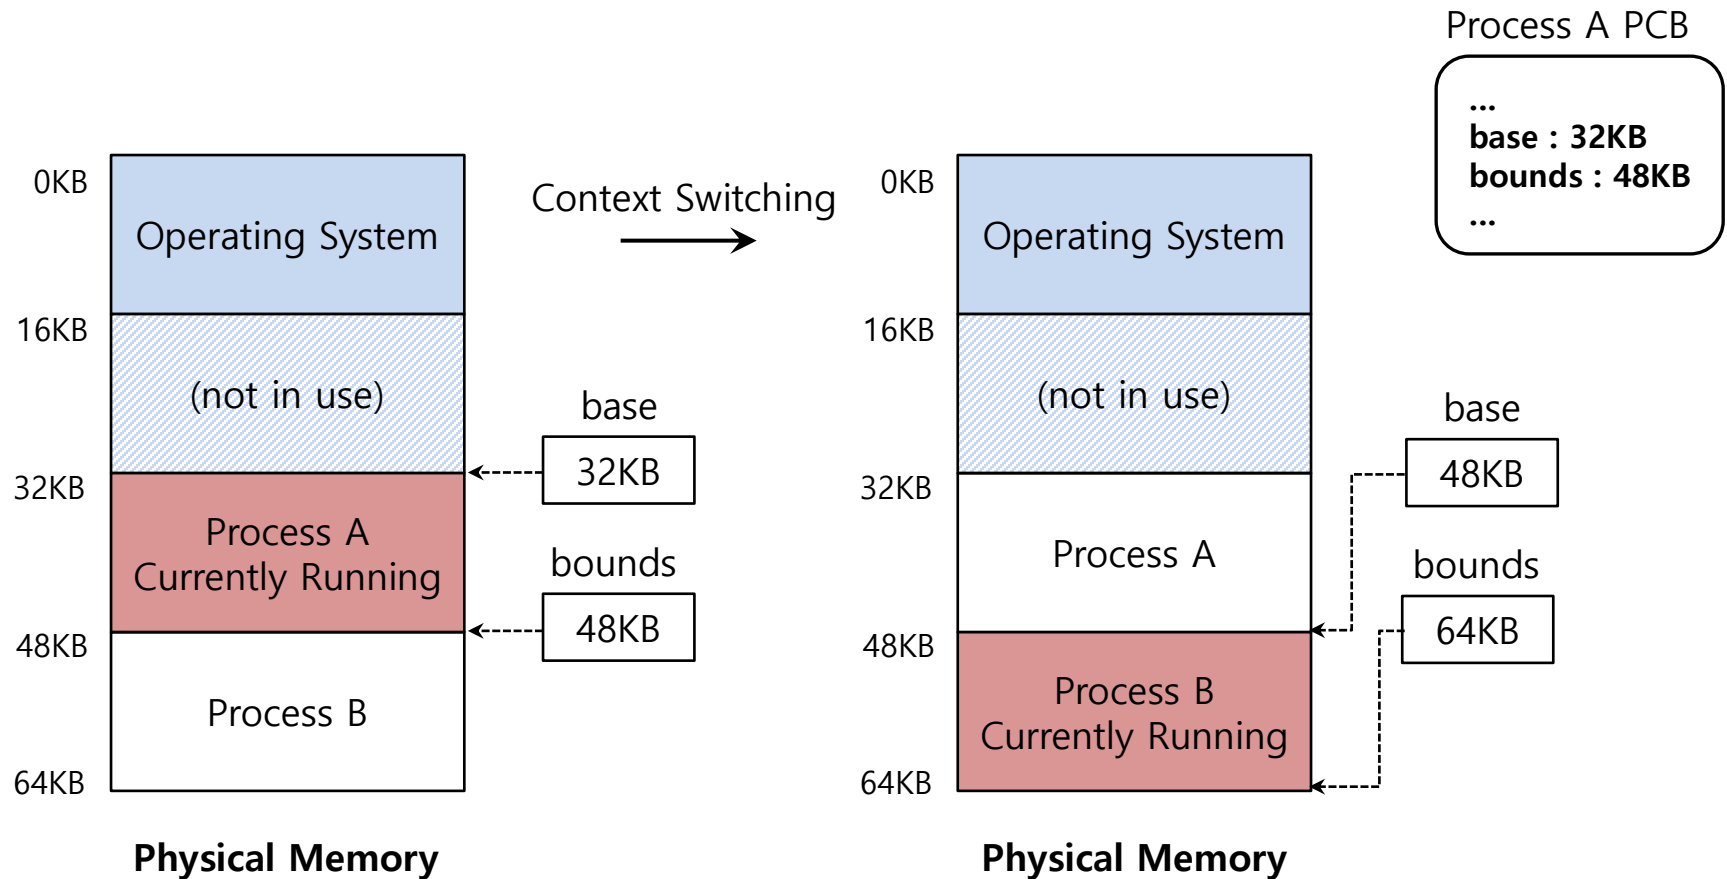
OS Issues: When a Process Is Terminated

- ▣ The OS must **put the memory back** on the free list.



OS Issues: When Context Switch Occurs

- ❑ The OS must **save and restore** the base-and-bounds pair.
 - ◆ In **process structure** or **process control block(PCB)**



Summary

- ❑ Address Space – an **abstraction** of physical memory created by OS
 - ◆ The address space contains all about a running process.
 - ◆ That is consist of program code, data, heap, and stack
- ❑ Address Translation – Hardware-based Address Translation
 - ◆ The hardware transforms each memory access (e.g., an instruction fetch, load, or store), changing the virtual address provided by the instruction to a physical address where the desired information is actually located .
 - ◆ OS performs memory management – keeping track of which locations are free and which are in use, and judiciously intervening to maintain control over how memory is used.
 - ◆ The Base and Bounds Approach
 - When a program starts running, the OS decides where in physical memory a process should be loaded
 - The hardware checks whether virtual addresses are in bounds.
- ❑ Next: Memory Management (Chapters [16](#), [17](#), [18](#), [19](#), [20](#), [21](#), [22](#), [23](#))