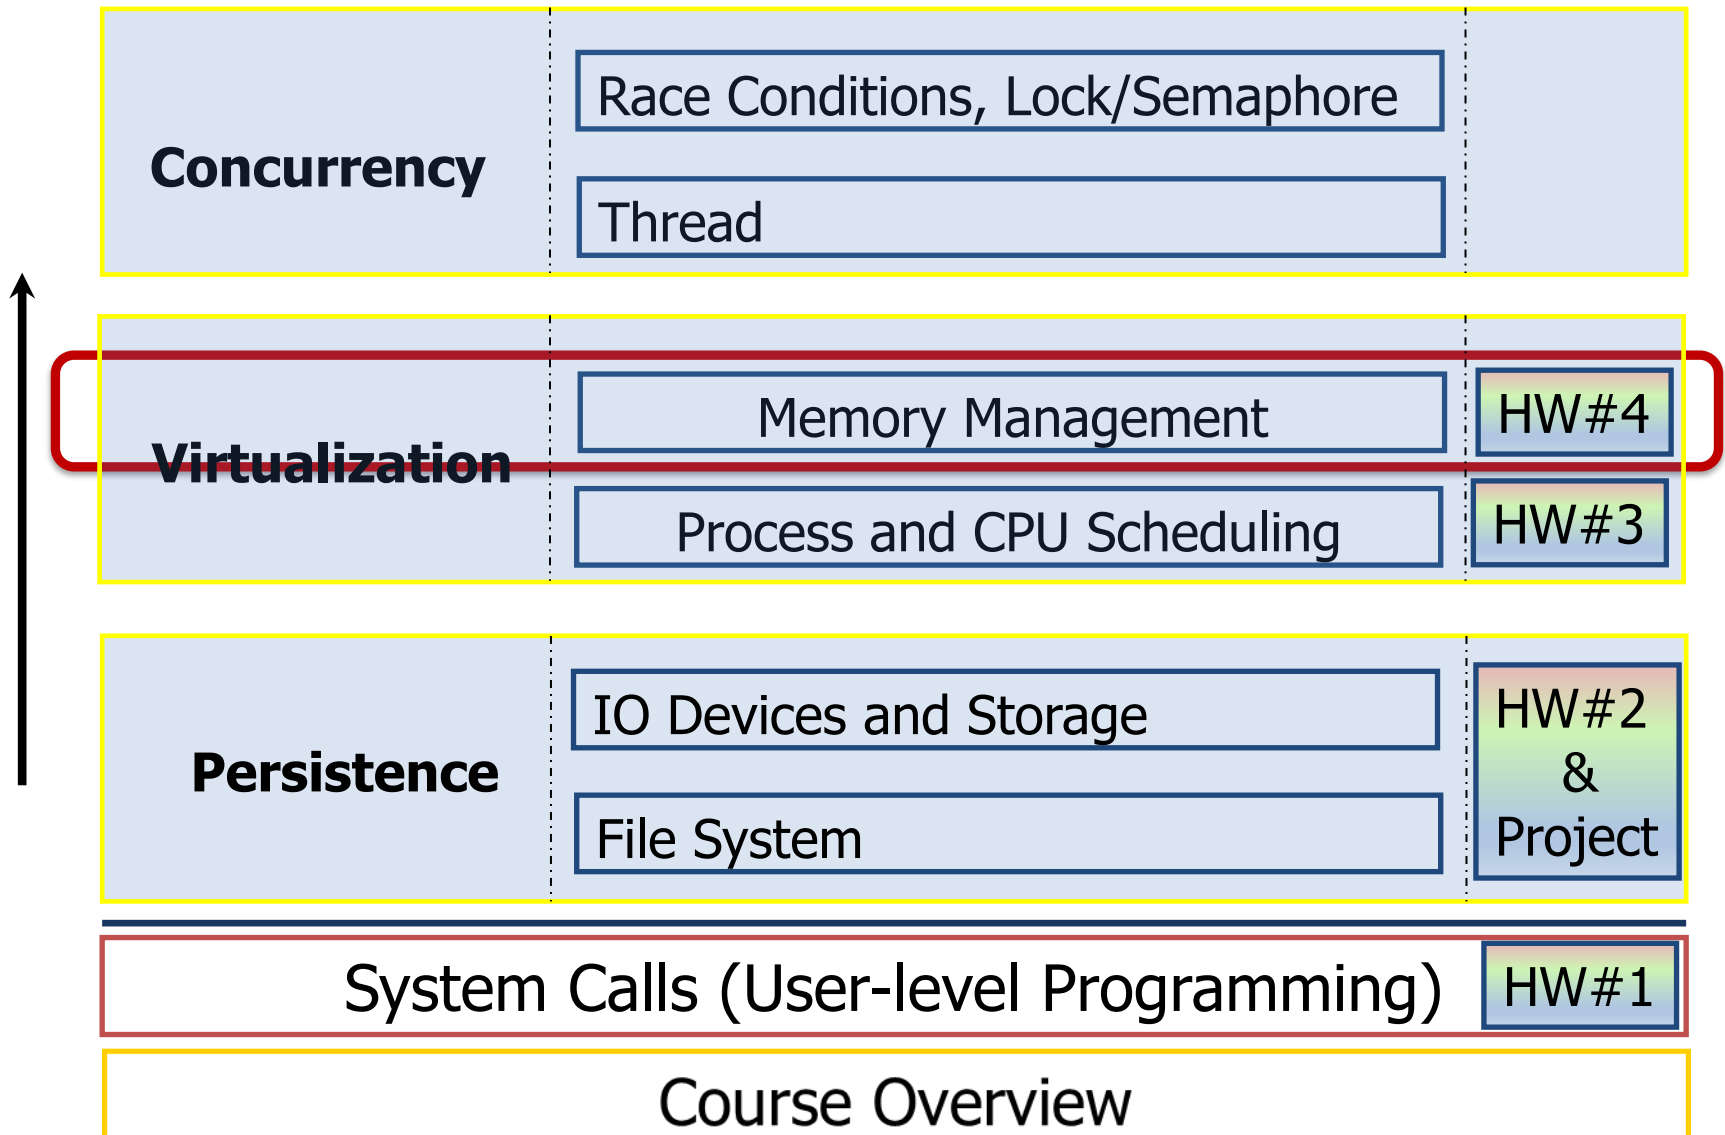


Lecture 10: Virtualizing Memory – Segmentation and Paging

The Course Organization (Bottom-up)

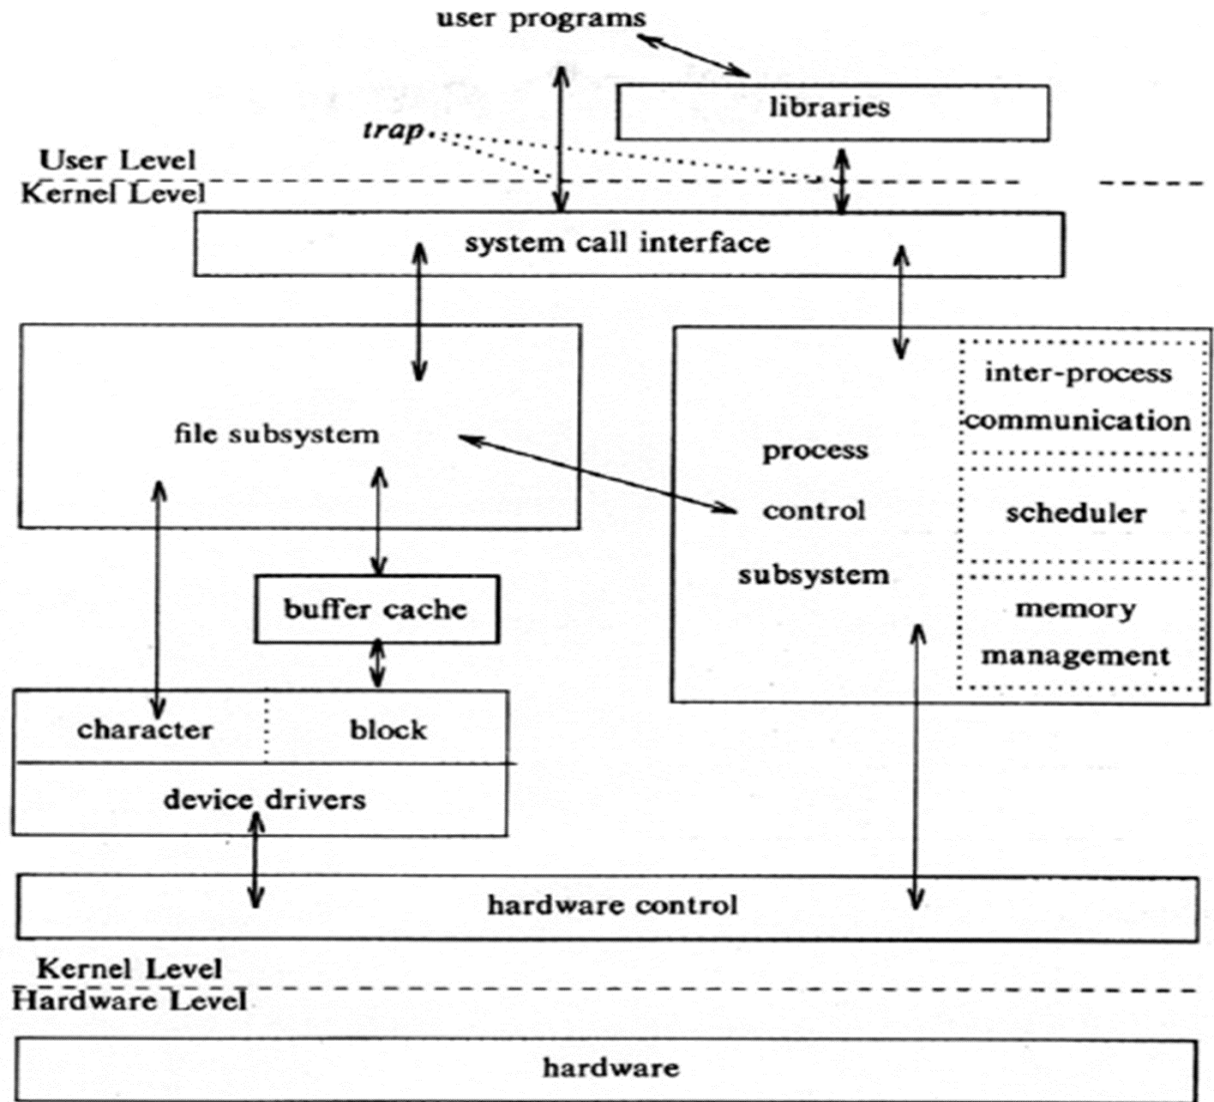


OS – Resource management via virtualization

OS provides services via **System Call** (typically a few hundred) to run **process**, access memory/devices/files, etc.

The OS **manages resources** such as *CPU*, *memory* and *disk* via **virtualization**.


- many programs to run (processes) → Sharing the CPU
- many processes to *concurrently* access their own instructions and data → Sharing memory
- many processes to access devices → Sharing disks



The Design Of The Unix Operating System (Maurice Bach, 1986)

Part I: Segmentation

Inefficiency of the Base and Bound Approach




- ❑ **Big chunk of “free” space**
- ❑ “free” space **takes up** physical memory.
- ❑ Hard to run when an address space **does not fit** into physical memory

Segmentation

- Segment is just a **contiguous portion** of the address space of a particular length.
 - ◆ Logically-different segment: code, stack, heap
- Each segment can be **placed** in **different part of physical memory**.
 - ◆ **Base** and **bounds** exist **per each segment**.

Placing Segment In Physical Memory

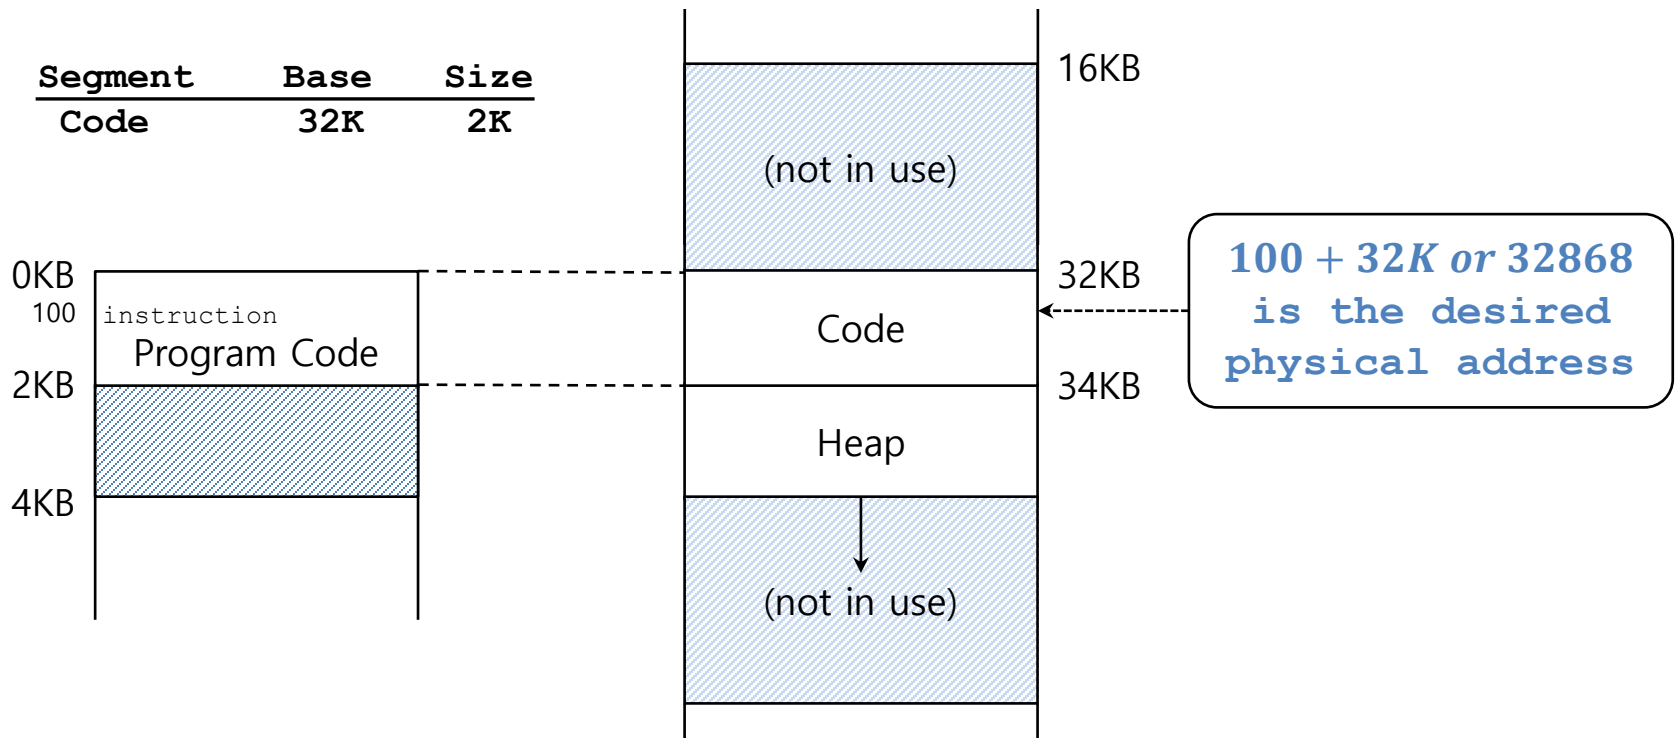


Segment	Base	Size
Code	32K	2K
Heap	34K	2K
Stack	28K	2K

Address Translation on Segmentation

$$\text{physical address} = \text{offset} + \text{base}$$

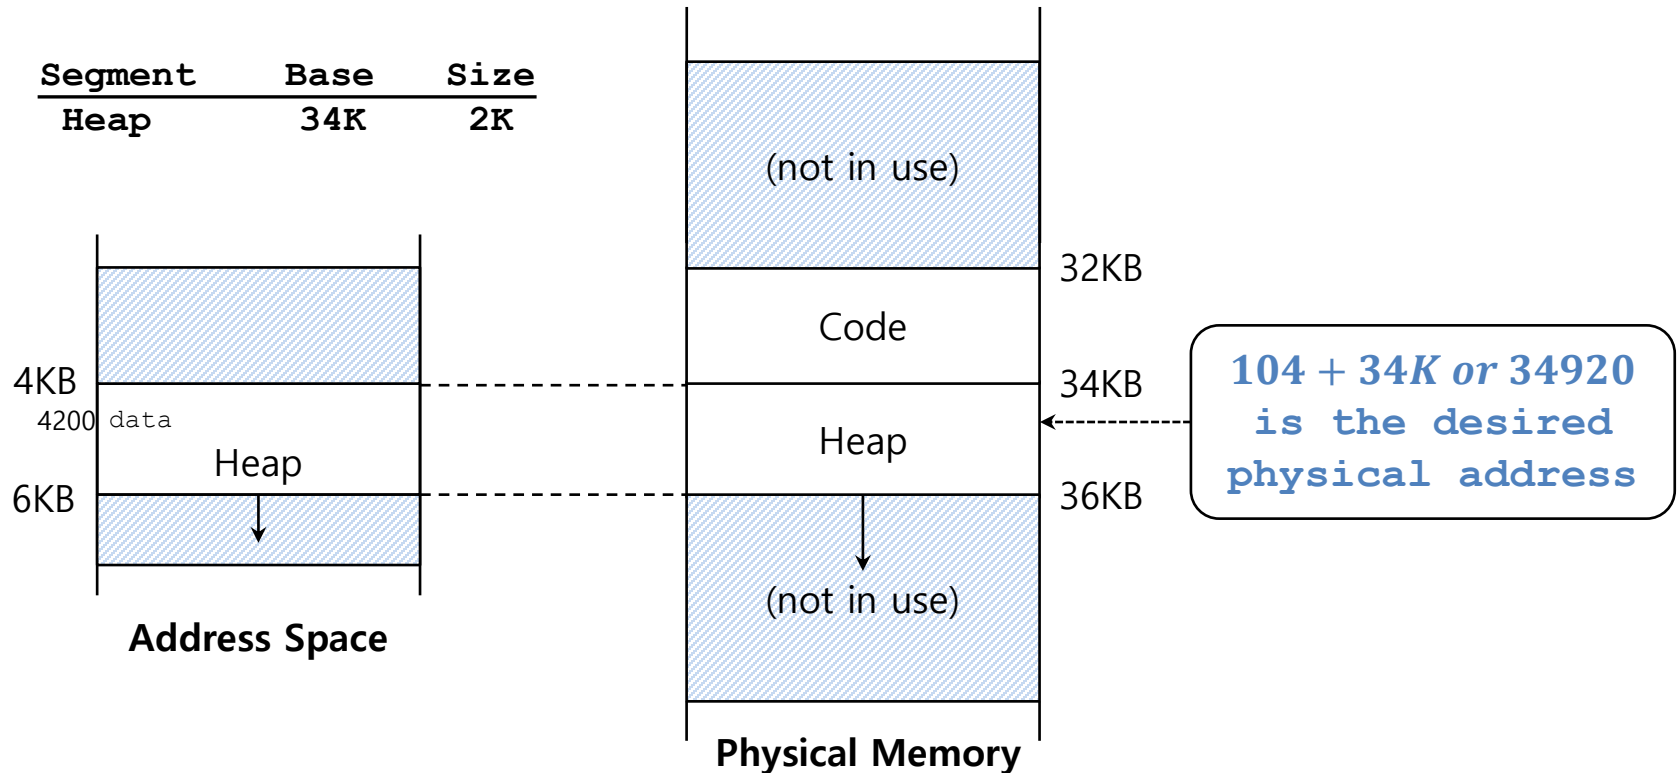
- The offset of virtual address 100 is 100.
 - ◆ The code segment **starts at virtual address 0** in address space.



Address Translation on Segmentation(Cont.)


Virtual address + base is not the correct physical address.

- The offset of virtual address 4200 is 104.
- ◆ The heap segment **starts at virtual address 4096** in address space.



Segmentation Fault or Violation


- ❑ If an **illegal address** such as 7KB which is beyond the end of heap is referenced, the OS occurs **segmentation fault**.
 - ◆ The hardware detects that the address is **out of bounds**.



Referring to Segment


□ Explicit approach

- ◆ Chop up the address space into segments based on the **top few bits** of virtual address.



- Example: virtual address 4200 (010000001101000)

Segment	bits
Code	00
Heap	01
Stack	10
-	11




Referring to Segment(Cont.)

```
1  // get top 2 bits of 14-bit VA
2  Segment = (VirtualAddress & SEG_MASK) >> SEG_SHIFT
3  // now get offset
4  Offset = VirtualAddress & OFFSET_MASK
5  if (Offset >= Bounds[Segment])
6      RaiseException(PROTECTION_FAULT)
7  else
8      PhysAddr = Base[Segment] + Offset
9      Register = AccessMemory(PhysAddr)
```

- ◆ `SEG_MASK = 0x3000 (11000000000000)`
- ◆ `SEG_SHIFT = 12`
- ◆ `OFFSET_MASK = 0xFFF (00111111111111)`

Referring to Stack Segment

- ❑ Stack grows **backward**.
- ❑ **Extra hardware support** is need.
 - ◆ The hardware checks which way the segment grows.
 - ◆ 1: positive direction, 0: negative direction



Segment Register(with Negative-Growth Support)

Segment	Base	Size	Grows Positive?
Code	32K	2K	1
Heap	34K	2K	1
Stack	28K	2K	0

Support for Sharing

- Segment can be **shared between address** space.
 - ◆ **Code sharing** is still in use in systems today.
 - ◆ by extra hardware support.
- Extra hardware support is need for form of **Protection bits**.
 - ◆ **A few more bits** per segment to indicate **permissions** of **read**, write and **execute**.

Segment Register Values(with Protection)

Segment	Base	Size	Grows Positive?	Protection
Code	32K	2K	1	Read-Execute
Heap	34K	2K	1	Read-Write
Stack	28K	2K	0	Read-Write

Fine-Grained and Coarse-Grained



- ❑ **Coarse-Grained** means segmentation in a small number.
 - ◆ e.g., code, heap, stack.
- ❑ **Fine-Grained** segmentation allows **more flexibility** for address space in some early system.
 - ◆ To support many segments, Hardware support with a **segment table** is required.

OS support: Fragmentation

- ❑ **External Fragmentation:** little holes of **free space** in physical memory that make difficulty to allocate new segments.
 - ◆ There is **X KB** free, but **not in one contiguous** segment.
 - ◆ The OS **cannot** satisfy the **X KB** request.

- ❑ **Compaction:** **rearranging** the exiting segments in physical memory.
 - ◆ Compaction is **costly**.
 - **Stop** running process.
 - **Copy** data to somewhere.
 - **Change** segment register value.

Memory Compaction




Part II: Paging


Concept of Paging

- Paging **splits up** address space into **fixed-sized** unit called a **page**.
 - ◆ Segmentation: variable size of logical segments(code, stack, heap, etc.)
- With paging, **physical memory** is also **split** into some number of pages called a **page frame**.
- **Page table** per process is needed **to translate** the virtual address to physical address.

Paging



Fill in Page Table



Page Tables:

P1	P2	P3
3	0	8
1	4	5
7	2	9
10	6	11


Advantages Of Paging

- ❑ **Flexibility:** Supporting the abstraction of address space effectively
 - ◆ Don't need assumption how heap and stack grow and are used.

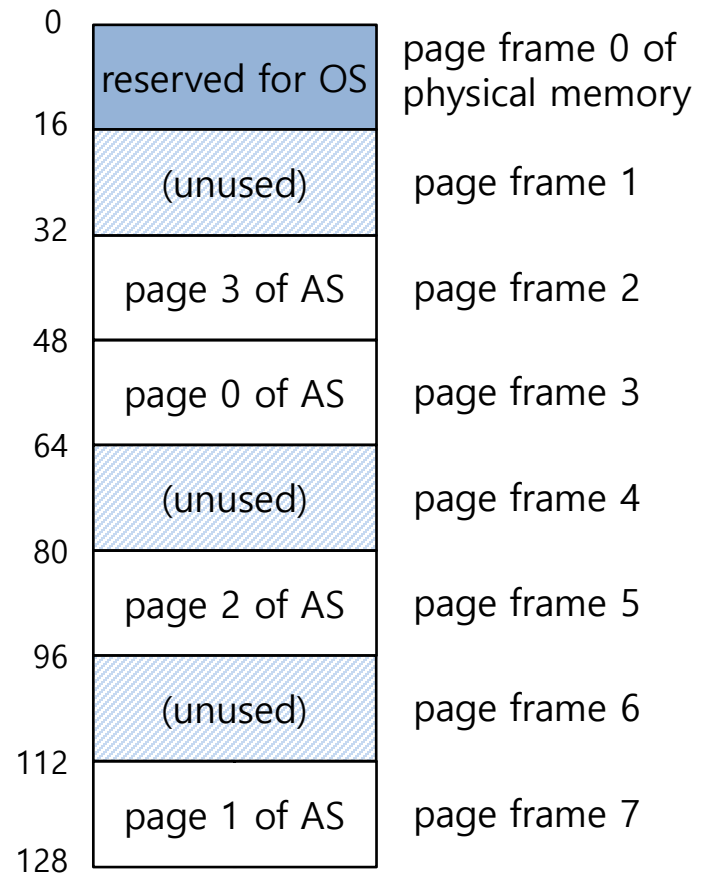
- ❑ **Simplicity:** ease of free-space management
 - ◆ The page in address space and the page frame are the same size.
 - ◆ Easy to allocate and keep a free list

Example: A Simple Paging

- 128-byte physical memory with 16 bytes page frames
- 64-byte address space with 16 bytes pages




A Simple 64-byte Address Space




64-Byte Address Space Placed In Physical Memory

Address Translation

- Two components in the virtual address
 - VPN: virtual page number
 - Offset: offset within the page




- Example: virtual address 21 in 64-byte address space



Example: Address Translation

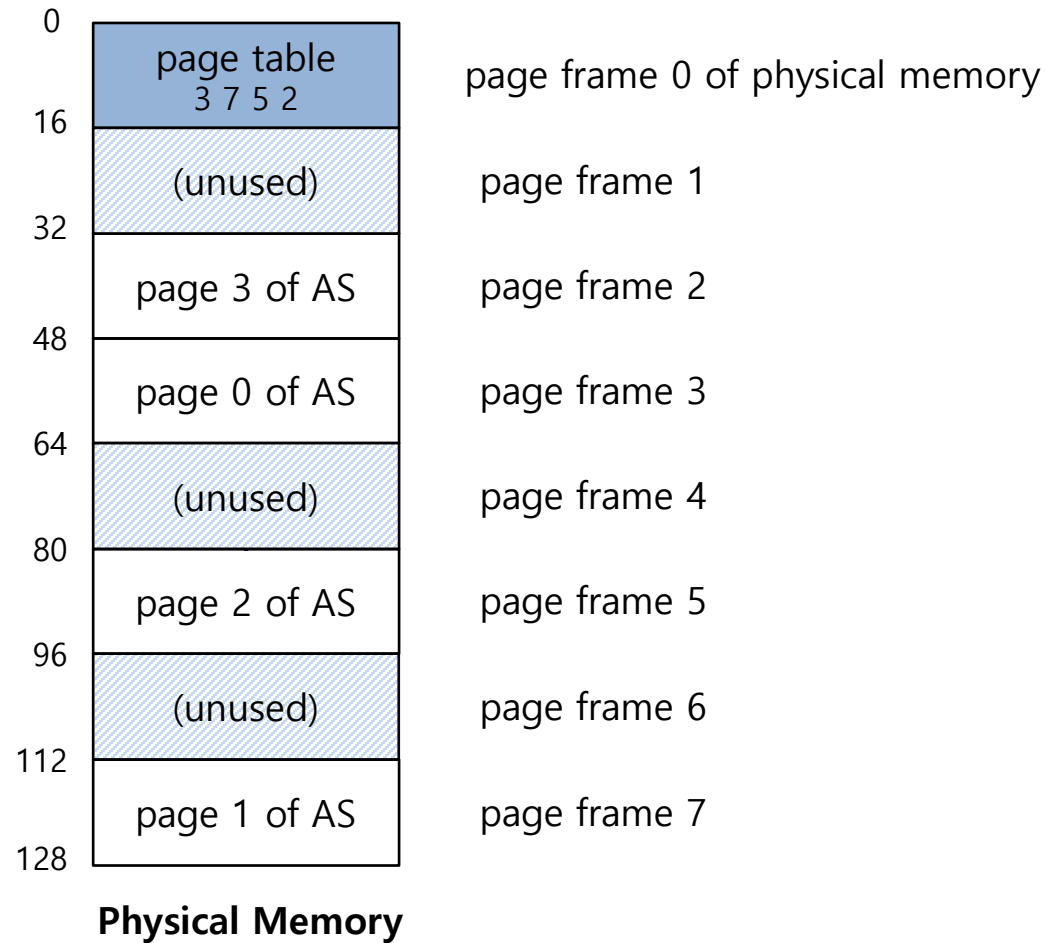
- ▣ The virtual address 21 in 64-byte address space



Where Are Page Tables Stored?

- ❑ Page tables can get awfully large
 - ◆ 32-bit address space with 4-KB pages, 20 bits for VPN
 - $4MB = 2^{20} \text{ entries} * 4 \text{ Bytes per page table entry}$
- ❑ Page tables for peach process are stored in memory.

Example: Page Table in Kernel Physical Memory




What is in the page table?

- ❑ The page table is just a **data structure** that is used to map the virtual address to physical address.
 - ◆ Simplest form: a linear page table, an array
- ❑ The OS **indexes** the array by VPN, and looks up the page-table entry.

Common Flags Of Page Table Entry

- ❑ **Valid Bit:** Indicating whether the particular translation is valid.
- ❑ **Protection Bit:** Indicating whether the page could be read from, written to, or executed from
- ❑ **Present Bit:** Indicating whether this page is in physical memory or on disk(swapped out)
- ❑ **Dirty Bit:** Indicating whether the page has been modified since it was brought into memory
- ❑ **Reference Bit(Accessed Bit):** Indicating that a page has been accessed

Example: x86 Page Table Entry



An x86 Page Table Entry(PTE)

- ❑ P: present
- ❑ R/W: read/write bit
- ❑ U/S: supervisor
- ❑ A: accessed bit
- ❑ D: dirty bit
- ❑ PFN: the page frame number

Paging: Too Slow

- ❑ To find a location of the desired PTE, the **starting location** of the page table is **needed**.
- ❑ For every memory reference, paging requires the OS to perform one **extra memory reference**.

Accessing Memory With Paging

```
1      // Extract the VPN from the virtual address
2      VPN = (VirtualAddress & VPN_MASK) >> SHIFT
3
4      // Form the address of the page-table entry (PTE)
5      PTEAddr = PTBR + (VPN * sizeof(PTE))
6
7      // Fetch the PTE
8      PTE = AccessMemory(PTEAddr)
9
10     // Check if process can access the page
11     if (PTE.Valid == False)
12         RaiseException(SEGMENTATION_FAULT)
13     else if (CanAccess(PTE.ProtectBits) == False)
14         RaiseException(PROTECTION_FAULT)
15     else
16         // Access is OK: form physical address and fetch it
17         offset = VirtualAddress & OFFSET_MASK
18         PhysAddr = (PTE.PFN << PFN_SHIFT) | offset
19         Register = AccessMemory(PhysAddr)
```


Summary

❑ Segmentation

- ◆ A program contains logically-different **segments**: code, stack, heap
- ◆ Each segment can be **placed** in **different part of physical memory**.
 - **Base** and **bounds** exist **per each segment**.

❑ Paging

- ◆ **physical memory** is **split** into some number of pages called a **page frame**.
- ◆ **Page table** per process is used **to translate** the virtual address to physical address.

❑ Next: Free Space Management, Translation Lookaside Buffers, Advanced Page Tables, and Swapping ([Chapters 17](#), [19](#), [20](#), [21](#), [22](#), [23](#))