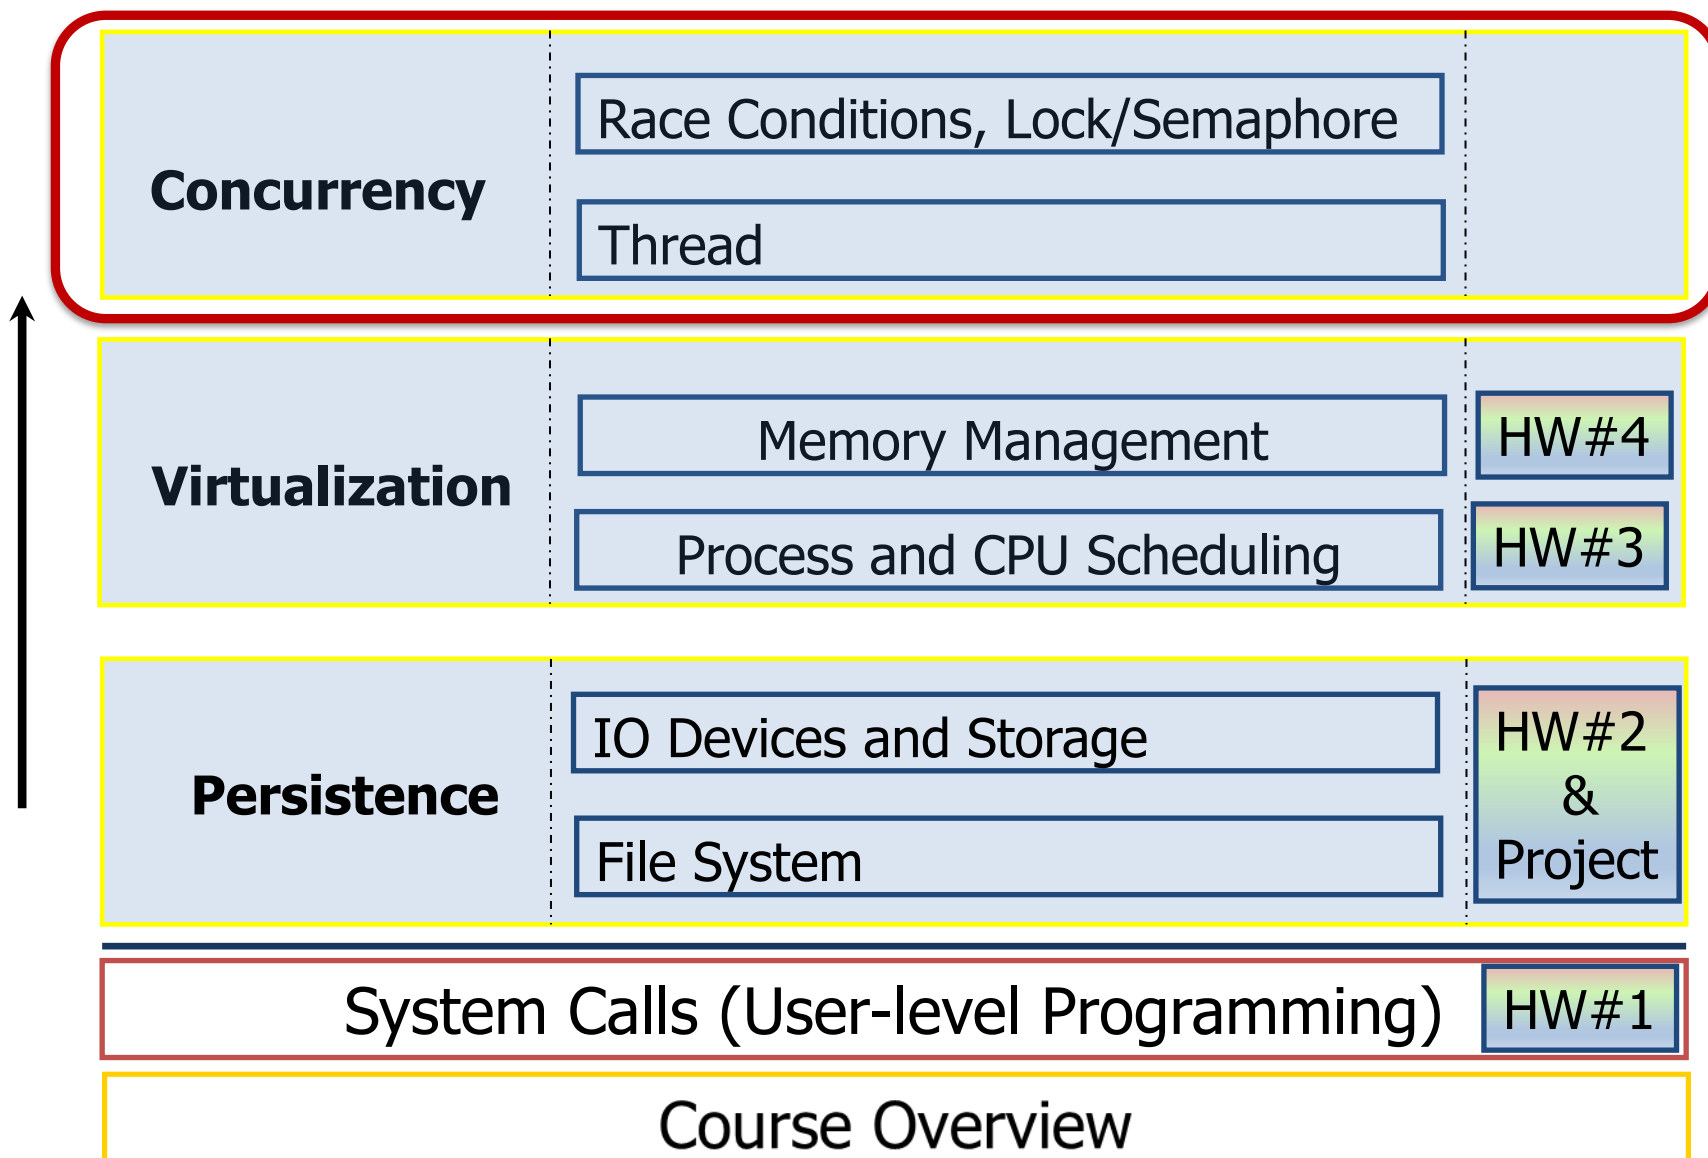


# **Lecture 13: Concurrency – Thread and Thread API**

---

# The Course Organization (Bottom-up)



# **Part I: Introduction to Thread**

---

# Thread

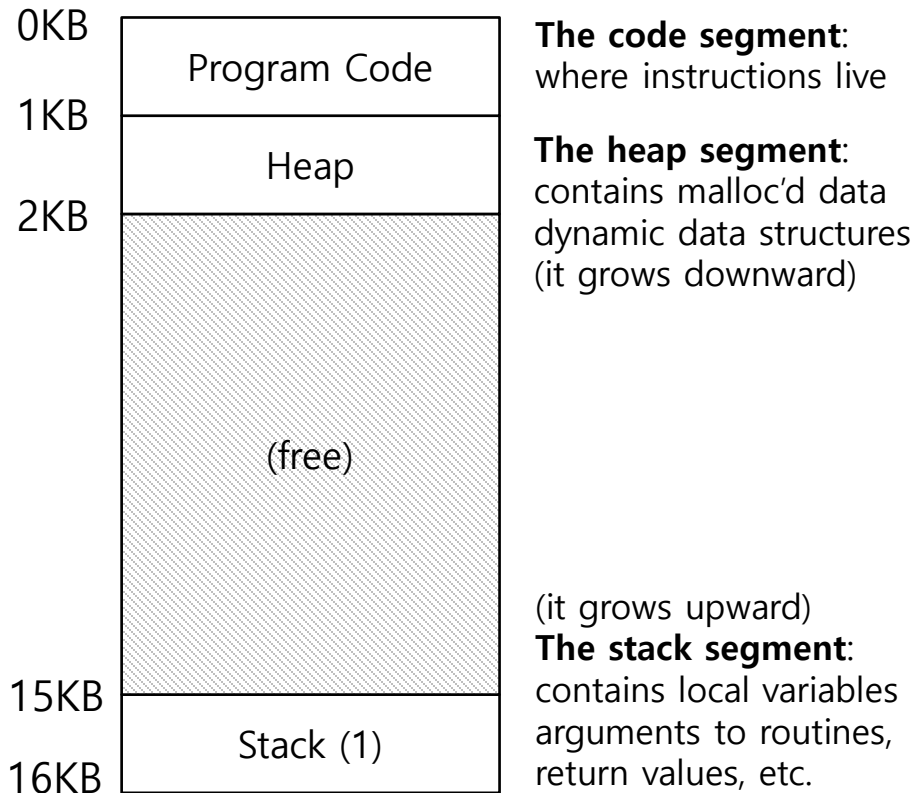
- ❑ A new abstraction for a single running process
- ❑ Multi-threaded program
  - ◆ A multi-threaded program has more than one point of execution.
  - ◆ Multiple PCs (Program Counter)
  - ◆ They **share** the same **address space**.

# Context switch between threads


- Each thread has its own program counter and set of registers.
  - ◆ One or more **thread control blocks(TCBs)** are needed to store the state of each thread.
  
- When switching from running one (T1) to running the other (T2),
  - ◆ The register state of T1 be saved.
  - ◆ The register state of T2 restored.
  - ◆ The **address space remains** the same.

# The stack of the relevant thread

- There will be **one stack per thread**.




**A Single-Threaded  
Address Space**




**Two threaded  
Address Space**

# Concurrency Problem

- ❑ The OS is juggling **many things at once**, first running one process, then another, and so forth.
- ❑ Modern **multi-threaded programs** also exhibit the concurrency problem.



single-threaded process



multithreaded process

# Concurrency Example

## ▣ A Multi-threaded Program (thread.c)

```
1      #include <stdio.h>
2      #include <stdlib.h>
3      #include <pthread.h>
4
5      volatile int counter = 0;
6      int loops;
7
8      void *worker(void *arg) {
9          int i;
10         for (i = 0; i < loops; i++) {
11             counter++;
12         }
13         return NULL;
14     }
```



# Concurrency Example (Cont.)

```
16     int
17     main(int argc, char *argv[])
18     {
19         if (argc != 2) {
20             fprintf(stderr, "usage: threads <value>\n");
21             exit(1);
22         }
23         loops = atoi(argv[1]);
24         pthread_t p1, p2;
25         printf("Initial value : %d\n", counter);
26
27         pthread_create(&p1, NULL, worker, NULL);
28         pthread_create(&p2, NULL, worker, NULL);
29         pthread_join(p1, NULL);
30         pthread_join(p2, NULL);
31         printf("Final value : %d\n", counter);
32         return 0;
33     }
```

- ◆ The main program creates **two threads**.
  - Thread: a function running within the same memory space. Each thread start running in a routine called `worker()`.
  - `worker()`: increments a counter

## Concurrency Example (Cont.)

- `loops` determines how many times each of the two workers will **increment the shared counter** in a loop.

- ◆ `loops: 1000.`


```
prompt> gcc -o thread thread.c -Wall -pthread
prompt> ./thread 1000
Initial value : 0
Final value : 2000
```

- ◆ `loops: 1000000000`

```
prompt> ./thread 1000000000
Initial value : 0
Final value : 1997974414 // huh??
prompt> ./thread 1000000000
Initial value : 0
Final value : 1997940107 // what the??
```

# Why is this happening?

- ▣ Increment a shared counter → take three instructions.
  1. Load the value of the counter from the memory into a register.
  2. Increment it
  3. Store it back into the memory
  
- ▣ These three instructions do not execute **atomically**. → Problem of **concurrency** happen.



# What happened?

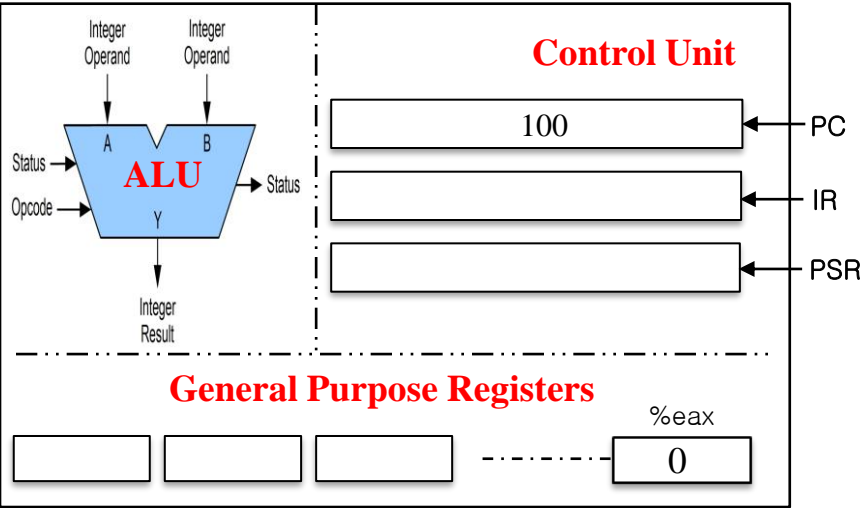
## □ Example with two threads

- ♦  $\text{counter} = \text{counter} + 1$  (**initial value: 50**)
- ♦ We expect the result is **52**. However,

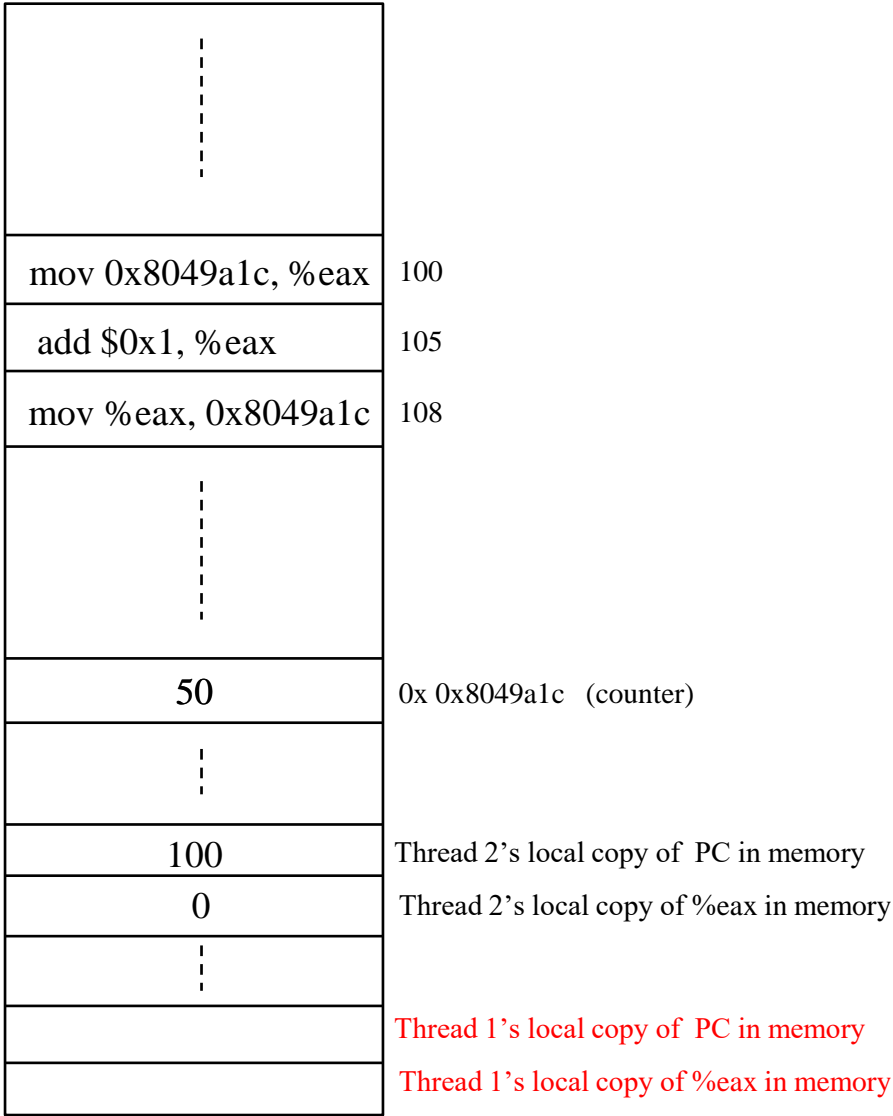
OS	Thread1	Thread2	(after instruction)		
			PC	%eax	counter
	<b>Initial value</b>		100	0	50
	mov 0x8049a1c, %eax		105	50	50
	add \$0x1, %eax		108	51	50
interrupt	save T1's state				
	restore T2's state		100	0	50
		mov 0x8049a1c, %eax	105	50	50
		add \$0x1, %eax	108	51	50
		mov %eax, 0x8049a1c	113	51	51
interrupt	save T2's state				
	restore T1's state		108	51	51
	mov %eax, 0x8049a1c		113	51	<b>51</b>

# Step 0: Initial Value (PC=100)

## CPU



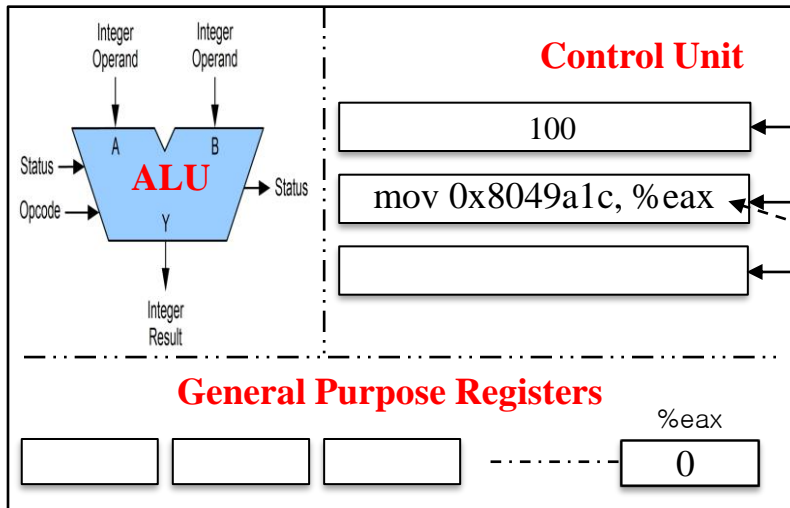
## Memory




OS	Thread1	Thread2	(after instruction)		
			PC	%eax	counter
	<b>Initial value</b>		100	0	50
	mov 0x8049a1c, %eax		105	50	50
	add \$0x1, %eax		108	51	50
interrupt					
save T1's state					
restore T2's state			100	0	50
	mov 0x8049a1c, %eax		105	50	50
	add \$0x1, %eax		108	51	50
	mov %eax, 0x8049a1c		113	51	51
interrupt					
save T2's state					
restore T1's state			108	51	51
	mov %eax, 0x8049a1c		113	51	51

# Step 1: Fetch instruction from the memory (PC->Address=100)

## CPU



## Memory



100

105

108

0x 0x8049a1c (counter)

Thread 2's local copy of PC in memory

Thread 2's local copy of %eax in memory

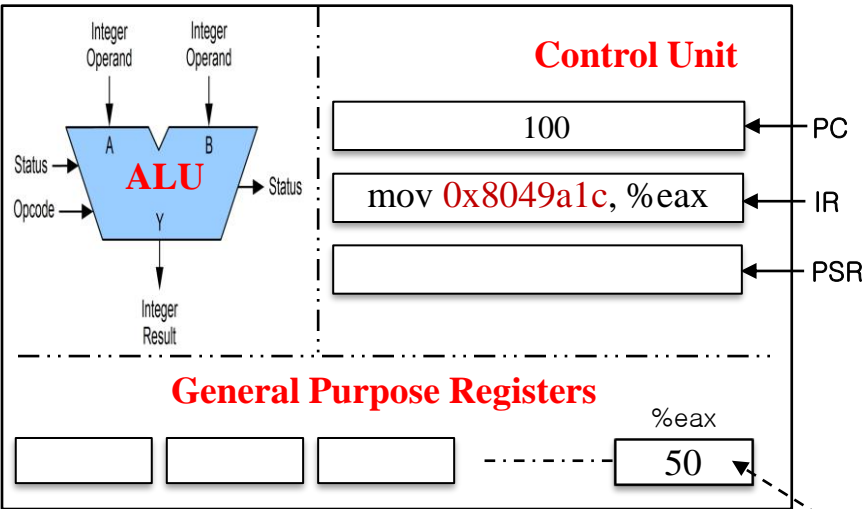
Thread 1's local copy of PC in memory

Thread 1's local copy of %eax in memory


OS	Thread1	Thread2	(after instruction)		
			PC	%eax	counter
	<b>Initial value</b>		100	0	50
	mov 0x8049a1c, %eax		105	50	50
	add \$0x1, %eax		108	51	50
interrupt					
save T1's state					
restore T2's state			100	0	50
	mov 0x8049a1c, %eax		105	50	50
	add \$0x1, %eax		108	51	50
	mov %eax, 0x8049a1c		113	51	51
interrupt					
save T2's state					
restore T1's state			108	51	51
	mov %eax, 0x8049a1c		113	51	51

# Steps 2 & 3: Decode & Execution

## CPU



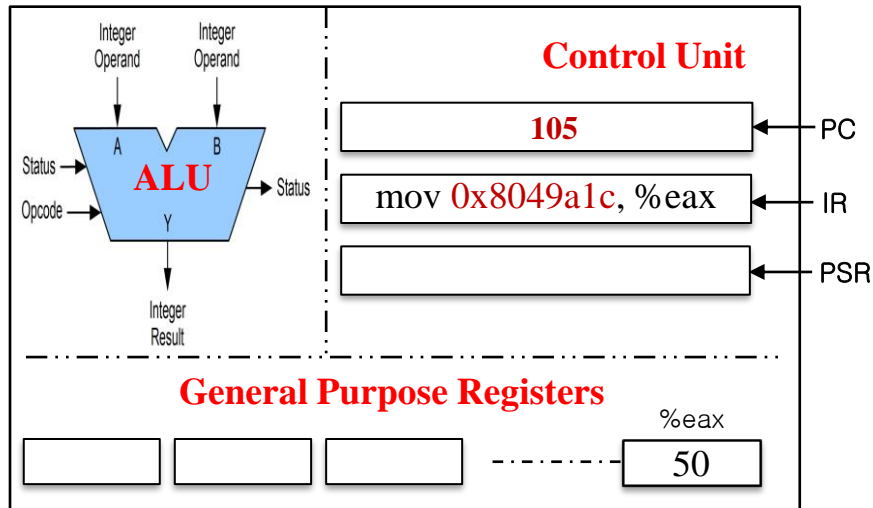
## Memory



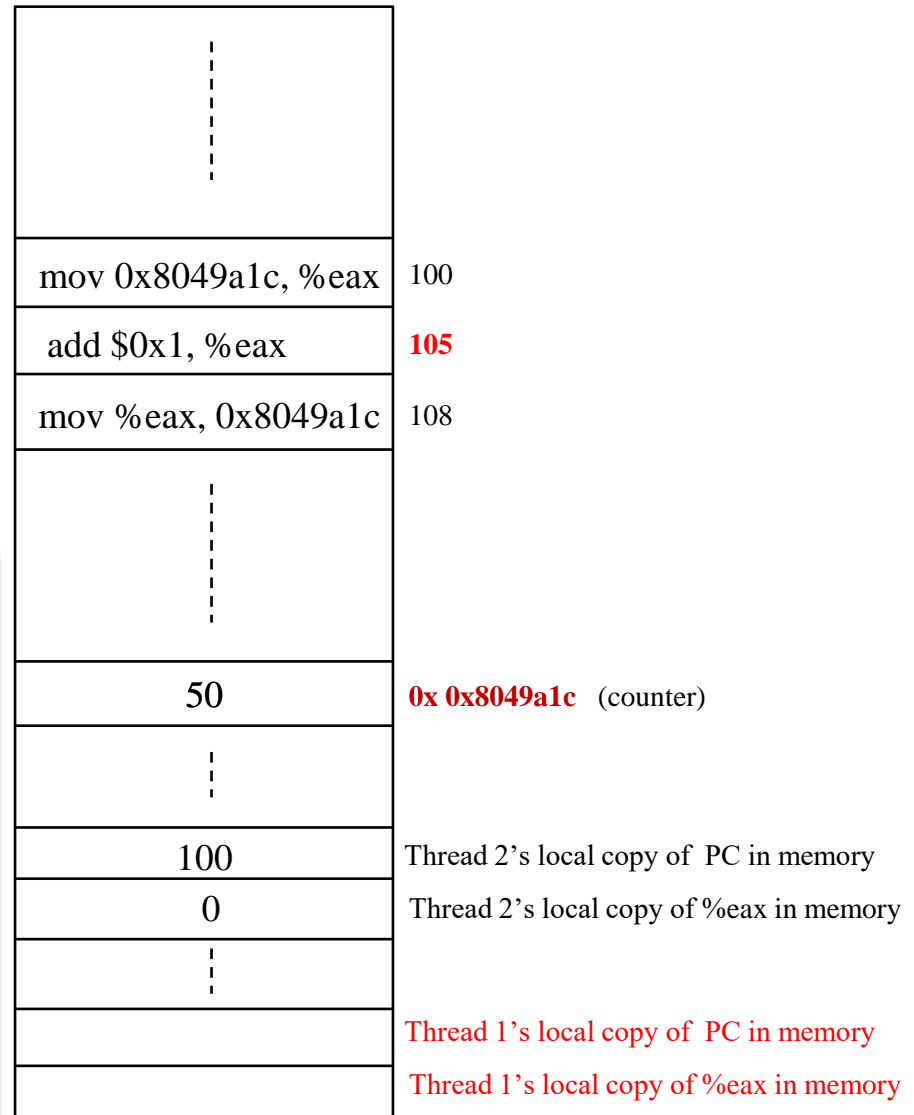
OS	Thread1	Thread2	(after instruction)		
			PC	% <u>eax</u>	counter
	<b>Initial value</b>		100	0	50
	mov 0x8049a1c, % <u>eax</u>		105	50	50
	add \$0x1, % <u>eax</u>		108	51	50
interrupt					
save T1's state					
restore T2's state			100	0	50
		<u>mov</u> 0x8049a1c, % <u>eax</u>	105	50	50
		<u>add</u> \$0x1, % <u>eax</u>	108	51	50
		<u>mov</u> % <u>eax</u> , 0x8049a1c	113	51	51
interrupt					
save T2's state					
restore T1's state			108	51	51
	<u>mov</u> % <u>eax</u> , 0x8049a1c		113	51	51

# Step 4: Increase PC (pointed to the next instruction in the memory)

## CPU



## Memory

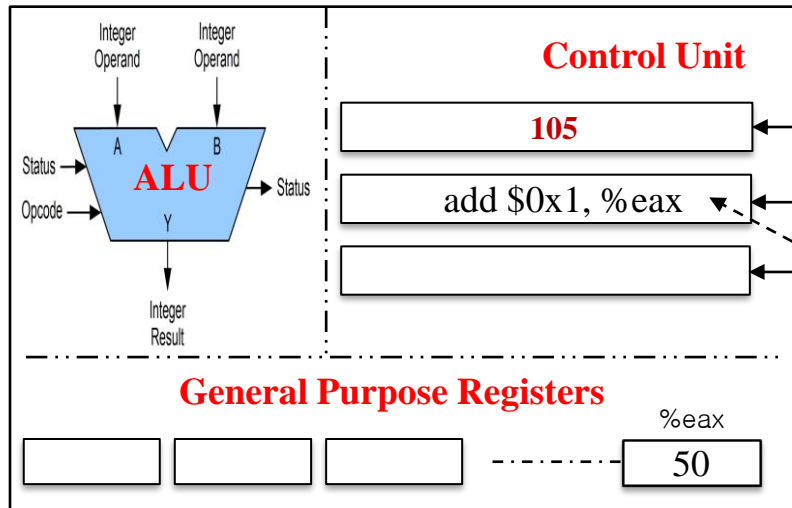


OS	Thread1	Thread2	(after instruction)		
			PC	%eax	counter
	Initial value		100	0	50
	mov 0x8049a1c, %eax		105	50	50
	add \$0x1, %eax		108	51	50
interrupt	save T1's state				
	restore T2's state		100	0	50
	mov 0x8049a1c, %eax		105	50	50
	add \$0x1, %eax		108	51	50
	mov %eax, 0x8049a1c		113	51	51
interrupt	save T2's state				
	restore T1's state		108	51	51
	mov %eax, 0x8049a1c		113	51	51




# Step 1: Fetch instruction from the memory (PC->Address=105)

## CPU

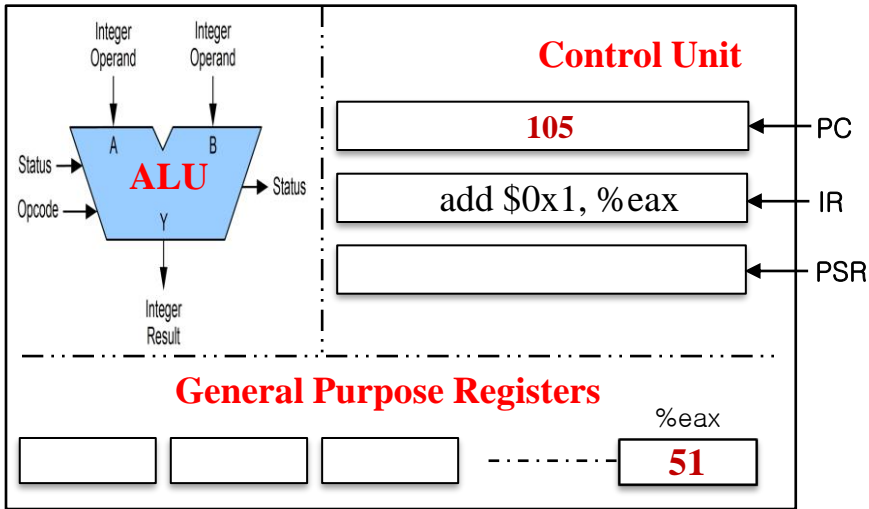


## Memory

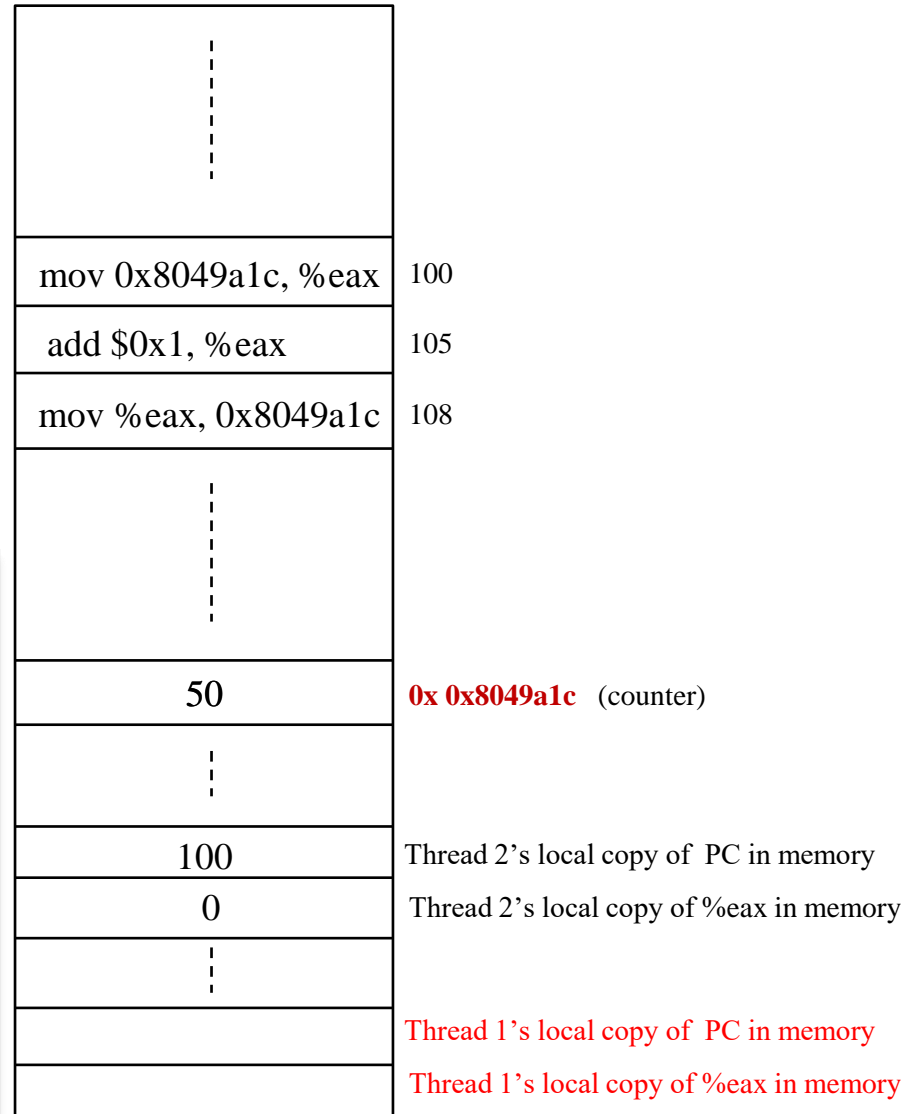


OS	Thread1	Thread2	(after instruction)		
			PC	%eax	counter
	Initial value		100	0	50
	mov 0x8049a1c, %eax		105	50	50
	add \$0x1, %eax		108	51	50
interrupt	save T1's state				
	restore T2's state		100	0	50
	mov 0x8049a1c, %eax		105	50	50
	add \$0x1, %eax		108	51	50
	mov %eax, 0x8049a1c		113	51	51
interrupt	save T2's state				
	restore T1's state		108	51	51
	mov %eax, 0x8049a1c		113	51	51

## Steps 2 & 3: Decode & Execution



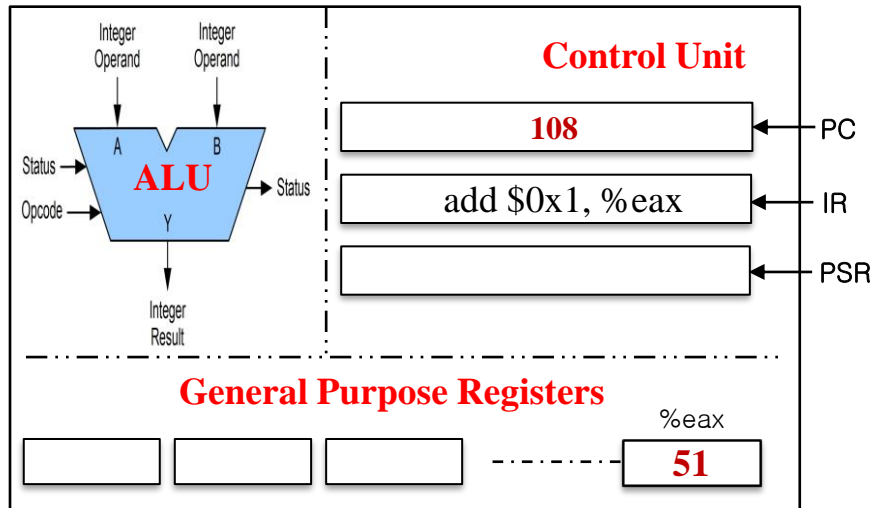
# Memory



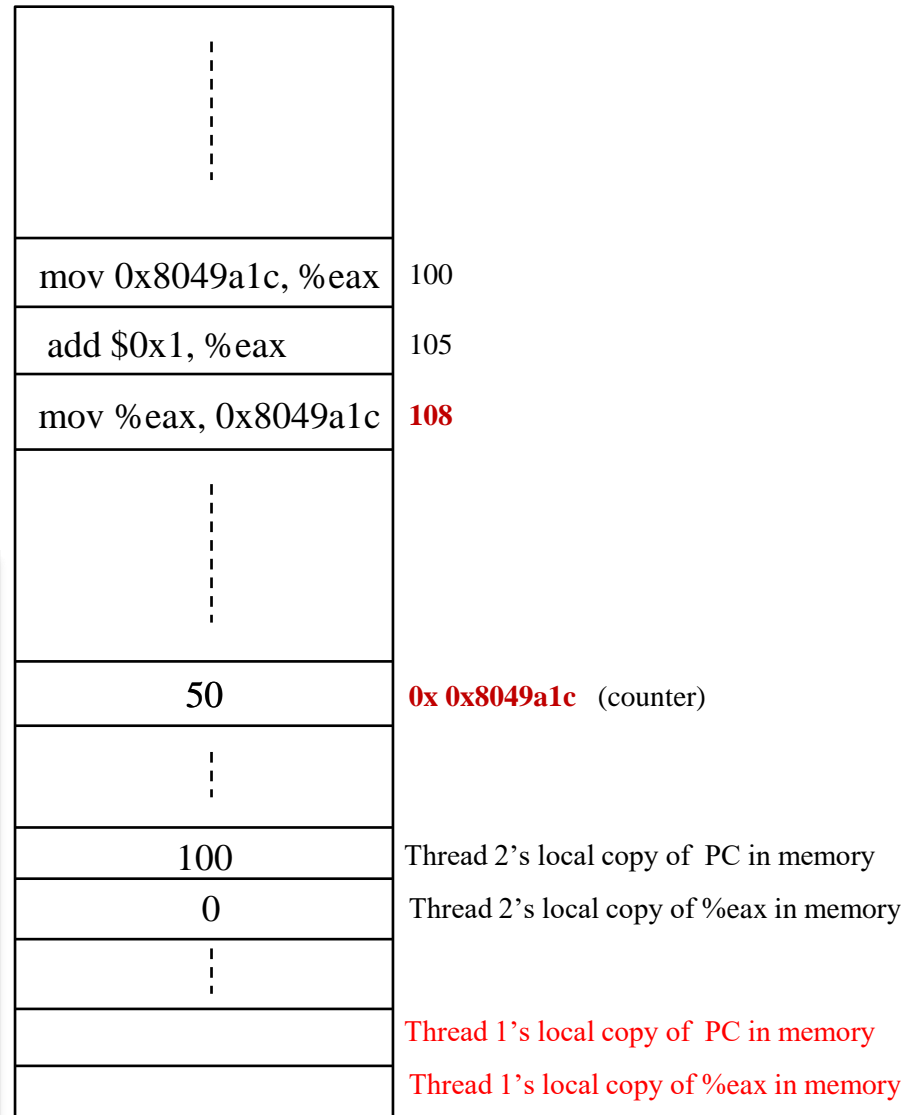
OS	Thread1	Thread2	PC	(after instruction) %eax	counter
	<b>Initial value</b>		100	0	50
	mov 0x8049a1c, %eax		105	50	50
	add \$0x1, %eax		108	51	50
interrupt					
save T1's state					
restore T2's state			100	0	50
		mov 0x8049a1c, %eax	105	50	50
		add \$0x1, %eax	108	51	50
		mov %eax, 0x8049a1c	113	51	51
interrupt					
save T2's state					
restore T1's state			108	51	51
	mov %eax, 0x8049a1c		113	51	51

# Step 3: Increase PC (pointed to the next instruction in the memory)

## CPU



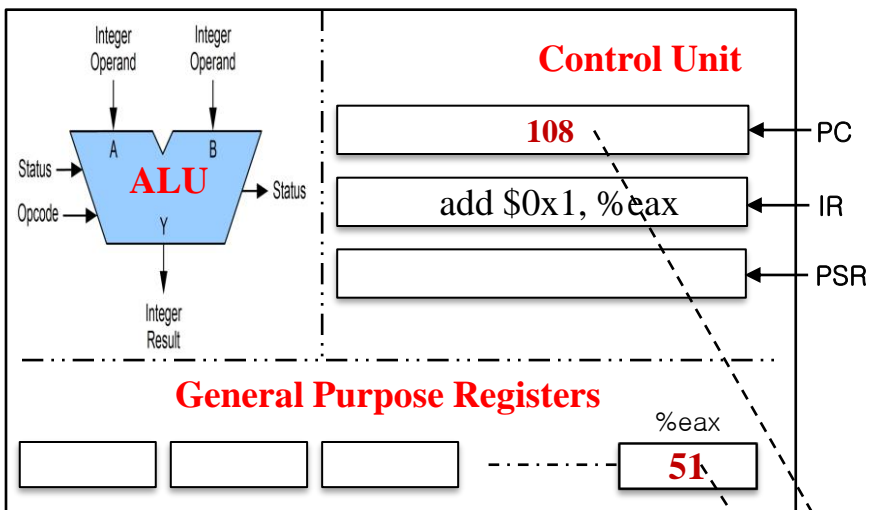
## Memory




OS	Thread1	Thread2	(after instruction)		
			PC	%eax	counter
	<b>Initial value</b>		100	0	50
	mov 0x8049a1c, %eax		105	50	50
	add \$0x1, %eax		108	51	50
interrupt	save T1's state				
	restore T2's state		100	0	50
	mov 0x8049a1c, %eax		105	50	50
	add \$0x1, %eax		108	51	50
	mov %eax, 0x8049a1c		113	51	51
interrupt	save T2's state				
	restore T1's state		108	51	51
	mov %eax, 0x8049a1c		113	51	51

# Interrupt occurred; Save Thread 1's state

## CPU



## Memory



OS	Thread1	Thread2	(after instruction)		
			PC	%eax	counter
	<b>Initial value</b>		100	0	50
	<u>mov</u> 0x8049a1c, <u>%eax</u>		105	50	50
	<u>add</u> \$0x1, <u>%eax</u>		108	51	50

interrupt

save T1's state

restore T2's state

<u>mov</u> 0x8049a1c, <u>%eax</u>	105	50	50
<u>add</u> \$0x1, <u>%eax</u>	108	51	50
<u>mov</u> %eax, 0x8049a1c	113	51	51

interrupt

save T2's state

restore T1's state

<u>mov</u> %eax, 0x8049a1c	108	51	51
	113	51	51

**0x 0x8049a1c** (counter)

Thread 2's local copy of PC in memory

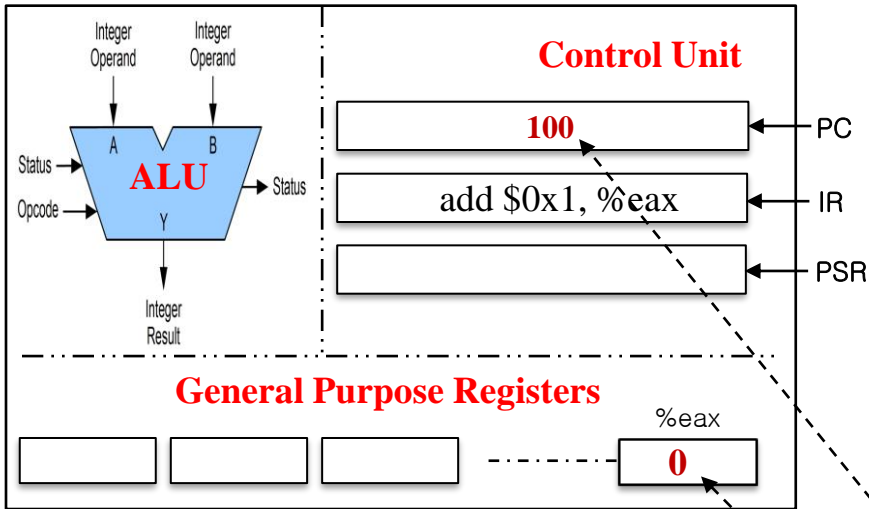
Thread 2's local copy of %eax in memory

Thread 1's local copy of PC in memory


Thread 1's local copy of %eax in memory

# Restore Thread 2's state

## CPU



## Memory



OS	Thread1	Thread2	PC	<u>%eax</u>	(after instruction) counter
	<b>Initial value</b>		100	0	50
	<u>mov</u> 0x8049a1c, <u>%eax</u>		105	50	50
	<u>add</u> \$0x1, <u>%eax</u>		108	51	50
interrupt	save T1's state				
	<b>restore T2's state</b>		100	0	50
		<u>mov</u> 0x8049a1c, <u>%eax</u>	105	50	50
		<u>add</u> \$0x1, <u>%eax</u>	108	51	50
		<u>mov</u> <u>%eax</u> , 0x8049a1c	113	51	51
interrupt	save T2's state				
	<b>restore T1's state</b>		108	51	51
		<u>mov</u> <u>%eax</u> , 0x8049a1c	113	51	51

100

105

108

0x 0x8049a1c (counter)

Thread 2's local copy of PC in memory

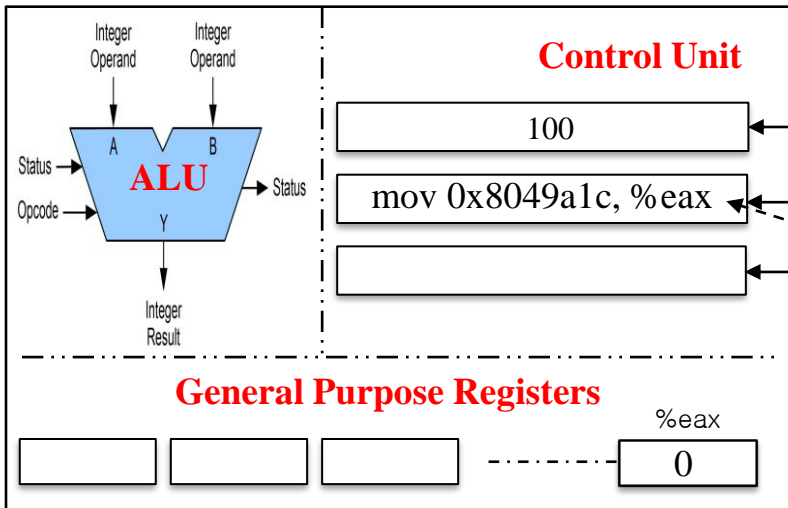
Thread 2's local copy of %eax in memory

Thread 1's local copy of PC in memory


Thread 1's local copy of %eax in memory

# Step 1: Fetch instruction from the memory (PC->Address=100)

## CPU



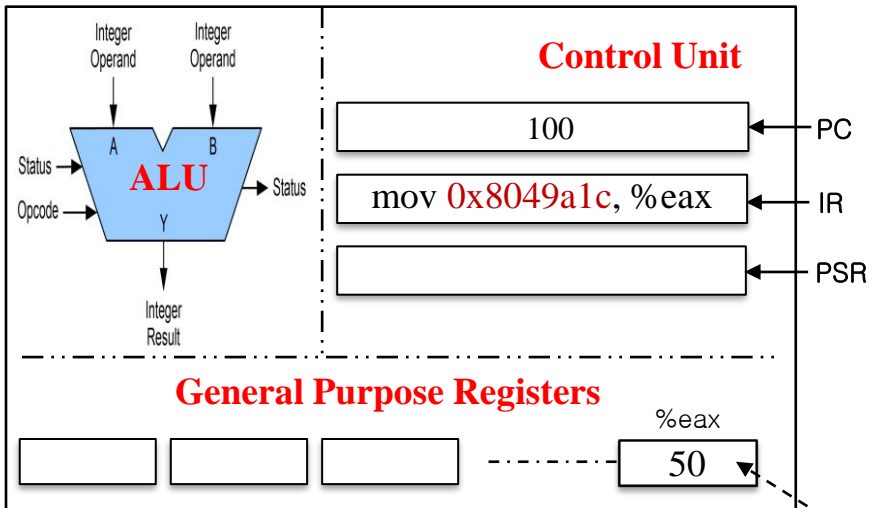
## Memory




OS	Thread1	Thread2	(after instruction)		
			PC	%eax	counter
	<b>Initial value</b>		100	0	50
	<u>mov</u> 0x8049a1c, <u>%eax</u>		105	50	50
	<u>add</u> \$0x1, <u>%eax</u>		108	51	50
interrupt	save T1's state				
	restore T2's state		100	0	50
	<u>mov</u> 0x8049a1c, <u>%eax</u>		105	50	50
	<u>add</u> \$0x1, <u>%eax</u>		108	51	50
	<u>mov</u> %eax, 0x8049a1c		113	51	51
interrupt	save T2's state				
	restore T1's state		108	51	51
	<u>mov</u> %eax, 0x8049a1c		113	51	51

# Steps 2 & 3: Decode & Execution

## CPU



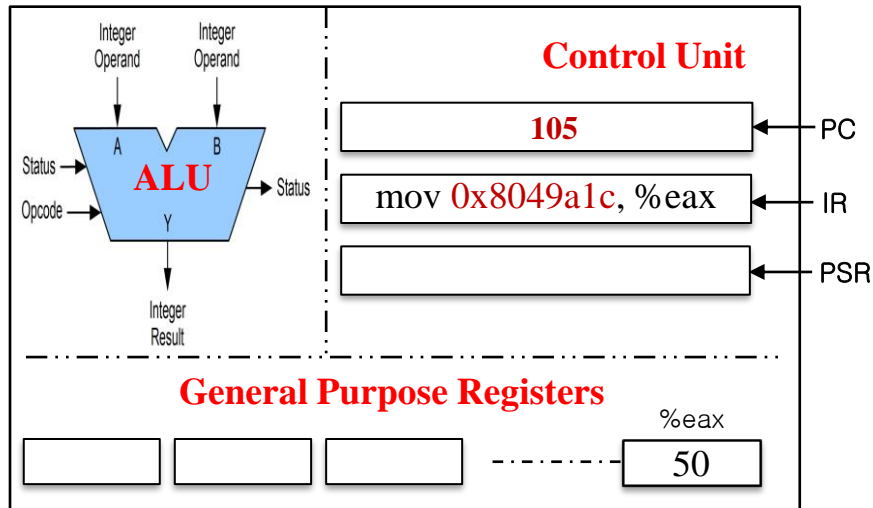
## Memory



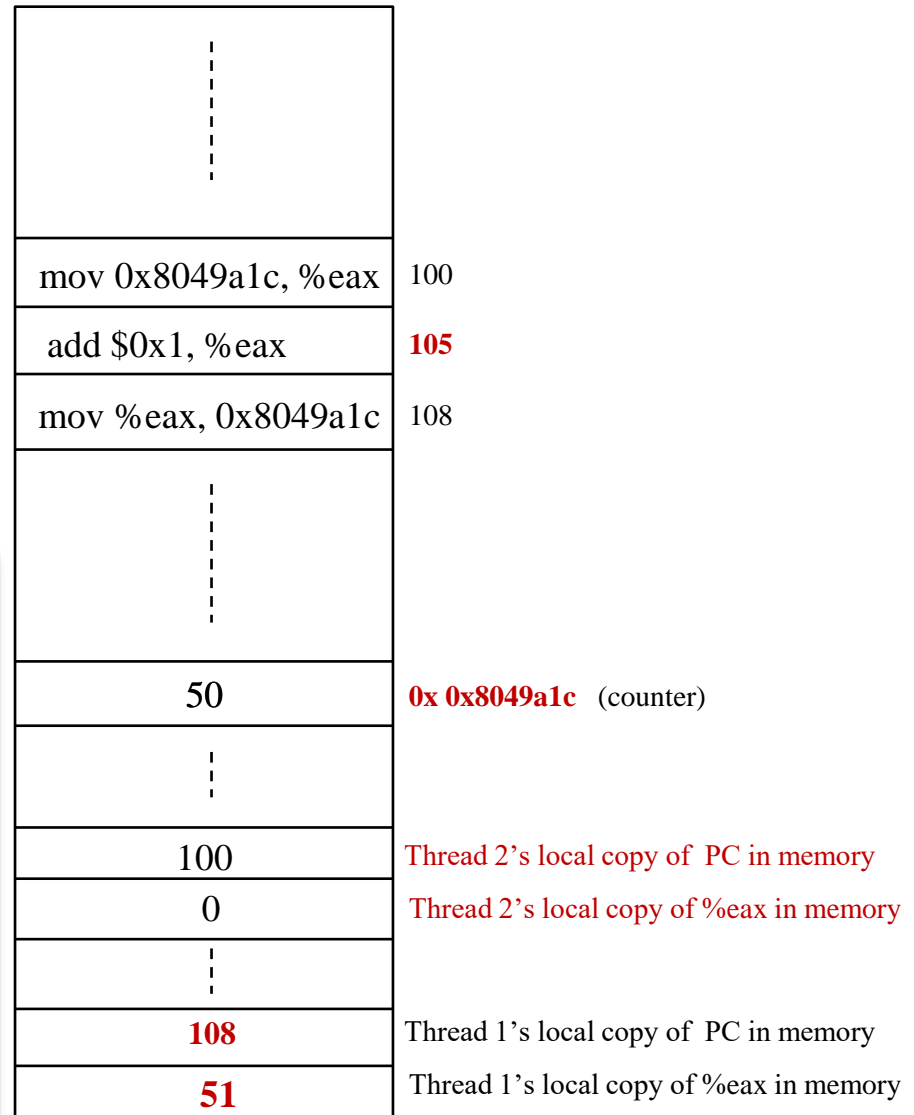
OS	Thread1	Thread2	(after instruction)		
			PC	%eax	counter
	<b>Initial value</b>		100	0	50
	<u>mov</u> 0x8049a1c, <u>%eax</u>		105	50	50
	<u>add</u> \$0x1, <u>%eax</u>		108	51	50
interrupt	save T1's state				
	restore T2's state		100	0	50
	<u>mov</u> 0x8049a1c, <u>%eax</u>		105	50	50
	<u>add</u> \$0x1, <u>%eax</u>		108	51	50
	<u>mov</u> %eax, 0x8049a1c		113	51	51
interrupt	save T2's state				
	restore T1's state		108	51	51
	<u>mov</u> %eax, 0x8049a1c		113	51	51

# Step 4: Increase PC (pointed to the next instruction in the memory)

## CPU



## Memory

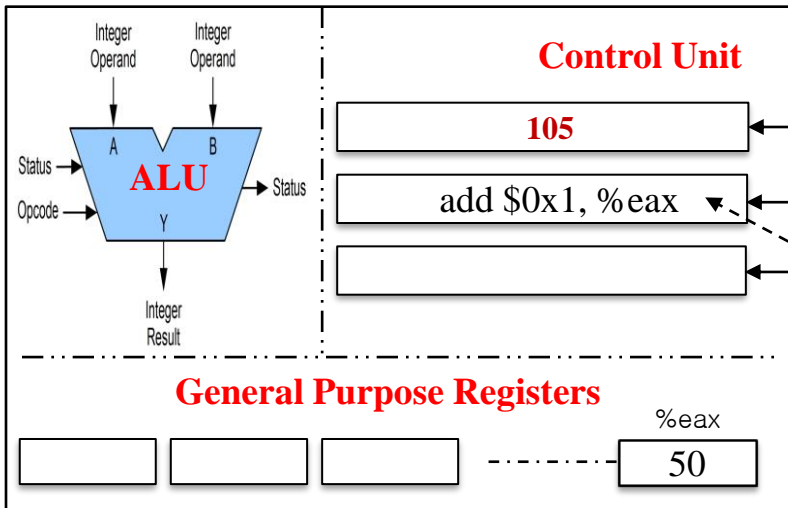


OS	Thread1	Thread2	(after instruction)		
			PC	<u>%eax</u>	counter
	<b>Initial value</b>		100	0	50
	<u>mov</u> 0x8049a1c, <u>%eax</u>		105	50	50
	add \$0x1, <u>%eax</u>		108	51	50
interrupt	save T1's state				
	restore T2's state		100	0	50
	<u>mov</u> 0x8049a1c, <u>%eax</u>		105	50	50
	add \$0x1, <u>%eax</u>		108	51	50
	<u>mov</u> <u>%eax</u> , 0x8049a1c		113	51	51
interrupt	save T2's state				
	restore T1's state		108	51	51
	<u>mov</u> <u>%eax</u> , 0x8049a1c		113	51	51




# Step 1: Fetch instruction from the memory (PC->Address=105)

## CPU



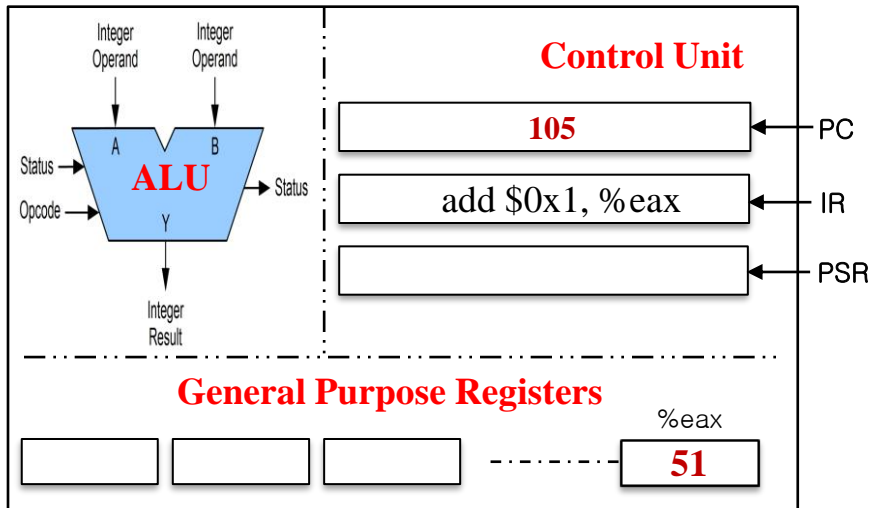
## Memory




OS	Thread1	Thread2	(after instruction)		
			PC	%eax	counter
	<b>Initial value</b>		100	0	50
	<u>mov</u> 0x8049a1c, <u>%eax</u>		105	50	50
	<u>add</u> \$0x1, <u>%eax</u>		108	51	50
interrupt	save T1's state				
	restore T2's state		100	0	50
		<u>mov</u> 0x8049a1c, <u>%eax</u>	105	50	50
		<u>add</u> \$0x1, <u>%eax</u>	108	51	50
		<u>mov</u> <u>%eax</u> , 0x8049a1c	113	51	51
interrupt	save T2's state				
	restore T1's state		108	51	51
		<u>mov</u> <u>%eax</u> , 0x8049a1c	113	51	51

# Steps 2 & 3: Decode & Execution

## CPU



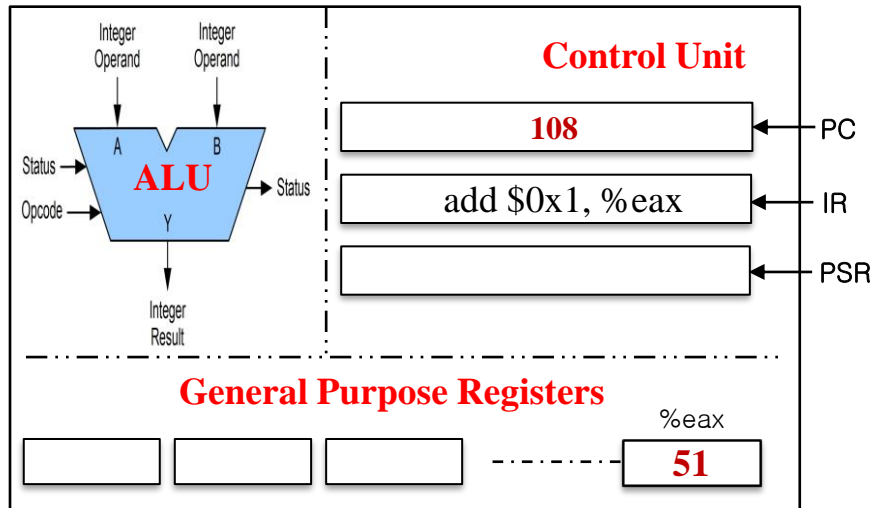
## Memory




OS	Thread1	Thread2	(after instruction)		
			PC	%eax	counter
	<b>Initial value</b>		100	0	50
	<u>mov</u> 0x8049a1c, <u>%eax</u>		105	50	50
	<u>add</u> \$0x1, <u>%eax</u>		108	51	50
interrupt	save T1's state				
	restore T2's state		100	0	50
		<u>mov</u> 0x8049a1c, <u>%eax</u>	105	50	50
		<u>add</u> \$0x1, <u>%eax</u>	108	51	50
		<u>mov</u> <u>%eax</u> , 0x8049a1c	113	51	51
interrupt	save T2's state				
	restore T1's state		108	51	51
		<u>mov</u> <u>%eax</u> , 0x8049a1c	113	51	51

# Step 4: Increase PC (pointed to the next instruction in the memory)

## CPU



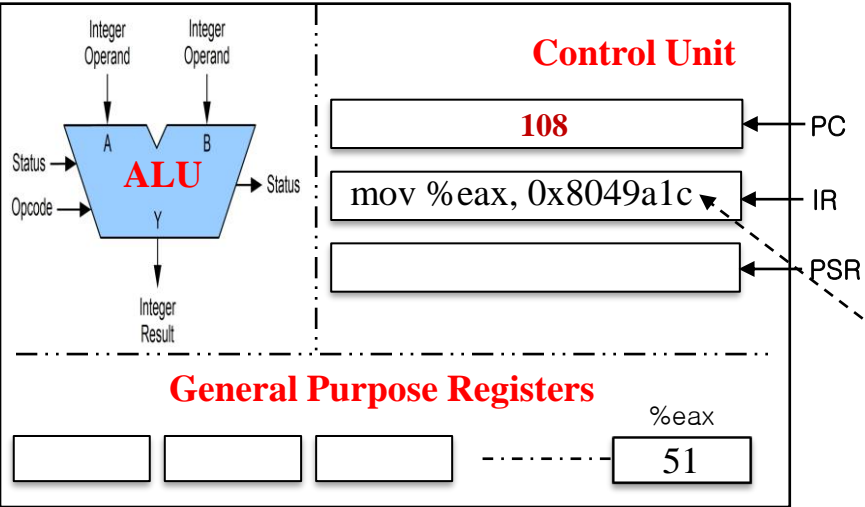
## Memory




OS	Thread1	Thread2	(after instruction)		
			PC	%eax	counter
	<b>Initial value</b>		100	0	50
	<u>mov</u> 0x8049a1c, <u>%eax</u>		105	50	50
	<u>add</u> \$0x1, <u>%eax</u>		108	51	50
interrupt	save T1's state				
	restore T2's state				
			100	0	50
		<u>mov</u> 0x8049a1c, <u>%eax</u>	105	50	50
		<u>add</u> \$0x1, <u>%eax</u>	108	51	50
		<u>mov</u> <u>%eax</u> , 0x8049a1c	113	51	51
interrupt	save T2's state				
	restore T1's state				
			108	51	51
	<u>mov</u> <u>%eax</u> , 0x8049a1c		113	51	51

# Step 1: Fetch instruction from the memory (PC->Address=108)

## CPU



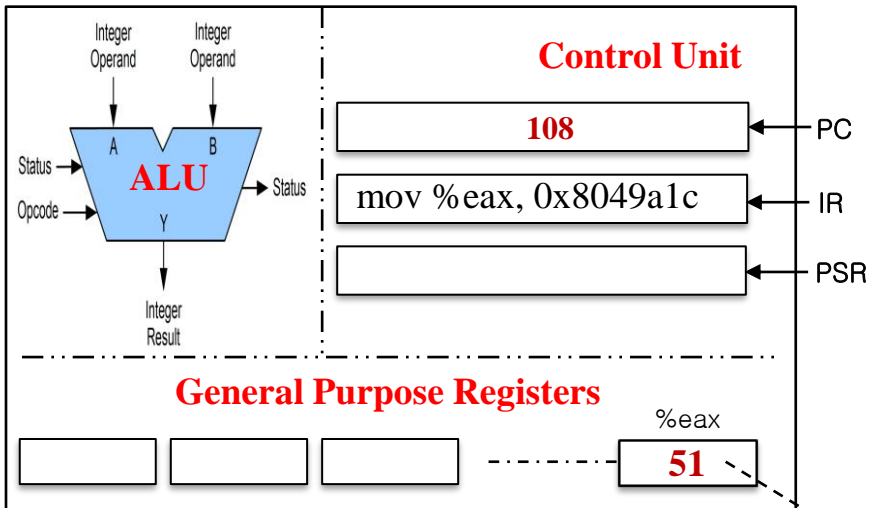
## Memory



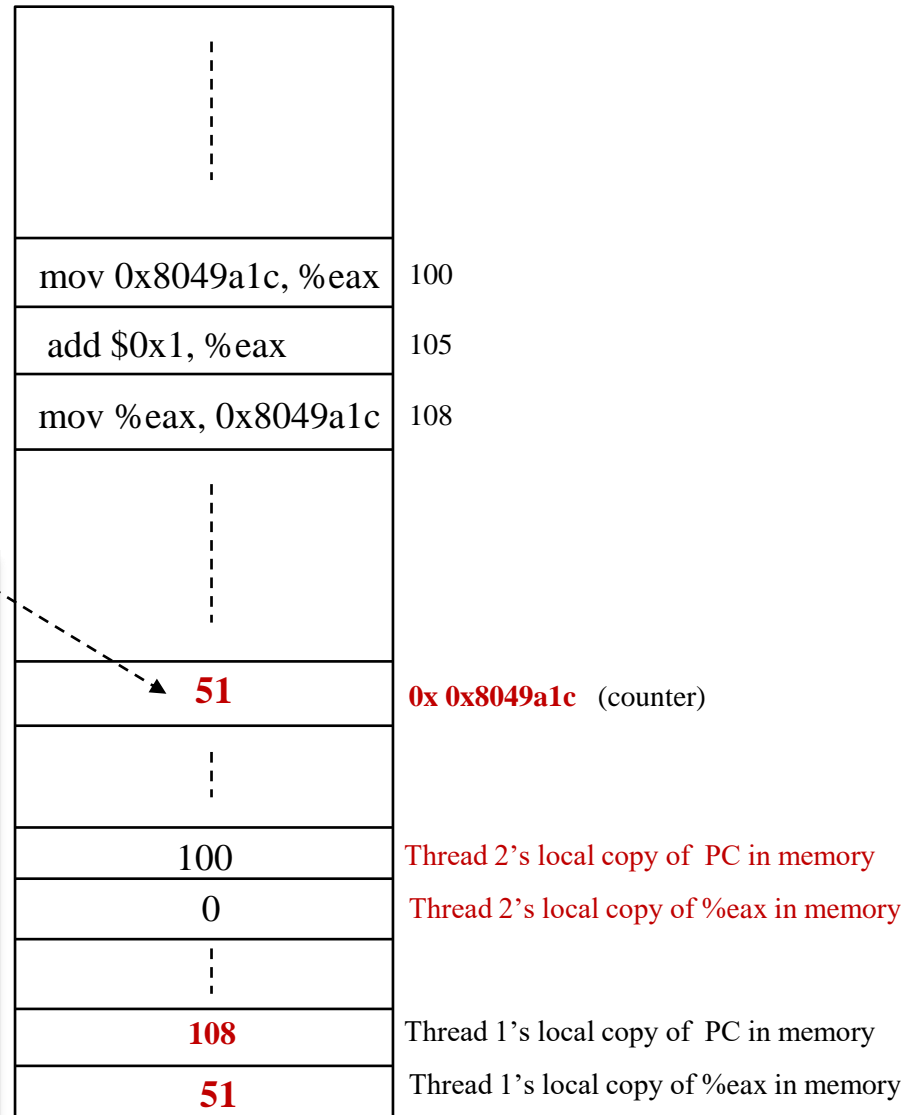
OS	Thread1	Thread2	(after instruction)		
			PC	%eax	counter
Initial value			100	0	50
mov 0x8049a1c, %eax			105	50	50
add \$0x1, %eax			108	51	50
interrupt					
save T1's state					
restore T2's state					
			100	0	50
mov 0x8049a1c, %eax			105	50	50
add \$0x1, %eax			108	51	50
mov %eax, 0x8049a1c			113	51	51
interrupt					
save T2's state					
restore T1's state					
			108	51	51
mov %eax, 0x8049a1c			113	51	51

# Steps 2 & 3: Decode & Execution

## CPU



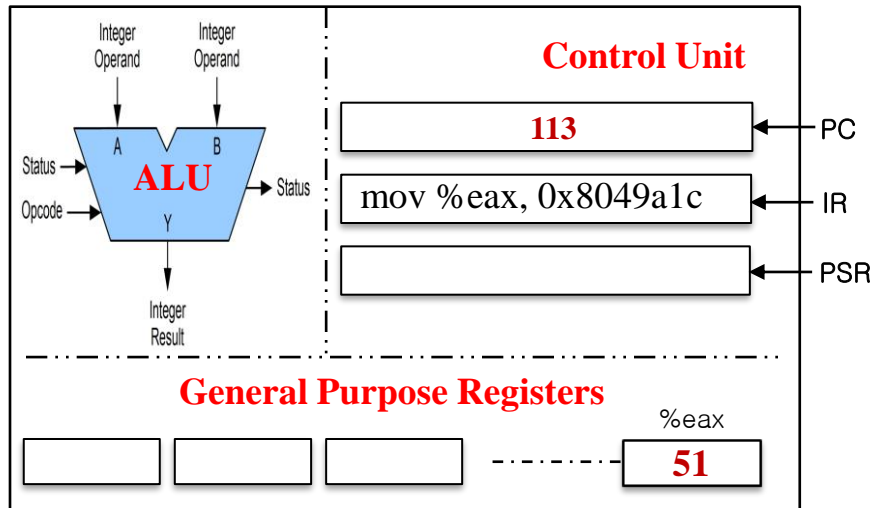
## Memory



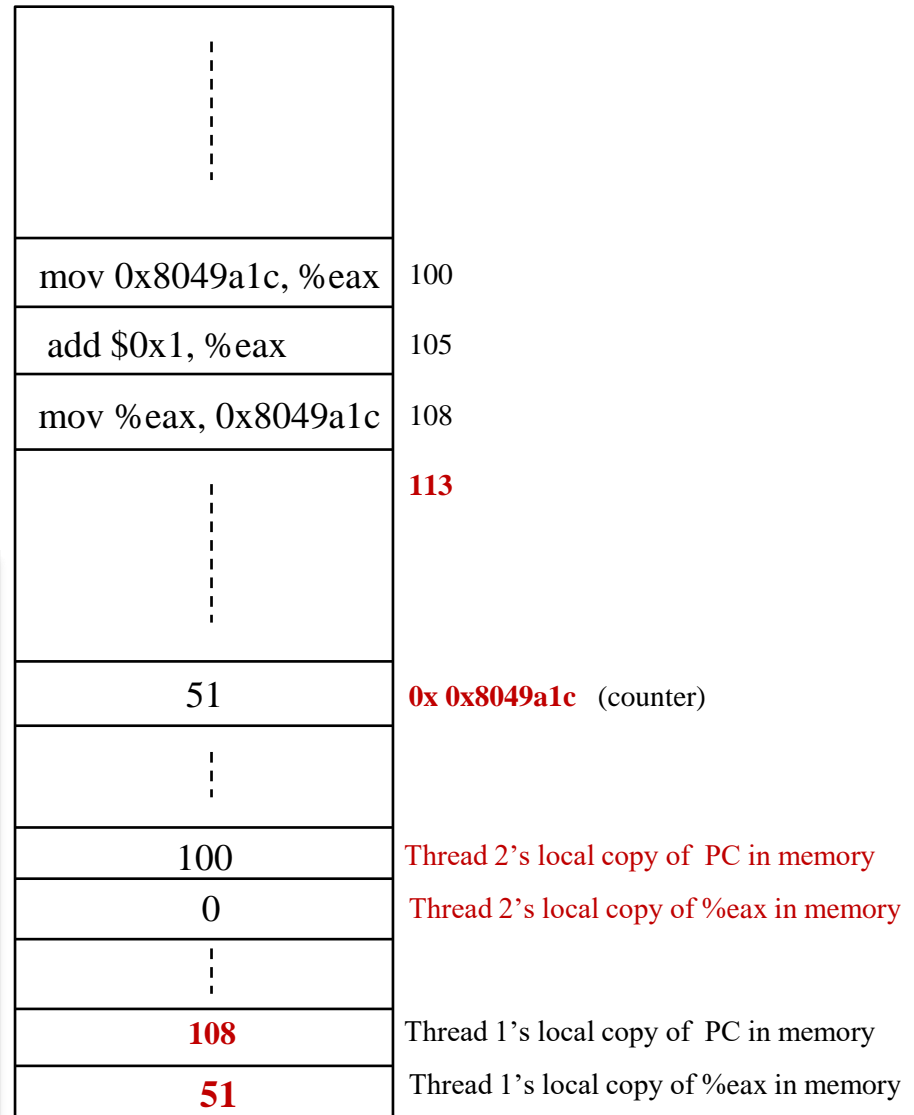
OS	Thread1	Thread2	(after instruction)		
			PC	% <u>eax</u>	counter
	<b>Initial value</b>		100	0	50
	<code>mov 0x8049a1c, %<u>eax</u></code>		105	50	50
	<code>add \$0x1, %<u>eax</u></code>		108	51	50
interrupt	save T1's state				
	restore T2's state				
			100	0	50
	<code>mov 0x8049a1c, %<u>eax</u></code>		105	50	50
	<code>add \$0x1, %<u>eax</u></code>		108	51	50
	<code>mov %<u>eax</u>, 0x8049a1c</code>		113	51	51
interrupt	save T2's state				
	restore T1's state				
			108	51	51
	<code>mov %<u>eax</u>, 0x8049a1c</code>		113	51	51

# Step 4: Increase PC (pointed to the next instruction in the memory)

## CPU



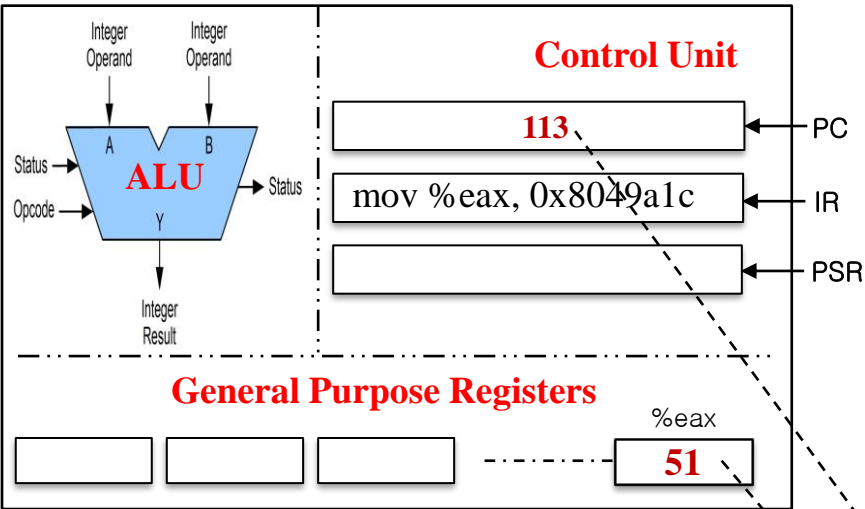
## Memory




OS	Thread1	Thread2	(after instruction)		
			PC	<u>%eax</u>	counter
	<b>Initial value</b>		100	0	50
	<code>mov 0x8049a1c, <u>%eax</u></code>		105	50	50
	<code>add \$0x1, <u>%eax</u></code>		108	51	50
interrupt	save T1's state				
	restore T2's state				
			100	0	50
	<code>mov 0x8049a1c, <u>%eax</u></code>		105	50	50
	<code>add \$0x1, <u>%eax</u></code>		108	51	50
	<code>mov <u>%eax</u>, 0x8049a1c</code>		113	51	51
interrupt	save T2's state				
	restore T1's state				
			108	51	51
	<code>mov <u>%eax</u>, 0x8049a1c</code>		113	51	51

# Interrupt occurred; Save Thread 2's state

## CPU

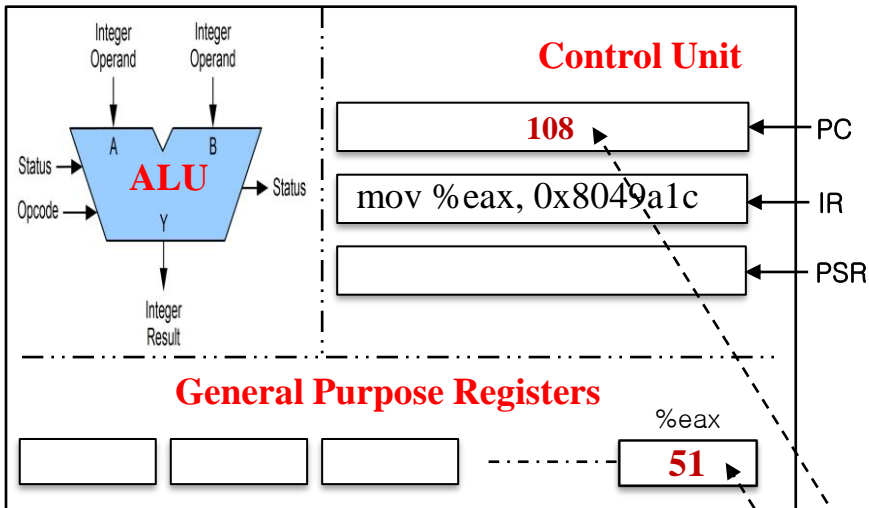


## Memory




OS	Thread1	Thread2	(after instruction)		
			PC	%eax	counter
	Initial value		100	0	50
	mov 0x8049a1c, %eax		105	50	50
	add \$0x1, %eax		108	51	50
Interrupt					
	save T1's state				
	restore T2's state		100	0	50
		mov 0x8049a1c, %eax	105	50	50
		add \$0x1, %eax	108	51	50
		mov %eax, 0x8049a1c	113	51	51
Interrupt					
	save T2's state				
	restore T1's state		108	51	51
	mov %eax, 0x8049a1c		113	51	51

## Restore Thread 1's state



## Memory

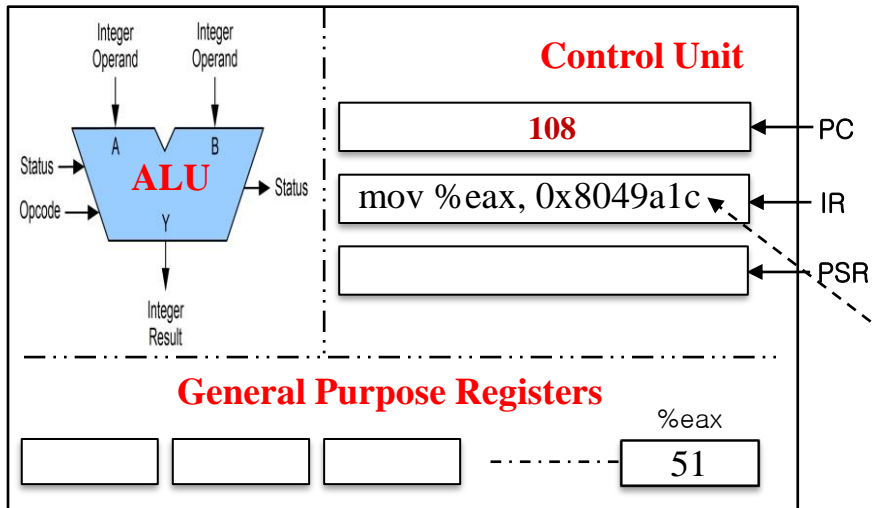


OS	Thread1	Thread2	(after instruction)		
			PC	%eax	counter
	<b>Initial value</b>		100	0	50
	<u>mov</u> 0x8049a1c, <u>%eax</u>		105	50	50
	<u>add</u> \$0x1, <u>%eax</u>		108	51	50
interrupt					
	save T1's state				
	restore T2's state		100	0	50
		<u>mov</u> 0x8049a1c, <u>%eax</u>	105	50	50
		<u>add</u> \$0x1, <u>%eax</u>	108	51	50
		<u>mov</u> <u>%eax</u> , 0x8049a1c	113	51	51
interrupt					
	save T2's state				
	restore T1's state		108	51	51
	<u>mov</u> <u>%eax</u> , 0x8049a1c		113	51	51

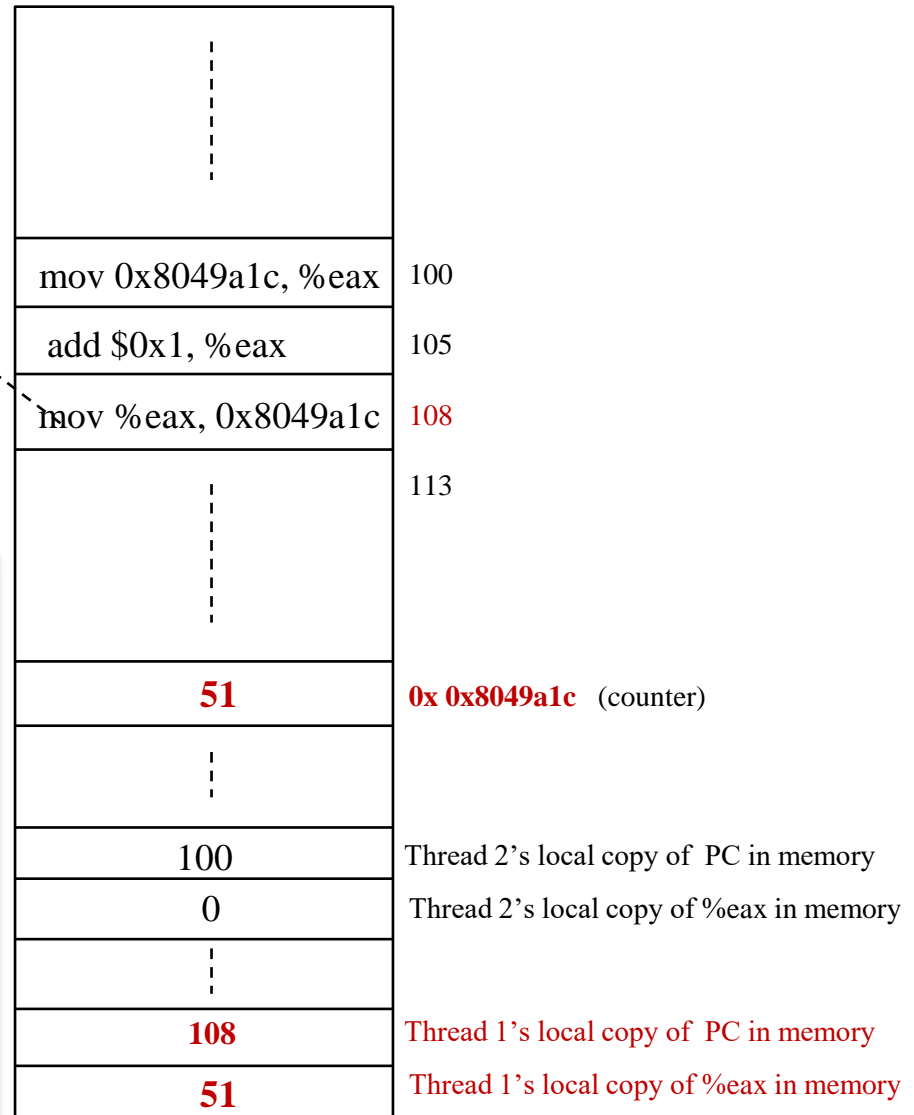


# Step 1: Fetch instruction from the memory (PC->Address=108)

## CPU



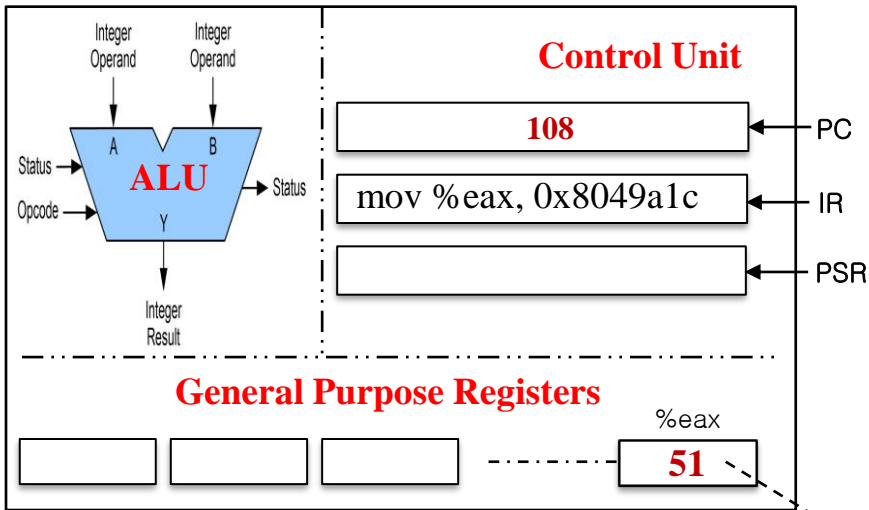
## Memory



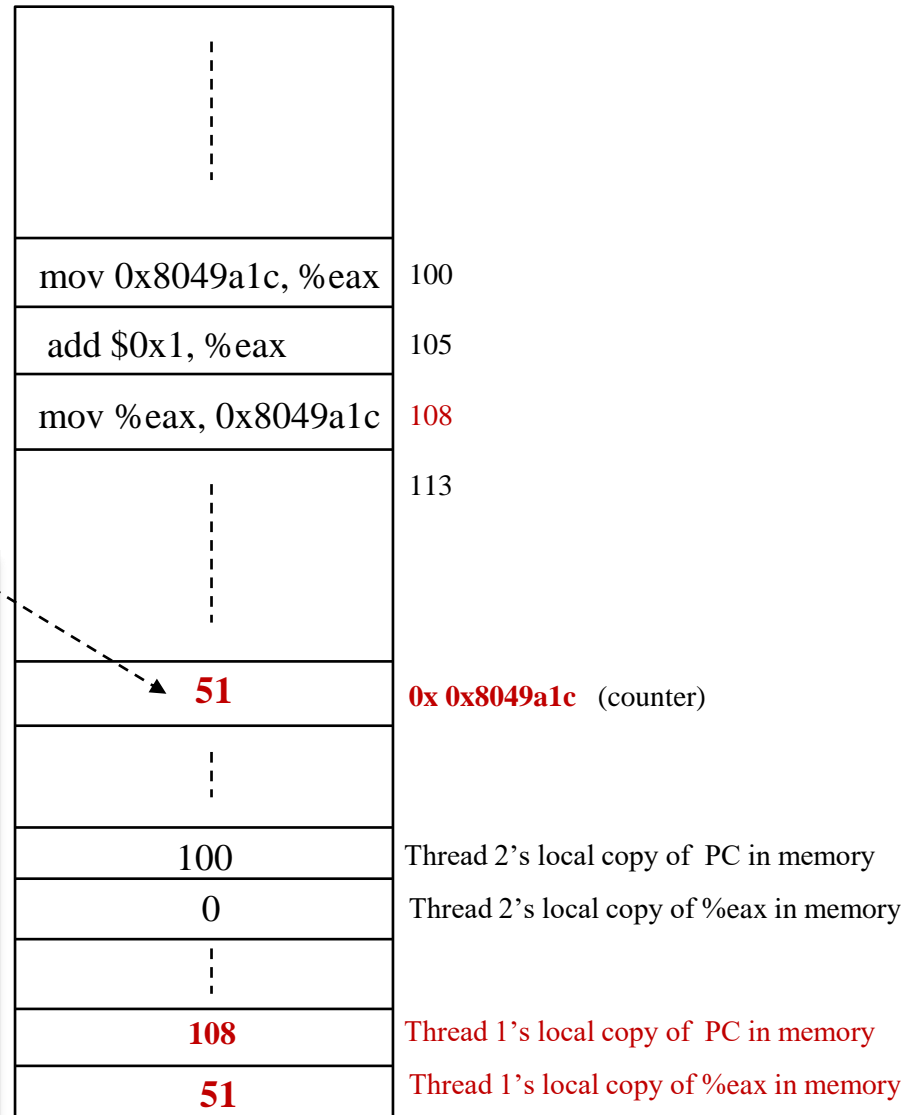
OS	Thread1	Thread2	(after instruction)		
			PC	%eax	counter
	<b>Initial value</b>		100	0	50
	<code>mov 0x8049a1c, %eax</code>		105	50	50
	<code>add \$0x1, %eax</code>		108	51	50
interrupt					
save T1's state					
restore T2's state			100	0	50
	<code>mov 0x8049a1c, %eax</code>		105	50	50
	<code>add \$0x1, %eax</code>		108	51	50
	<code>mov %eax, 0x8049a1c</code>		113	51	51
interrupt					
save T2's state					
restore T1's state			108	51	51
	<code>mov %eax, 0x8049a1c</code>		113	51	51

# Step 2 & 3: Decode & Execution

## CPU



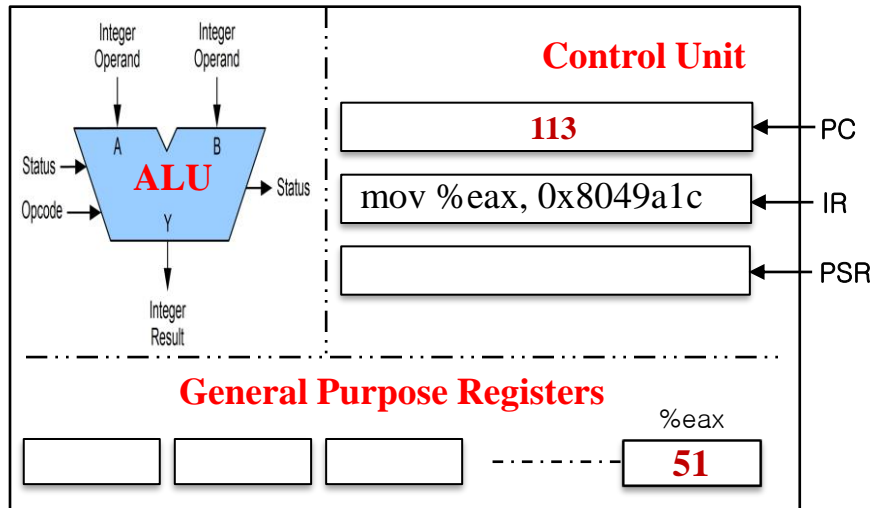
## Memory




OS	Thread1	Thread2	(after instruction)		
			PC	%eax	counter
	<b>Initial value</b>		100	0	50
	<code>mov 0x8049a1c, %eax</code>		105	50	50
	<code>add \$0x1, %eax</code>		108	51	50
interrupt					
save T1's state					
restore T2's state			100	0	50
	<code>mov 0x8049a1c, %eax</code>		105	50	50
	<code>add \$0x1, %eax</code>		108	51	50
	<code>mov %eax, 0x8049a1c</code>		113	51	51
interrupt					
save T2's state					
restore T1's state			108	51	51
	<code>mov %eax, 0x8049a1c</code>		113	51	51

# Step 4: Increase PC (pointed to the next instruction in the memory)

## CPU



## Memory



OS	Thread1	Thread2	(after instruction)		
			PC	<u>%eax</u>	counter
	<b>Initial value</b>		100	0	50
	<u>mov</u> 0x8049a1c, <u>%eax</u>		105	50	50
	<u>add</u> \$0x1, <u>%eax</u>		108	51	50
interrupt					
save T1's state					
restore T2's state			100	0	50
	<u>mov</u> 0x8049a1c, <u>%eax</u>		105	50	50
	<u>add</u> \$0x1, <u>%eax</u>		108	51	50
	<u>mov</u> <u>%eax</u> , 0x8049a1c		113	51	51
interrupt					
save T2's state					
restore T1's state			108	51	51
	<u>mov</u> <u>%eax</u> , 0x8049a1c		113	51	51

# Critical section

- ❑ A piece of code that **accesses a shared variable** and must not be concurrently executed by more than one thread.
  - ◆ Multiple threads executing critical section can result in a race condition.
  - ◆ Need to support **atomicity** for critical sections (**mutual exclusion**)

# Locks

- Ensure that any such critical section executes as if it were a single atomic instruction (**execute a series of instructions atomically**).

```
1  lock_t mutex;  
2  . . .  
3  lock(&mutex);  
4  balance = balance + 1;  
5  unlock(&mutex);
```



Critical section

## **Part II: Thread API**

---

# Thread Creation

## □ How to create and control threads?

```
#include <pthread.h>

int
pthread_create(      pthread_t*      thread,
                    const pthread_attr_t* attr,
                    void*             (*start_routine)(void*),
                    void*             arg);
```

- ◆ **thread**: Used to interact with this thread.
- ◆ **attr**: Used to specify any attributes this thread might have.
  - Stack size, Scheduling priority, ...
- ◆ **start\_routine**: the function this thread start running in.
- ◆ **arg**: the argument to be passed to the function (start routine)
  - *a void pointer* allows us to pass in *any type of* argument.

# Example: Creating a Thread

```
#include <pthread.h>

typedef struct __myarg_t {
    int a;
    int b;
} myarg_t;

void *mythread(void *arg) {
    myarg_t *m = (myarg_t *) arg;
    printf("%d %d\n", m->a, m->b);
    return NULL;
}

int main(int argc, char *argv[]) {
    pthread_t p;
    int rc;

    myarg_t args;
    args.a = 10;
    args.b = 20;
    rc = pthread_create(&p, NULL, mythread, &args);
    ...
}
```



# Wait for a thread to complete

```
int pthread_join(pthread_t thread, void **value_ptr);
```

- ◆ thread: Specify which thread *to wait for*
- ◆ value\_ptr: A pointer to the return value
  - Because pthread\_join() routine changes the value, the return value needs to be **passed in a pointer** to that value.

# Example: Waiting for Thread Completion

```
1  #include <stdio.h>
2  #include <pthread.h>
3  #include <assert.h>
4  #include <stdlib.h>
5
6  typedef struct __myarg_t {
7      int a;
8      int b;
9  } myarg_t;
10
11 typedef struct __myret_t {
12     int x;
13     int y;
14 } myret_t;
15
16 void *mythread(void *arg) {
17     myarg_t *m = (myarg_t *) arg;
18     printf("%d %d\n", m->a, m->b);
19     myret_t *r = malloc(sizeof(myret_t));
20     r->x = 1;
21     r->y = 2;
22     return (void *) r;
23 }
24
```

## Example: Waiting for Thread Completion (Cont.)

```
25  int main(int argc, char *argv[]) {
26      int rc;
27      pthread_t p;
28      myret_t *m;
29
30      myarg_t args;
31      args.a = 10;
32      args.b = 20;
33      pthread_create(&p, NULL, mythread, &args);
34      pthread_join(p, (void **) &m); // this thread has been
                                     // waiting inside of the
                                     // pthread_join() routine.
35      printf("returned %d %d\n", m->x, m->y);
36      return 0;
37  }
```

## Example: Dangerous code

- ❑ Be careful with how values are returned from a thread.

```
1  void *mythread(void *arg) {
2      myarg_t *m = (myarg_t *) arg;
3      printf("%d %d\n", m->a, m->b);
4      myret_t r; // ALLOCATED ON STACK: BAD!
5      r.x = 1;
6      r.y = 2;
7      return (void *) &r;
8  }
```

- ◆ When the variable `r` returns, it is automatically **de-allocated**.

# Example: Simpler Argument Passing to a Thread

## ▣ Just passing in a single value

```
1  void *mythread(void *arg) {
2      int m = (int) arg;
3      printf("%d\n", m);
4      return (void *) (arg + 1);
5  }
6
7  int main(int argc, char *argv[]) {
8      pthread_t p;
9      int rc, m;
10     pthread_create(&p, NULL, mythread, (void *) 100);
11     pthread_join(p, (void **) &m);
12     printf("returned %d\n", m);
13     return 0;
14 }
```

# Locks

## ▣ Provide **mutual exclusion** to a critical section

### ◆ Interface

```
int pthread_mutex_lock(pthread_mutex_t *mutex);  
int pthread_mutex_unlock(pthread_mutex_t *mutex);
```

### ◆ Usage (w/o *lock initialization* and *error check*)

```
pthread_mutex_t lock;  
pthread_mutex_lock(&lock);  
x = x + 1; // or whatever your critical section is  
pthread_mutex_unlock(&lock);
```

- No other thread holds the lock → the thread will acquire the lock and **enter the critical section**.
- If another thread hold the lock → the thread will **not return from the call** until it has acquired the lock.

# Locks (Cont.)

- All locks must be properly initialized.
  - ◆ One way: using `PTHREAD_MUTEX_INITIALIZER`

```
pthread_mutex_t lock = PTHREAD_MUTEX_INITIALIZER;
```

- ◆ The dynamic way: using `pthread_mutex_init()`

```
int rc = pthread_mutex_init(&lock, NULL);  
assert(rc == 0); // always check success!
```

# Locks (Cont.)

- ▣ Check errors code when calling lock and unlock
  - ◆ An example wrapper

```
// Use this to keep your code clean but check for failures
// Only use if exiting program is OK upon failure
void Pthread_mutex_lock(pthread_mutex_t *mutex) {
    int rc = pthread_mutex_lock(mutex);
    assert(rc == 0);
}
```



# Locks (Cont.)

- ▣ These two calls are also used in **lock acquisition**

```
int pthread_mutex_trylock(pthread_mutex_t *mutex);  
int pthread_mutex_timelock(pthread_mutex_t *mutex,  
                           struct timespec *abs_timeout);
```

- ◆ trylock: return failure if the lock is already held
- ◆ timelock: return after a timeout or after acquiring the lock

# Condition Variables

- ❑ **Condition variables** are useful when some kind of **signaling** must take place between threads.

```
int pthread_cond_wait(pthread_cond_t *cond,  
                      pthread_mutex_t *mutex);  
int pthread_cond_signal(pthread_cond_t *cond);
```

- ◆ `pthread_cond_wait`:
  - Put the calling thread to sleep.
  - Wait for some other thread to signal it.
- ◆ `pthread_cond_signal`:
  - Unblock at least one of the threads that are blocked on the condition variable

# Condition Variables (Cont.)

## ▣ A thread calling wait routine:

```
pthread_mutex_t lock = PTHREAD_MUTEX_INITIALIZER;  
pthread_cond_t init = PTHREAD_COND_INITIALIZER;  
  
pthread_mutex_lock(&lock);  
while (initialized == 0)  
    pthread_cond_wait(&init, &lock);  
pthread_mutex_unlock(&lock);
```

- ◆ The wait call **releases the lock** when putting said caller to sleep.
- ◆ Before returning after being woken, the wait call **re-acquires the lock**.

## ▣ A thread calling signal routine:

```
pthread_mutex_lock(&lock);  
initialized = 1;  
pthread_cond_signal(&init);  
pthread_mutex_unlock(&lock);
```

## Condition Variables (Cont.)

- The waiting thread **re-checks** the condition **in a while loop**, instead of a simple **if** statement.

```
pthread_mutex_t lock = PTHREAD_MUTEX_INITIALIZER;
pthread_cond_t init = PTHREAD_COND_INITIALIZER;

pthread_mutex_lock(&lock);
while (initialized == 0)
    pthread_cond_wait(&init, &lock);
pthread_mutex_unlock(&lock);
```

- ◆ Without rechecking, the waiting thread will continue thinking that the condition has changed even though it has not (the condition can be changed before an awakened thread returns from `pthread_cond_wait()`).

# Condition Variables (Cont.)

- ❑ Don't ever to this.

- ◆ A thread calling wait routine:

```
while(initialized == 0)  
    ; // spin
```

- ◆ A thread calling signal routine:

```
initialized = 1;
```

- ◆ It performs poorly in many cases. → just wastes CPU cycles.
- ◆ It is error prone.

# Compiling and Running

- ▣ To compile them, you must include the header `pthread.h`
  - ◆ Explicitly link with the **pthread library**, by adding the `-pthread` flag.

```
prompt> gcc -o main main.c -Wall -pthread
```

- ◆ For more information,

```
man -k pthread
```

# Summary

- ❑ A multi-threaded program
  - ◆ More than one point of execution (with multiple PCs)
  - ◆ Multiple threads **share** the same **address space** but with different **stacks**
- ❑ Race condition may occur when multiple threads enter critical sections.
- ❑ Thread API
  - ◆ Thread creation: `pthread_create`
  - ◆ Thread wait: `pthread_join()`
  - ◆ Lock & unlock: `pthread_mutex_lock()` / `pthread_mutex_unlock()`
  - ◆ Conditional variable: `pthread_cond_wait()` / `pthread_cond_signal()`
- ❑ Next: Semaphore